Genesis OS / Relationship Scores

Cluster Rule Pack Files - Batch 1

Implementation sheets for the first governed clusters named by Cluster Rule Pack Specification v1

What this pack is

This document is the first cluster-by-cluster companion set that sits beside the governing specification. It converts the six initial governed clusters into implementable rule files with trigger posture, severity mapping, write-back behavior, decay, readiness impact, forecast impact, drift interlocks, and QA test expectations.

Status	Implementation-ready rule-pack batch for governed v1 launch
Depends on	Cluster Rule Pack Specification v1, Readiness & Forecast Contract v1 and Part 2, DriftEligibility appendix, workbook-governed threshold references
Clusters included	Vulnerability Loop, Control Avoidance, Slow Drift, Push-Pull, Repair Gap, Readiness Overstatement
Versioning rule	Any trigger, threshold, or write-back change requires a version tag, migration note, and parity test against historical data

Batch intent and implementation notes

- These files assume the universal severity ladder and write-back classes already defined by the governing cluster specification.
- Thresholds below are intentionally explicit enough for engineering, but they should still be cross-referenced to the canonical threshold workbook once the service constants are frozen.
- A cluster may open in memory before it writes back. Public or partner-visible surfaces must only react to persisted write-back states allowed by policy.
- No cluster in this batch is allowed to rewrite baseline archetype or baseline shadow assignments. They govern current operational state, not identity recasting.

Cluster	Core signal	Primary scope	Main downstream systems
CL-101	closeness sought, then contact retreats when vulnerability becomes real	Individual and pair	Readiness, forecast, Sanctuary, AI coaching
CL-102	autonomy defense pattern that blocks repair, structure, or mutual planning	Individual and pair	Readiness, forecast, pacing, Knock / messaging guidance
CL-103	slow-moving deterioration in connection quality across a governed rollup window	Pair-state with user write-back mirrors	Forecast, pair guidance, drift ledger, guarded mode surfaces
CL-104	alternating pursuit and withdrawal cycle that destabilizes reciprocity	Pair with mirrored user-state traces	Forecast, messaging interventions, Sanctuary reflections
CL-105	conflict occurs but repair behavior is absent, delayed, or collapses	Pair and individual behavior trace	Forecast, readiness confidence, AI intervention timing
CL-106	claimed readiness exceeds evidence-supported readiness or current behavioral stability	Individual	Public readiness display, AI coaching, eligibility gates

CL-101 - Vulnerability Loop

Family	Contradiction / attachment reflection cluster
Scope	Individual first; pair-visible effect allowed only

	through softened downstream signals
Purpose	Detect a recurring pattern where the user expresses desire for closeness, depth, honesty, or reassurance, then pulls back when intimacy becomes emotionally real or demands stable reciprocity.
Primary sources	Assessment contradictions (for example openness vs retreat answers), message-response retreat events, bid acceptance followed by avoidance, reflection completions, repair participation.
Default posture	High-confidence contradiction cluster with controlled early warning; should open cautiously and escalate only when repetition or cross-source agreement exists.
User exposure	Softened user-facing language allowed; raw cluster ID remains internal.

Canonical detection logic

- Open S1 when a workbook-approved contradiction pair suggests closeness-seeking plus retreat-protective posture, or when one behavioral retreat event follows a strong closeness bid within the recent window.
- Escalate to S2 when at least two qualifying events occur in 30 days, or one assessment contradiction aligns with one behavioral retreat event.
- Escalate to S3 when the pattern repeats after reassurance, after a vulnerable conversation, or after the user explicitly asks for closeness then reduces continuity.
- Escalate to S4 when the cycle is repeated across at least two separate connection episodes or when it materially drags repair, reciprocity, or pair stability.
- S5 is reserved for persistent and reinforced patterning that keeps producing rupture-level instability despite reflection prompts and sufficient opportunity for repair.

Severity ladder for this cluster

Lvl	Trigger Posture	Default Write-Back	Primary Effect
S1	watch only	ephemeral in-memory flag or low-weight user trace	nudges confidence, no public penalty
S2	light contradiction	user_cluster_state opens;	small readiness drag,

	confirmed	reflection prompt eligible	forecast confidence watch
S3	active repetition	persistent user write-back with 30d TTL	moderate readiness penalty and pair caution
S4	repeating destabilizer	persistent user write-back plus pair caution trace	stronger readiness drag; forecast guarded
S5	critical unresolved loop	persistent user and pair protection flags	suppress numeric exposure; intervention required

Area	Rule
Write-back	Store cluster state on the user record with source hash references, contradiction pair IDs, last three reinforcing event IDs, confidence band, and allowed consumer services. Pair write-back is permitted only from S4 upward and must use softened labels rather than the raw cluster name.
Decay / clear	Downgrade one level after 21 days with no reinforcing event plus one positive continuity or repair signal. Hard clear after 60 quiet days or after two approved reflection completions plus stable reciprocity evidence. Do not clear during active slow drift.
Readiness effect	S2 may reduce readiness confidence without moving the public band. S3 applies a moderate contradiction penalty. S4-S5 suppress raw numeric readiness unless confidence remains high and policy explicitly allows it.
Forecast effect	From S3 upward, increase contradiction load and reciprocity strain. S4-S5 may downgrade pair forecast to guarded mode even when compatibility remains strong.
Drift interlock	Do not convert a short-lived loop into drift by itself. Vulnerability Loop can reinforce drift only when the same retreat pattern persists across the governed drift rollup window and DriftEligibility is satisfied.
AI / Sanctuary use	AI may translate this cluster into coaching about wanting closeness but feeling overloaded once closeness arrives. Sanctuary can offer grounding, pacing, and repair reflection prompts, but must not

	shame or diagnose.

QA examples to lock before production

Test at least four cases: (1) assessment contradiction only, no behavior; should not exceed S2. (2) repeated post-vulnerability retreat; should reach S3. (3) healthy pause followed by clear repair; should decay rather than escalate. (4) one bad week during acute grief; should lower confidence and avoid over-penalizing.

CL-102 - Control Avoidance

Family	Autonomy defense / pacing conflict cluster
Scope	Individual with pair-state consequences when planning, structure, or repair repeatedly fail
Purpose	Detect the pattern where the user experiences mutual planning, accountability, or relational definition as threat or pressure and protects autonomy by deflection, delay, ambiguity, or selective disengagement.
Primary sources	Assessment items around commitment, planning, definition, structure, exclusivity, accountability; messaging deferral events; calendar or plan cancellations; repair refusal; mismatch between stated goals and tolerated depth.
Default posture	Meaningful but not automatically pathological. The rule must distinguish healthy pacing from defensive anti-structure behavior.
User exposure	Softened user-facing labels allowed, such as “pressure mismatch” or “structure resistance”.

Canonical detection logic

- S1 requires either an assessment posture that strongly resists structure or one behavioral episode where planning or clarity is dodged after previously signaling interest.
- S2 requires two qualifying events in 45 days or one assessment contradiction plus one behavior event that blocks structure or repair.
- S3 requires repeated ambiguity maintenance that materially harms pacing alignment or leaves the other person carrying all definition labor.
- S4 requires sustained patterning across multiple attempts at clarification, planning, or accountability, especially when accompanied by excuse churn or selective availability.
- S5 is reserved for chronic and reinforced anti-structure behavior that drives repeated guarded or suppressed forecast outputs.

Severity ladder for this cluster

Lvl	Trigger Posture	Default Write-Back	Primary Effect
S1	healthy caution or early signal	watch trace only	no score move; pacing watch

S2	pattern starts to repeat	user write-back with low persistence	small readiness/pacing drag
S3	active anti-structure pattern	persistent user write-back and pair pacing caution	moderate readiness and forecast effect
S4	sustained structure resistance	persistent user + pair guarded trace	strong forecast drag; gated product prompts
S5	chronic destabilizer	persistent protection state	suppress overconfident readiness claims

Area	Rule
Write-back	Persist planning/clarity avoidance metrics, recent qualifying event IDs, stated goal mismatch markers, and the trigger version. Pair write-back is allowed from S3 upward when the pattern directly harms pacing or reciprocity.
Decay / clear	Downgrade after 28 quiet days plus at least one completed plan, one direct clarification, or one successful repair conversation. Clear after 75 quiet days with no renewed resistance. Keep active longer when slow drift is also active.
Readiness effect	Control Avoidance mostly affects readiness through behavior-responsibility and readiness overstatement interaction. S3-S5 can limit numeric readiness display if the user claims serious intent while repeatedly avoiding the structure required for it.
Forecast effect	Adds pressure mismatch and pacing drag. At S4-S5 it can materially lower pacing alignment and pair stability even if attraction and compatibility remain high.
Drift interlock	Can feed slow drift when repeated resistance gradually erodes continuity. It should not be recast as drift if the issue is a single planning dispute inside an otherwise stable connection.
AI / Sanctuary use	AI can ask whether structure feels like pressure, whether the user is preserving optionality, and whether the other person is carrying definition work alone. Sanctuary prompts should target directness,

	clean pacing, and non-defensive language.

QA examples to lock before production

Lock examples for healthy slow pacing versus avoidant anti-structure behavior. A cautious user who communicates clearly should not rise above S1-S2. A user who asks for closeness yet repeatedly evades clarity should reach S3 or higher.

CL-103 - Slow Drift

Family	Longitudinal drift cluster
Scope	Pair-state first, with mirrored user traces for downstream coaching and guarded display rules
Purpose	Detect sustained deterioration in connection quality across a governed rollup window without overreacting to one conflict, one pause, or one acute life event.
Primary sources	DriftEligibility rollups, response latency shift, unanswered bid density, ritual drop-off, reciprocity decay, repair failure trend, messaging continuity, shared routine loss, acute-state overlays.
Default posture	Governed by DriftEligibility. Slow Drift is a medium-lane operational state and should never behave like an instant mood spike or an identity rewrite.
User exposure	Internal first. User-facing surfaces should use guarded language such as “recent connection slippage” or “stability softening”.

Canonical detection logic

- Do not open above S1 unless DriftEligibility deems the pair eligible for slow-state evaluation using the approved rollup windows and minimum evidence count.
- S1 may open when rollup metrics soften but remain inside watch range; no public number should change yet.
- S2 requires a confirmed shift in at least two governed drift metrics, such as unanswered bid density plus ritual drop-off, across the active rollup window.
- S3 requires qualified deterioration that persists across the current window and remains directionally weaker than the previous window after acute-state adjustments.
- S4 requires persistent connection erosion with weak repair recovery or continued reciprocity loss despite opportunity to stabilize.
- S5 is reserved for prolonged and reinforced instability that materially threatens pair continuity and keeps forecast in guarded or high-risk posture.

Severity ladder for this cluster

Lvl	Trigger Posture	Default Write-Back	Primary Effect

S1	watch range only	pair watch state; no public alert	confidence watch, no penalty
S2	eligible softening confirmed	pair drift write-back opens	small forecast drag; guarded copy only
S3	active slow drift	persistent pair + mirrored user traces	moderate forecast penalty
S4	sustained erosion	persistent guarded state	strong forecast penalty; pair interventions
S5	terminal instability pressure	persistent high-risk state	suppression and escalation actions

Area	Rule
Write-back	Persist rollup window IDs, metric deltas, eligibility posture, acute-state adjustment markers, confidence band, and decay checkpoints. Mirrored user traces may store only the soft downstream meaning, not raw pair-sensitive details.
Decay / clear	Requires governed improvement across the next qualified rollup window. Downgrade one level only after metric recovery plus repair stabilization. Hard clear needs two stable windows. Acute grief or anger overlays may pause escalation rather than force immediate decay.
Readiness effect	Slow Drift affects readiness modestly and mainly through confidence and pair-state caution. It should not heavily re-score an individual unless paired with user-specific clusters like Repair Gap or Vulnerability Loop.
Forecast effect	This is one of the strongest forecast inputs in batch 1. S3-S5 directly increase drift load, lower pair stability, and can force guarded or suppressed forecast surfaces.
Drift interlock	This cluster is the drift family anchor. It must carry the exact rollup metrics, eligibility state, move entropy support, and state transitions defined by the DriftEligibility appendix.
AI / Sanctuary use	AI may use Slow Drift to time check-ins, suggest repair

	rituals, or explain that the issue looks gradual rather than explosive. AI may not describe it as fate or final outcome.

QA examples to lock before production

Run simulations for: stable pair with one argument; should stay below S2. Pair with steadily rising response latency and ritual drop-off; should reach S3 after eligibility. Pair in acute grief; should use adjustment markers and avoid false S4 escalation.

CL-104 - Push-Pull

Family	Pair reciprocity destabilization cluster
Scope	Pair-state with mirrored user pattern traces where permitted
Purpose	Detect cyclical alternation where one or both users pursue, intensify, then withdraw, causing unstable reciprocity, mixed signals, and emotionally expensive continuity.
Primary sources	Mutual initiation imbalance, pursuit spikes followed by retreat, conversation bursts followed by silence, repeated ambiguity after intensity, bid-response rhythm swings, pair-level cluster overlaps.
Default posture	More pair-centric than Vulnerability Loop. It focuses on the destabilizing cycle rather than the private reason underneath it.
User exposure	User-facing exposure should stay softened, for example “hot-cold rhythm” or “unstable reciprocity pattern”.

Canonical detection logic

- S1 opens when the system sees one early pursue-retreat swing or one reciprocity pattern that hints at cyclic instability.
- S2 requires two swings inside the configured window or one strong swing plus a supporting contradiction cluster such as Vulnerability Loop or Control Avoidance.
- S3 requires a repeatable hot-cold rhythm that visibly lowers conversation continuity or pacing alignment.
- S4 requires chronic cycling that repeatedly destabilizes the pair and prevents stable repair or stable progression.
- S5 is reserved for entrenched cycling that keeps producing rupture risk and forecast suppression despite interventions.

Severity ladder for this cluster

Lvl	Trigger Posture	Default Write-Back	Primary Effect
S1	pattern hint	pair watch only	no direct penalty

S2	cycle confirmed	pair write-back opens	small reciprocity drag
S3	active cycle	persistent pair state; mirrored user traces	moderate pair stability penalty
S4	chronic hot-cold instability	persistent guarded pair state	strong rupture-risk lift
S5	entrenched destabilizer	persistent high-risk pair state	suppression / safety escalation

Area	Rule
Write-back	Store swing count, cadence variance, response asymmetry, overlapping user clusters, and the most recent reinforcing event block. Do not attribute blame in storage; use pattern-state language.
Decay / clear	Requires at least 21 days of steadier reciprocity plus one successful repair or one stable progression marker. Clear after 60 stable days or two qualified stable windows, whichever is stricter.
Readiness effect	Usually indirect. Push-Pull alone should not heavily lower an individual score unless mirrored user traces show that one user is repeatedly generating the cycle.
Forecast effect	Strong forecast relevance. S3-S5 materially lower pair stability and can raise rupture risk. This cluster should often co-travel with Repair Gap when the cycle never stabilizes after conflict.
Drift interlock	Repeated Push-Pull can mature into Slow Drift if the net direction is gradual deterioration rather than volatile repetition. Keep the two clusters distinct in storage even when they co-occur.
AI / Sanctuary use	AI can translate the cycle into pacing and reciprocity coaching, suggest slowing intensity jumps, and prompt users to stop outsourcing clarity to chemistry spikes.

QA examples to lock before production

Two cases must stay separate: flirtation with normal pacing variance should stay low. Repeated intensity-then-silence cycles that harm continuity should reach S3+. Also test mirrored-user fairness

so one-sided blame is not inferred from incomplete data.

CL-105 - Repair Gap

Family	Repair failure / conflict recovery cluster
Scope	Pair-state first, with user accountability traces when one side repeatedly avoids or collapses repair
Purpose	Detect the difference between conflict itself and the inability or unwillingness to recover from it. The cluster is about what happens after rupture, not about having disagreements.
Primary sources	Conflict events, apology follow-through, post-conflict latency, reflection completion, plan-to-repair adherence, repeated unresolved issue reopening, message and ritual recovery metrics.
Default posture	High-value operational cluster because it predicts whether tension is survivable. Conflict without repair is expensive; conflict with repair can still be healthy.
User exposure	User-facing language may say “repair strain”, “slow recovery”, or “unfinished conflict pattern”.

Canonical detection logic

- S1 opens when one conflict or rupture event lacks timely repair but evidence is sparse.
- S2 requires two unresolved or poorly repaired events in the active window, or one significant rupture followed by no qualifying repair attempt.
- S3 requires repeated repair failure, apology without follow-through, or unresolved tension that drags continuity and forecast quality.
- S4 requires chronic recovery collapse, issue looping, or repeated rupture without functional repair participation.
- S5 is reserved for persistent inability to recover that materially threatens connection survival and keeps rupture risk elevated.

Severity ladder for this cluster

Lvl	Trigger Posture	Default Write-Back	Primary Effect
S1	single repair miss	watch trace	no major score move
S2	repair pattern forming	pair write-back opens	small forecast and

			confidence drag
S3	active repair deficit	persistent pair state	moderate forecast penalty
S4	chronic recovery collapse	persistent guarded pair state	strong rupture-risk lift
S5	critical unresolved state	persistent high-risk state	suppression / intervention escalation

Area	Rule
Write-back	Persist conflict-event references, repair latency, apology-follow-through signals, whether repair was attempted by one side or both, and resolution status. User accountability traces may note repeated repair avoidance without exposing pair-private detail in public surfaces.
Decay / clear	Downgrade after one successful repair sequence plus 14 quiet days. Hard clear after two successful repairs or 45 stable days with no reopened unresolved issue. Do not clear solely because the conversation went silent if Slow Drift or Hold Unstable is active.
Readiness effect	Can lower readiness confidence and behavior-responsibility components for users who repeatedly avoid repair. At S4-S5, it can also cap public readiness optimism even if self-reported readiness remains high.
Forecast effect	One of the strongest negative forecast contributors in batch 1. It directly lowers repair capacity, pair stability, and can lift rupture risk quickly if repeated.
Drift interlock	Repair Gap often feeds Slow Drift. Use distinct storage: Repair Gap captures the failure to recover after conflict; Slow Drift captures the longer directional softening that follows.
AI / Sanctuary use	AI can encourage concrete repair behavior: acknowledging the issue, following through, returning after cool-down, and making the next repair attempt specific rather than performative.

QA examples to lock before production

Test healthy conflict with successful repair; it should not stick. Test apology-without-follow-through; should escalate. Test silence after rupture; should not auto-clear. Test one-sided repair effort so the trace records asymmetry without assuming motive.

CL-106 - Readiness Overstatement

Family	Readiness integrity / self-report mismatch cluster
Scope	Individual
Purpose	Detect the gap between claimed relational readiness and the evidence-supported readiness posture. This protects the public readiness surface from becoming a confidence costume.
Primary sources	Declared intent, assessment answers, contradiction clusters, active shadow pressure, drift traces where relevant, repair behavior, pacing behavior, moderation or safety holds.
Default posture	Governance-heavy cluster. It is not punishment for optimism; it exists to stop the system from presenting inflated certainty or maturity that the evidence does not support.
User exposure	Use softened phrases such as “readiness still stabilizing”, “signal confidence not yet matched”, or “current actions trail stated intention”.

Canonical detection logic

- S1 opens when self-reported readiness or serious-intent posture is materially stronger than evidence confidence, but there is not yet enough contradiction to call it a persistent mismatch.
- S2 requires either repeated mismatch across refreshed evidence states or one clear mismatch combined with an active contradiction or repair-related cluster.
- S3 requires durable overstatement: the user continues to present high readiness while behavior, active clusters, or unresolved pressure clearly pull lower.
- S4 requires repeated overstatement across multiple evidence windows or product surfaces, especially if public numeric readiness would be misleading.
- S5 is reserved for severe mismatch combined with protected states where public exposure must be suppressed until evidence stabilizes.

Severity ladder for this cluster

Lvl	Trigger Posture	Default Write-Back	Primary Effect
S1	confidence mismatch watch	internal watch only	copy softening only

S2	mismatch confirmed	user write-back opens	numeric display caution
S3	durable overstatement	persistent user state	public band restriction
S4	repeating inflated posture	persistent restriction state	suppress raw number or upgrade prompts
S5	critical mismatch	protection state	internal only until stabilized

Area	Rule
Write-back	Persist declared readiness posture, last evidence refresh window, delta between self-report and evidence-supported readiness, active supporting clusters, and display restrictions in force. This write-back may directly drive UI suppression logic.
Decay / clear	Downgrade after the evidence-supported readiness converges upward or the user posture converges downward for one full refresh cycle. Clear after two aligned cycles with no active supporting cluster above S2.
Readiness effect	This cluster directly governs the readiness surface. It can downgrade numeric display to band-only, soften public copy, or suppress the number entirely while leaving internal readiness calculations intact.
Forecast effect	Indirect only. Forecast should not use this cluster as a major pair penalty except through supporting clusters that created the mismatch.
Drift interlock	Readiness Overstatement does not itself create drift. It may reference drift as evidence when the user presents stable-high readiness during active unresolved slow drift or hold-unstable conditions.
AI / Sanctuary use	AI may explain the difference between intention and current capacity without shaming the user. Sanctuary prompts should focus on alignment, not accusation.

QA examples to lock before production

Test optimistic but improving users; should stay low and decay as evidence catches up. Test high self-

claim plus unresolved Repair Gap and Control Avoidance; should escalate and suppress raw readiness number. Test low confidence sparse-evidence users; should use caution rather than severe mismatch penalties.

Batch 1 implementation checklist

- Freeze immutable cluster IDs and threshold constants in configuration with version tags.
- Map workbook contradiction pairs and reflection clusters to these cluster IDs before UI work begins.
- Build event capture for repair, reciprocity, planning, and retreat signals before enabling forecast penalties in production.
- Attach softened user-facing copy and suppression rules before any public readiness or forecast exposure.
- Run historical backtests, false-positive review, and paired QA simulations for each cluster before opening guarded-mode surfaces.

Next companion artifacts after this batch

Maintain the threshold workbook references, the user-facing copy library, admin review playbooks, and migration notes beside this pack. The cluster spec is the law, and these batch files are the first set of working statutes.