Event Labeling Heuristics v1.0

This is the “autopilot labeler” that turns raw chat text into structured **trigger_labels** + **user_move** + (optional) **tone flags**. It’s not perfect—its job is to bootstrap data *until* you add human annotation and learned classifiers.

You’ll implement this as a service: Labeler.extract(text, context) -> labels.

1) Inputs and outputs

Inputs

- text (current user message)
- prev_text (previous turn, optional)
- metadata (optional): response latency, message length, punctuation count, caps ratio
- genome_snapshot (optional): helps disambiguate
- state_pre (optional): arousal/safety estimates

Outputs

- trigger_labels[] (multi-label)
- user_move (single label)
- markers[] (supporting evidence tags; internal)
- confidence (0–1)

2) Preprocessing (do this first)

Normalize:

- lowercase
- strip extra whitespace
- keep punctuation counts
- detect all-caps ratio
- detect repeated punctuation (“!!!”, “???”)
- detect profanity count (simple list)
- detect “always/never” count
- detect time pressure phrases (“right now”, “answer me”)

Also extract:

- message length buckets: short (<12 tokens), medium, long (>80 tokens)

3) Trigger labels: rule sets (multi-label)

Each trigger gets a score s (0–1). Add triggers where s >= 0.55.

3.1 Rejection cue (rejection_cue)

High-signal phrases:

- “you don’t care”
- “ignored”
- “left me / leaving”
- “ghost”
- “forgot about me”
- “i’m not important”
- “fine, i get it”
- “whatever, i’ll stop”

Scoring:

- +0.45 if any high-signal phrase appears
- +0.20 if “fine” + short message (<10 words) + cold punctuation (period)

- +0.15 if contains “don’t care / not important”
- +0.10 if mentions “always/never” plus “you”

Clamp to 1.

3.2 Ambiguity (ambiguity)

Signals:

- “what do you mean”
- “i don’t understand”
- “unclear”
- “make it make sense”
- “so…?”
- repeated questions: ?? or multiple question marks
- “which is it”
- “you said X then Y”

Scoring:

- +0.35 for question density (question marks / tokens > 0.03)
- +0.35 for contradiction callout (“you said… then…”)

- +0.25 for “what do you mean / explain”
- +0.10 for “so…” / “ok…” with ??/???

3.3 Disrespect (disrespect_cue)

Signals:

- direct insults: “stupid, dumb, crazy”
- contempt: “pathetic, loser”
- belittling: “grow up”
- dismissive: “shut up”
- sarcasm markers: “sure.” “yeah right”
- name-calling

Scoring:

- +0.60 if insult/name-calling present
- +0.30 if “shut up”/“whatever” + hostile punctuation
- +0.20 if sarcasm marker + “you”

3.4 Control loss / pressure (control_loss)

Signals:

- “don’t tell me what to do”
- “stop trying to control me”

- “you can’t make me”
- “i’m not doing this”
- “leave me alone”
- “back off”

Scoring:

- +0.55 if any direct control-loss phrase
- +0.20 if withdraw move detected + pressure words (“right now”, “answer”)
- +0.15 if imperative push detected (“you need to”, “you must”) in prev_text (if available)

3.5 Unfairness / scorekeeping (unfairness)

Signals:

- “it’s not fair”
- “i always / you never”
- “after everything i’ve done”
- “i did X and you did Y”
- “you owe me”
- “keep that same energy”

- “where was that for me”

Scoring:

- +0.40 if “always/never” count >= 2
- +0.35 if “after everything” / “i did… you did…”
- +0.25 if “owe me” / “fair”
- +0.15 if lists receipts (multiple clauses with “and” + past actions)

3.6 Shame exposure (shame_exposure)

Signals:

- “i feel stupid”
- “i’m embarrassed”
- “i hate that i said that”
- “why am i like this”
- “i’m sorry i’m like this”
- “i feel disgusting / ashamed”
- self-attack: “i’m worthless / i’m a mess”

Scoring:

- +0.55 if explicit shame words

- +0.25 if self-attack present
- +0.20 if apology + self-attack combo

3.7 Abandonment panic (abandonment_panic) (optional but useful)

Signals:

- “please don’t leave”
- “are you leaving me”
- “say you’re not leaving”
- “promise”
- “i can’t lose you”

Scoring:

- +0.60 for “don’t leave / are you leaving”
- +0.25 for “promise” + high question density
- +0.15 for reassurance begging (“please answer”)

3.8 Betrayal / trust rupture (betrayal)

Signals:

- “you lied”
- “you betrayed me”
- “i can’t trust you”
- “you went behind my back”
- “who is she/he”
- secrecy accusations

Scoring:

- +0.60 if lie/trust phrases
- +0.30 if third-party suspicion terms

4) User move classification (single label)

Pick the highest score among these moves. If tie, choose the more conservative/high-risk class.

Moves:

- withdraw
- press_for_clarity
- escalate
- appease

- repair
- joke_deflect

4.1 Escalate

Signals:

- profanity count high
- insults/contempt
- caps ratio high
- repeated exclamation
- “done with you”, “f*ck this”
- threats/ultimatums: “if you don’t X, I’m leaving”

Score:

- +0.45 disrespect trigger score
- +0.25 profanity>0
- +0.20 caps_ratio>0.25
- +0.20 exclamations>=3
- +0.25 ultimatum phrase

4.2 Withdraw

Signals:

- “i’m done”

- “leave me alone”
- “whatever”
- “k.” “fine.”
- short cold messages
- “i’m not talking about this”
- silence marker (if metadata shows long delay + short reply)

Score:

- +0.50 control_loss score if includes “leave me alone/back off”
- +0.35 short message (<8 tokens) + (“fine” or “whatever”)
- +0.25 latency high + short reply
- +0.20 “i’m done / not talking”

4.3 Press for clarity

Signals:

- multiple questions
- “explain”
- “what do you mean”
- “which is it”
- contradictions called out

Score:

- +0.55 ambiguity score
- +0.25 question_density
- +0.20 “answer me” pressure phrasing (without insults)

4.4 Appease

Signals:

- over-apologizing
- “it’s my fault”
- “i’m sorry” repeated
- “i didn’t mean to”
- self-blame + fear of losing them

Score:

- +0.50 shame_exposure score if paired with apology
- +0.35 “sorry” count >=2
- +0.25 “it’s my fault” / self-blame
- +0.20 abandonment_panic present

4.5 Repair

Signals:

- ownership + solution focus
- “can we restart”
- “i want to fix this”
- “help me understand”
- “i’m here”
- “i get why you felt that”

Score:

- +0.45 phrases: “fix this / restart / talk it through”
- +0.30 validation phrases (“i get why…”, “that makes sense”)
- +0.20 low punctuation aggression + medium length
- −0.30 if disrespect/profanity present

4.6 Joke/deflect

Signals:

- “lol”, “lmao”
- memes, sarcasm
- changing topic
- “anyway…”

Score:

- +0.45 humor tokens
- +0.25 topic shift markers
- +0.20 sarcasm without direct insult

5) Confidence scoring (0–1)

Confidence should reflect signal strength and consistency.
Start at 0.35 and add:

- +0.20 if one move score ≥ 0.75
- +0.15 if top move beats second by ≥ 0.20
- +0.10 if triggers align with move (e.g., ambiguity + press)
- −0.15 if message is extremely short (<4 tokens)
- −0.10 if only punctuation/emoji

Clamp 0–1.
If confidence < 0.45, set user_move="unknown" and only keep the strongest 1–2 triggers.

6) Context overrides

(optional but powerful)

Use genome_snapshot to resolve ambiguity:

- High pursuit_withdraw_bias (pursuer): “fine.” might actually be *pressing* later—still label as withdraw but increase reassurance_hunger target.
- High autonomy_need: “why are you asking so much” likely withdraw + control_loss
- High clarity_appetite: question-heavy messages are “press_for_clarity” not “escalate” unless disrespect exists.

These overrides should adjust confidence, not rewrite labels unless strong.

7) Output format

{ "trigger_labels": ["ambiguity", "unfairness"], "user_move": "press_for_clarity", "markers": ["question_density_high",

"always_never_count_2",
"contradiction_callout"], "confidence": 0.71
}

8) Implementation order (so your team doesn't drown)

- Build trigger scoring (multi-label)
- Build move classifier (single label)
- Add confidence & unknown fallback
- Add genome/context overrides
- Log markers for audit/debug
- Replace heuristic labeler gradually with ML classifier trained on your own annotated data

One more "unfair advantage" add-on

Add a **Rupture Threat Detector** (binary) for phrases like:

- "i'm done"
- "i can't do this"

- “we’re over”
- “don’t talk to me”
- “block me” This becomes a high-priority label feeding the Governor and Forecast Engine.

If you want, the next step is the **Taxonomy Lock File**: a single authoritative list of all triggers, moves, strategies, and contradiction clusters (IDs + descriptions) so every service uses the same enums and your data never gets corrupted by “almost the same label.”