Genesis OS Anger Hex Labeling, Trigger, and Event Creation Appendix v2

Canonical source map, trigger law, subtype assignment, event schema, and partner-safe display rules

Companion to the Anger Hex official contract, aligned to the corrected master blueprint and scoring workbook.

Purpose. This appendix hard-locks the precise places anger evidence can originate, the concrete trigger families the engine may score, the exclusions that prevent false positives, the minimum evidence law required before the system is allowed to label an event as anger, the six-band Anger Hex subtype routing logic, and the write-back / display rules that keep the anger layer useful without letting it hijack baseline identity or compatibility truth.

1. Governing stance

Anger Hex is a companion runtime. It does not diagnose, moralize, or rewrite a user's baseline archetype, shadow type, or workbook-root scoring outputs. It tracks anger expression, pressure, recovery, and repair in governed time windows so downstream systems can understand relational safety and regulation quality without turning every hard moment into a permanent identity sentence.

- The corrected master blueprint places Anger Hex in the companion-system lane beside DriftEligibility, Shadow Activation, and E_GRIEF.
- The workbook remains the canonical source of baseline scored-answer truth; Anger Hex enriches runtime understanding through governed event packets and decaying overlays.
- A single emotional signal is not enough to permanently label a user as angry. The engine must score evidence, apply exclusions, meet threshold law, and only then create an anger event.

2. Canonical Anger Hex bands

This appendix uses the six-band Anger Hex ladder already implied by prior Anger Hex / Anger Map materials and re-expresses it in build language. The bands are event outcomes, not personality types.

Band	Label	Hex	Core meaning	Typical posture	Default persistence
R1	Disruption	#FFB36A	Small irritation or flow break	Short replies, micro-sighs, sharper tempo	minutes
R2	Violation	#FF894D	Boundary or standard crossed	Jaw set, scanning, corrective tone	minutes to hours
R3	Threat	#FF5A36	Pressure or identity threat felt	Compressed speech, grievance	hours

				stacking, tense silence	
R4	Defense	#E43117	Boundary violation feels personal or repeated	Sharp clarity, armor mode, demand for correction	hours to a day
R5	Fire	#B10F0F	Pressure discharge or outward rupture	Outburst, hard escalation, harsh release	episode based
R6	Overload	#5A0B0B	Anger implodes into shutdown or disappearance	Hostile withdrawal, deadened tone, ghosting	episode to multi-day

3. Canonical runtime flow

Every anger-labeled runtime passes through the same ordered pipeline. No service is allowed to skip steps 2 through 5 and jump straight to event creation.

Step	Runtime action	Required output
1	Emotion candidate detected	candidate_id, source_id, raw signal, timestamp
2	Candidate mapped to anger trigger family	trigger_family_id, initial source role (direct/support/context)
3	Evidence scored by source and strength	evidence points, confidence contribution, source provenance
4	Exclusions checked	exclusion pass/fail, downgraded emotion class if not anger
5	Threshold crossed or not	create_event flag, confidence band, insufficiency reason if no event
6	Anger event created	anger_event_id, severity_band, subtype provisional
7	Event classified into Anger Hex subtype	R1-R6 final band, output style, repair risk
8	Write-back rules applied	session writeback, overlay writeback, partner-safe summary if allowed
9	Partner-safe display filtered	display label, privacy scope, suppression flags

4. Source registry: where anger evidence is allowed to come from

The engine may only source anger evidence from the following governed lanes. Each lane must declare whether it is direct evidence, supporting evidence, or contextual evidence, and whether it may create an anger event alone or only strengthen one.

Source lane	Exact fields / objects	Evidence role	When captured	Can create event alone?	Why it matters
Onboarding answers	intake answers, dealbreakers, history prompts, regulation prompts	supporting; direct only when a user explicitly self-reports repeating anger behaviors	initial profile build and onboarding refresh	usually no	historical / declared context, not primary runtime
Scored assessment answers	AnswerBank selections, shadow items, readiness items, conflict items	supporting; direct when answer text itself describes recurring anger behavior	assessment completion, rescoring windows	yes for declared recurrent anger traits with confidence cap	baseline pattern propensity
Polarity prompts	Polarity_QuestionMap and Polarity_AnswerBank selections	supporting/ context	polarity assessment completion and pair recalculation	no	interaction style pressure patterns, pursuit/withdraw heat
Messaging behavior	message text, delay patterns, punctuation load, conversation events, repair bids, block / leave / mute events	direct + supporting	live and batch message processing	yes	strongest runtime lane
Sanctuary sessions	Mirror, Whisper, Stillness check-ins, Micro Rituals, Alignment Compass, Last Light, DM Reflection	supporting; direct when user explicitly describes anger episode or the system captures anger markers during	session completion or significant step transitions	yes with confidence cap unless supported by another lane	regulated self-report and post-event recovery context

		reflective processing			
ExsAnonymous sessions	ex story capture, residue pass, closure reflections, pattern extraction	supporting; direct if recurring anger narrative is explicit and repeated	session save / completion	yes with confidence cap	historic anger residue and unresolved rupture style
Game behavior	session quits, trash talk labels, surrender patterns, teammate blame, rage exits, cooldown refusal	direct + supporting	game event stream and post-session reflection	yes	pressure-behavior lane
Reflection clusters	cluster hits such as vulnerability loop, control avoidance, push-pull, repair gap, readiness overstatement	context + supporting	cluster writeback generation	no	explains structure beneath anger expression
Manual/admin review	moderation labels, safety reviews, abuse review notes	supporting or override-only depending role	post-review only	no autonomous creation; can confirm or suppress existing events	safety and dispute handling

5. Trigger family catalog: what counts as anger evidence

The engine must use concrete trigger families rather than vague impressions. Below is the canonical catalog for v2. Each trigger family can be detected from more than one lane, but each lane must use the lane-specific extraction notes later in this appendix.

ID	Trigger family	What it means	Primary lanes	Notes
AH_TF_01	Direct anger language	Explicit words or phrases indicating anger, rage, being	Messaging, Sanctuary, ExsAnonymous,	Direct first-person ownership counts stronger than

		mad, pissed, heated, furious, done, fed up, or equivalent	game chat, admin review	general venting
AH_TF_02	Blame-heavy phrasing	Statements centering fault assignment, repeated accusation, you-always / you-never attacks, responsibility dumping	Messaging, ExsAnonymous, Sanctuary, assessments	One sentence can be signal; pattern requires repetition or stronger companions
AH_TF_03	Contempt markers	Belittling tone, disgust framing, ridicule, eye-roll language, derision, humiliation framing	Messaging, game chat, admin review	High-risk trigger family; often escalates severity faster
AH_TF_04	Control-through-intensity	Pressure used to force pace, compliance, emotional surrender, or immediate correction	Messaging, polarity context, game behavior	Often reads as urgency plus dominance rather than simple irritation
AH_TF_05	Protest escalation	Escalation after felt neglect, rejection, or unanswered bids; repeated pushing when the bond feels threatened	Messaging, drift lane, Sanctuary, clusters	Must be distinguished from anxiety without hostility
AH_TF_06	Repeated hostile withdrawal	Silence, disappearance, short deadening replies, or refusal to engage after rupture, with hostility markers attached	Messaging, drift, Sanctuary, games	Not every withdrawal counts; hostility component required
AH_TF_07	Retaliatory behavior	Behavior intended to punish, sting, get even,	Messaging, games, admin review	Often shifts event toward R4-R6 faster

		embarrass, withhold, block, or provoke payback		
AH_TF_08	Scorekeeping / grievance stacking	Tallies of prior wrongs, itemized resentment, repeated ledger-style recall used as present-pressure	Messaging, ExsAnonymous, Sanctuary, assessments	Can begin as frustration and become anger after persistence
AH_TF_09	Rupture without repair	Rupture event occurs and the user refuses repair, rejects bids, or escalates instead of stabilizing	Messaging, Sanctuary DM reflection, pair events	Needs pair context and sequence timing
AH_TF_10	Aggression-adjacent action patterns	Threat-adjacent gestures, slam, spam, crowding, hard leave/rejoin loops, rage quit, intentionally destabilizing moves	Games, messaging events, moderation	Highest caution family; partner-safe displays must be softened

6. Source-by-source extraction law

6.1 Onboarding answers

Onboarding is not treated as real-time anger behavior; it is a declaration lane. It seeds propensity, history, and watch-list context. It can only create an anger event by itself when the user plainly self-reports a repeated anger pattern in present tense or near-present framing.

- Read fields: conflict history prompts, regulation prompts, prior relationship rupture prompts, self-description prompts, open-text red-flag fields, dealbreaker explanations, and any optional narrative responses where present.
- Capture time: first onboarding completion, major onboarding edits, and later governed profile refreshes.
- Direct anger signals: 'I get angry fast', 'when I feel disrespected I snap', 'I shut down hard when mad', 'I hold grudges', 'I can be cutting when pushed'.
- Supporting-only signals: references to strict standards, conflict avoidance, being intense, or disliking disrespect without explicit anger framing.
- Do not label anger from one vague word like 'passionate', 'strong personality', or 'protective'. Those stay as context unless combined with stronger evidence elsewhere.

Onboarding pattern	Maps to trigger families	Default role	Can create event?	Comments
Present-tense self-report of snapping / exploding / raging	AH_TF_01, AH_TF_04 or AH_TF_10	direct	yes, capped at moderate confidence	Create declared-pattern event, not live episode
Historic statement about past anger only	AH_TF_01 or AH_TF_08	supporting	no	Seeds watch-list and decays unless later confirmed
Boundary language without hostility	none or context only	context	no	May support R1-R2 if runtime event occurs later
Grudge / tally / scorekeeping statement	AH_TF_08	supporting or direct if repeated / current	conditional	Needs present or recurring framing to rise

6.2 Scored assessment answers

Assessment answers are stronger than onboarding because they live inside the workbook's scored-answer universe. Still, they remain declared evidence unless paired with runtime behavior. The engine should use the workbook answer record, dimension mapping, and answer text together.

- Read fields from AnswerBank: answer_id, question_number_int, answer_text, archetype, shadow_type, dimension_primary, dimension_secondary, readiness_norm, attachment_risk_norm, integrity_penalty_norm, cluster_penalty_norm, and pattern_risk_norm.
- Capture time: full assessment submission, monthly/quarterly retake, and any governed rescoring path.
- Direct anger signals come from answers that explicitly normalize snapping, scorekeeping, punishing withdrawal, retaliatory logic, pressure-based correction, contempt, or refusal to repair.
- Supporting signals come from adjacent regulation or attachment items that imply pressure sensitivity without naming anger.

Assessment evidence class	Examples of answer content that count	Likely trigger families	Weight stance
Declared sharpness	'When I feel disrespected my tone gets cold fast'	AH_TF_01, AH_TF_04	medium direct
Punitive withholding	'Sometimes people need to sit in what they caused'	AH_TF_06, AH_TF_07	high direct
Ledger mentality	'I remember patterns and people should answer for them'	AH_TF_02, AH_TF_08	medium direct
Refusal of repair	'If they cross the line enough times there is	AH_TF_09, AH_TF_06	medium direct

	nothing to talk about'		
Pressure sensitivity only	'I get overloaded when things feel chaotic'	context only unless paired	supporting

6.3 Polarity prompts

Polarity prompts do not exist to diagnose anger. They exist to explain pursuit, surrender, validation, structure, intensity, and pressure. For anger use, polarity is mainly a context lane. It answers whether the nervous system tends to convert tension into pursuit, control, reactive heat, deadening, or rigidity.

- Read fields from polarity workbook tabs: question_id, answer_slot, redesigned_role, question_weight, mapped axes, and any role/axis metadata tied to repair_endurance, flexibility_axis, tension_attraction, validation_dependency, initiation_axis, and structure_preference_axis.
- Capture time: polarity completion, pair recomputation, and any governed recalculation event.
- Polarity alone should not create anger events. It supplies preconditions and interaction style context that shape subtype routing and threshold calibration.
- Examples: high tension attraction plus low repair endurance may lower the barrier for AH_TF_05 protest escalation; high structure preference plus low flexibility can intensify AH_TF_04 control-through-intensity in pair conflict.

Polarity pattern	Possible anger relevance	Engine use
High initiation + high structure + low flexibility	Correction may arrive as pressure or forceful pace-setting	context multiplier for AH_TF_04
High validation dependency + low calm receptivity	Perceived neglect may escalate into protest loops	context multiplier for AH_TF_05
Low repair endurance + high tension attraction	Conflict may intensify before recovery attempts appear	context multiplier for AH_TF_09
Low surrender safety + high structure preference	Loss of control may convert quickly into defense posture	subtype routing aid R2-R4

6.4 Messaging behavior

Messaging is the strongest runtime lane because it provides sequence, timing, text, pacing, repair bids, and actual pair behavior. Messaging can create anger events by itself when the evidence is strong enough.

- Read objects: message text, send time, reply gap, deleted message markers, block / mute / leave events, reaction events, escalation tags from event labeler, repair attempt events, ignored bid counts, and moderation labels where present.
- Capture time: continuously in live stream and in batch repair / quality passes.
- The message processor should score both content and sequence. A harsh statement during a repair attempt carries more weight than the same statement in isolation.
- The engine must distinguish one sharp sentence from a patterned hostile exchange.

Messaging signal	Trigger family	Minimum event note

Explicit 'I am angry / pissed / done'	AH_TF_01	single strong direct candidate
Three or more accusation-heavy lines in one rupture window	AH_TF_02	pattern direct candidate
Belittling or disgust-laden wording	AH_TF_03	high-risk candidate even if brief
Repeated demands for immediate compliance + pressure punctuation	AH_TF_04	sequence-sensitive candidate
Escalating after unanswered bids across a narrow window	AH_TF_05	needs neglect/protest sequence
Disappearing after rupture and returning with hostile deadness	AH_TF_06	needs withdrawal + hostility
Punitive leave/block/unblock/spam cycle	AH_TF_07 or AH_TF_10	high-severity candidate
Long tally list resurrected during current conflict	AH_TF_08	moderate or high candidate
Refusal to answer repair bids across governed window after rupture	AH_TF_09	pair-sequence candidate

6.5 Sanctuary sessions

Sanctuary is a reflective lane, not a trap. Its job is to help interpret and recover, not to over-punish honesty. Sanctuary can create anger events only when the user directly describes a current or recurring anger episode, or when a session logs a structured anger marker with enough specificity.

- Read module events from Mirror, Whisper, Stillness check-ins, Micro Rituals, Alignment Compass, Truth Signal responses, Last Light, Sanctuary-DM Reflection, and Archive writebacks.
- Capture time: on step completion, session completion, and any structured mood / narrative save.
- Weight honest self-report positively for repair potential even while scoring the anger signal itself. Confession of anger is not equivalent to uncontrolled expression.
- Sanctuary can also downgrade candidate anger to grief, hurt, overwhelm, or fear when the reflective content makes the underlying primary emotion clearer.

Sanctuary pattern	Result
'I was angry and I snapped at them'	direct candidate + repair context
'I felt cornered and got sharp, then calmed down'	anger candidate with likely lower band and faster decay
'I was hurt, not angry' with no hostile markers	route away from anger
Ritual completion after rupture with sincere repair plan	apply recovery modifier; keep event but improve decay / repair posture

6.6 ExsAnonymous sessions

ExsAnonymous sits inside Sanctuary governance but deserves its own anger extraction note because it captures residue, old resentment, betrayal framing, and recurring ex-pattern scripts.

- Read objects: ex story capture, rupture narrative, unresolved resentment list, closure draft, residue pass, pattern extraction output, and optional AI reflections where allowed.
- Capture time: draft save, final submit, pattern extraction complete, and later revisit sessions.
- A single bitter story about an ex does not prove the user is presently in anger mode. The lane is strongest when repeated narratives show recurring blame, contempt, punishment logic, or unresolved revenge framing.

ExsAnonymous evidence	Trigger family	Rule
Repeated 'they ruined / wasted / destroyed' narratives with present heat	AH_TF_02	supporting or direct if recurring and current
Humiliation / disgust language toward exes	AH_TF_03	strong supporting
Payback fantasies or punitive withholding stories	AH_TF_07	strong supporting
Long historical ledger of wrongs without present hostility	AH_TF_08	supporting only
Grief-heavy story with sadness more than heat	route to E_GRIEF, not anger	exclusion candidate

6.7 Game behavior

Game behavior matters because pressure strips performance. It shows what happens when pace, loss, status, and control collide. This lane is especially useful for real-time anger output style.

- Read objects: rage quit events, surrender / leave timing, teammate blame labels, toxic chat tags, spam / ping bursts, refusal of cooldown, intentional throw markers if such labeling exists, and post-session reflection prompts.
- Capture time: in-session event stream and post-session recap.
- Games may create anger events when the pattern is strong, but partner-safe display must stay summarized and behavior-framed rather than embarrassing.

Game signal	Maps to	Band tendency
Fast blame after mistake or loss	AH_TF_02 / AH_TF_03	R2-R4
Rage quit / repeated exit loop	AH_TF_10 / AH_TF_06	R4-R6
Punitive sabotage or throw behavior	AH_TF_07 / AH_TF_10	R4-R6
Brief irritation without escalation	AH_TF_01	R1-R2
Cooldown refusal plus immediate re-engage with heat	AH_TF_04 / AH_TF_05	R3-R5

6.8 Reflection clusters and manual review

Clusters are not anger by themselves. They explain the architecture underneath it. Manual review is not a parallel engine; it confirms, suppresses, or annotates governed outputs.

- Clusters that often interact with anger: Vulnerability Loop, Control Avoidance, Push-Pull, Repair Gap, Slow Drift, Readiness Overstatement.
- Use clusters as context multipliers, subtype-routing aids, and threshold stabilizers. Do not let a cluster create an anger event on its own.
- Admin or moderation review may confirm risk, add provenance, suppress a false positive, or create a safety note, but should not invent anger labels without an underlying candidate.

7. Exclusion rules: when not to call it anger

False-positive protection is mandatory. The engine must stop anger labeling whenever the evidence is better explained by a non-anger family or by insufficient hostility.

Exclusion family	What it looks like	Engine action
Pure sadness / grief	Tears, loss, emptiness, yearning, mourning, heaviness without correctional or hostile posture	route to E_GRIEF or sadness family
Overwhelm without hostility	Flooded, overstimulated, needs space, sensory stress, shutdown without blame or punitive edge	route to overwhelm / depletion
Fear without correctional pressure	Anxious, scared, worried, unsure, freezing, hypervigilant but not attacking	route to fear / attachment alarm
Clean boundary without heat	Firm no, self-protection, leaving respectfully, setting pace limits, declining contact without sting	do not mark anger
One blunt line without sequence	Single short or curt reply without repetition, contempt, threat, or retaliatory pattern	keep as weak tension candidate only
Reflective ownership of anger after resolution	User calmly reports prior anger while focusing on repair and accountability	retain context but lower current-event likelihood

8. Minimum evidence law and threshold model

The system must not label an anger event unless minimum evidence law is met. Minimum evidence law uses a simple but deterministic structure so the implementation is stable even if future tuning changes point values.

Threshold posture. Use source-weighted evidence units rather than fragile one-off heuristics. The exact numeric values may be tuned later, but the structure below is mandatory.

Evidence class	Definition	Default strength
Weak direct	Mild irritation phrase, one sharp line, low-grade complaint with no sequence	1 unit
Strong direct	Explicit anger wording, contempt marker, retaliatory action, high-pressure correction, rupture with hostility	2 units
Supporting	Historical or declared pattern, contextual alignment, cluster support, polarity pressure pattern	0.5-1 unit
Recovery modifier	Immediate self-awareness, repair attempt, cooldown completion, ownership without excuse	-0.5 to -1 unit after event creation, not before
Exclusion hit	Evidence that better explains the state as grief, fear, overwhelm, or simple boundarying	blocks or downgrades event

Event type	Minimum evidence threshold	Required mix	Allowed source combinations	Confidence floor
R1-R2 irritation / violation event	2 total units	at least 1 direct signal	messaging alone; Sanctuary + assessment; game alone if explicit	low to moderate
R3-R4 threat / defense event	3 total units	at least 1 strong direct + 1 support or second direct	messaging + sequence; messaging + cluster; Exs + assessment + Sanctuary	moderate
R5-R6 fire / overload event	4 total units	at least 2 strong direct or 1 strong direct + 2 support with action marker	messaging + event actions; games + moderation; messaging + Sanctuary after rupture	moderate to high

Pattern / band writeback	3 created events in governed window or 2 high-severity events	multiple event history required	multi-session / multi-window only	moderate to high

Allowed source combinations must include at least one behavior or explicit self-report lane. Polarity alone, cluster alone, or admin review alone may never create an anger event.

- Valid high-confidence combination example: explicit contempt line in messaging + ignored repair bids + punitive leave event.
- Valid moderate-confidence combination example: Sanctuary admission of snapping + matching assessment answer about getting cutting when disrespected + recent hostile messaging.
- Invalid combination example: high tension-attraction polarity + repair gap cluster + strict standards in onboarding. This can create watch-list context, not an anger event.

9. Confidence bands

Band	Meaning	When used
Low	candidate meets minimum threshold but source diversity or sequence quality is limited	display only internally unless repeated
Moderate	solid direct signal plus supporting context or repeated direct signals	eligible for governed writeback and softened user view
High	multiple direct signals, clear action pattern, strong provenance, exclusions clean	eligible for stronger safety / forecast overlays
Locked	manual review or multi-window recurrence confirms pattern beyond ambiguity	used rarely for operational confidence, not public-facing bravado

10. Event creation: when tension becomes anger

A tension or frustration candidate becomes an anger event only after all of the following are true:

- The candidate maps to at least one canonical anger trigger family.
- Exclusion checks do not route it more cleanly to grief, fear, overwhelm, or ordinary boundarying.
- Minimum evidence threshold is crossed for the proposed severity band.
- At least one allowed source combination includes direct behavior or explicit self-report.
- Confidence band is low or higher and provenance is recorded.
- A subtype router can assign the event to one of the six Anger Hex bands without contradiction.

If a candidate has heat but not enough hostility or correctional pressure, keep it as tension or frustration context. Do not force anger just because the moment is emotionally negative.

11. Subtype assignment: how events route into the six-band Hex

Band	Best-fit markers	Most common trigger families	Routing notes
R1 Disruption	annoyance, shortness, micro-corrections, mild impatience	AH_TF_01 weak direct, minor AH_TF_02	brief; light recovery
R2 Violation	boundary-cross heat, standards breached, warmer body language, corrective focus	AH_TF_02, AH_TF_04 light, AH_TF_05 early	clear but still recoverable
R3 Threat	pressure felt personally, grievance stacking, compressed tone, tense silence	AH_TF_02 repeated, AH_TF_05, AH_TF_08	identity / order under pressure
R4 Defense	armor mode, sharp directness, demands for correction, pointed force	AH_TF_04, AH_TF_09, AH_TF_03 moderated	personalized protection
R5 Fire	verbal or behavioral discharge, outburst, harsh release, impulsive rupture	AH_TF_03 severe, AH_TF_07, AH_TF_10	episode severity
R6 Overload	imploded anger, hostile disappearance, cold shutdown, punitive deadening	AH_TF_06, AH_TF_09 severe, AH_TF_07	darkened post-rupture state

12. Pattern formation: when multiple anger events become a tracked pattern

One event does not equal a stable anger pattern. Pattern formation is a historical writeback rule.

Pattern type	Create when	Writeback target	Meaning
Irritation tendency	3 or more R1-R2 events in 30 days	light overlay	user gets annoyed quickly but may recover fast
Defensive pressure pattern	2 or more R3-R4 events plus weak repair in 45 days	medium overlay	anger expresses as correction or armor
Rupture fire pattern	2 or more R5 events in 60 days	heavy overlay and review flag	outward discharge risk needs governance
Hostile withdrawal pattern	2 or more R6 events in 60 days or repeated R4-	medium/heavy overlay	anger converts into disappearance or

	>R6 conversion		deadening
Mixed anger band	events spread across 3 or more bands with no stable center	analytic overlay	instability matters more than any single band

13. Write-back law

Writeback must be scoped so Anger Hex informs the system without taking over the score engine.

Writeback class	What is stored	Where it can write	What it must not do
Session writeback	anger_event_id, source mix, trigger families, band, confidence, decay timer	session log, event ledger, Sanctuary session if relevant	must not rewrite archetype / shadow baseline
Overlay writeback	current anger_state summary, recent severity, repair latency, recurrence count	companion overlay store, pair runtime context	must not directly override compatibility score
Pattern writeback	band history, recurrence bucket, partner-safe risk label	forecast support lane, admin review lane	must not create permanent public stigma label
Review writeback	manual suppression / confirmation reason	moderation and audit logs	must not silently alter source evidence history

14. Decay rules

Anger must decay. If the system can detect escalation but never recovery, it is not intelligence; it is paranoia with storage.

Band	Default decay posture	Decay accelerators	Decay blockers
R1	fast	cooldown complete, no repeat within 24h, repair attempt logged	repeat irritation chain same day
R2	fast to medium	acknowledgment + softened pace + no new hostile signals	continued blame or correction pressure
R3	medium	repair bid accepted, Sanctuary regulation completion, gap normalization	grievance stacking, repeated protest
R4	medium to slow	clear ownership, repair action, no retaliatory behavior	new rupture, contempt, punitive withdrawal

R5	slow	sustained repair and quiet period, no repeat events	secondary outburst, moderation events
R6	slowest	re-entry with ownership + repair + stability window	ghosting, cold reappearance, repeated deadening

15. Partner-safe display logic

Partner-safe display must soften and summarize. The system should never embarrass a user by surfacing private raw text, diagnostic phrasing, or a melodramatic scarlet-letter label.

Internal truth	Partner-safe display example	Use when
AH_TF_03 contempt + R4 defense	'Conflict can get sharp when this person feels disrespected. Best results come with calm pacing and direct repair.'	moderate confidence, pair context allowed
R6 hostile withdrawal pattern	'This person may go quiet or distant when conflict feels overloaded. Gentle clarity and repair structure help.'	pair-safe summary
R5 fire event with low confidence	no partner-facing display; keep internal only	insufficient repetition or privacy sensitivity
Historical ExsAnonymous resentment only	no partner-facing display	historic lane should stay private unless strongly current
Irritation tendency pattern	'This person responds best to clear expectations and low-chaos communication.'	soft trait guidance only

- Do not expose raw trigger-family IDs to partners.
- Do not reveal exact source lane if it would violate privacy expectations, especially Sanctuary or ExsAnonymous.
- Do not state 'this person has anger problems.' The display should focus on conflict handling and repair conditions.
- Any partner-facing surface must use softened, actionable language and must respect suppression rules from moderation / privacy governance.

16. Exact event labels and dependencies

To keep the implementation deterministic, the event labeler should emit a compact anger event catalog. Suggested production labels are below.

Event label	Meaning	Typical source
anger_candidate_detected	raw signal with possible anger	any lane

	relevance created	
anger_candidate_supported	supporting context attached	assessment, clusters, polarity, onboarding
anger_candidate_excluded	candidate routed away from anger	Sanctuary, moderation, classifier
anger_event_created	threshold crossed; event exists	messaging, Sanctuary, games, ExsAnonymous
anger_event_rebanded	event severity changed after more evidence	event processor
anger_event_recovered	repair / cooldown reduced active state	Sanctuary, messaging, games
anger_pattern_writeback	historical recurrence threshold crossed	batch / analytics pass
anger_display_suppressed	privacy or confidence rule blocks partner-facing view	display layer / moderation

17. Worked examples

Scenario	Correct result
User says in Sanctuary: 'I was angry and got sharp, then I apologized and fixed it that night.'	Create low/moderate anger event, likely R2-R3, apply recovery modifier; do not form pattern
Three accusation-heavy texts, one contempt line, ignored repair bid, then block/unblock cycle	Create strong anger event, likely R4-R5; consider retaliatory family and possible pattern if repeated
User writes in ExsAnonymous that an ex hurt them and they still feel sad	Route to grief / hurt unless strong current anger markers exist
Game session shows rage quit and blame spam after loss, followed by no cooldown and a second rage quit	Create anger event, likely R4-R6 depending hostility and shutdown posture
Polarity indicates low repair endurance and high structure preference, but no runtime hostility exists	No anger event; retain context only
User sets a firm boundary respectfully and stops engaging	No anger event; this is clean boundarying, not hostility

18. Implementation notes for developers

- Store provenance on every evidence packet: source lane, object ID, timestamp, and classifier version.
- Keep band assignment revisable during the same governed episode window; do not create duplicate events when one event merely escalates from R2 to R4.

- Use episode windows so clustered signals from one conflict become one event with accumulated evidence, not six separate panic records.
- Apply privacy hierarchy: raw source text stays in source system; Anger Hex stores normalized summaries and references.
- Never let partner-safe display pull directly from private reflective text.
- Keep batch recalculation idempotent. Reprocessing should not double-count old evidence.
- Manual review should confirm or suppress, not silently fabricate.

19. Final calibration note

This appendix is intentionally stricter than a conceptual anger overview. Its job is to make two engineers build the same anger engine, not two different interpretations of emotional weather. It is calibrated to the existing Genesis OS system: workbook-root scores remain primary, Anger Hex stays a governed companion layer, Sanctuary and ExsAnonymous are treated as reflective-but-real signal lanes, and partner-facing surfaces stay softened, actionable, and privacy-safe.