Anger Hex Official Companion Contract v1

Production event catalog, decay rules, write-back law, and partner-safe display logic

Purpose. This contract promotes Anger Hex from a light companion concept into a governed production subsystem calibrated to the current Genesis OS build set. It defines exactly how anger is captured, labeled, stored, decayed, written back, summarized for partners, and consumed by readiness, forecast, Sanctuary, Human Emotion Genome, clusters, drift, shadow activation, and AI guidance.

Document stance

- Workbook truth remains authoritative for score-engine math. Anger Hex does not freehand new score law; it contributes governed runtime evidence under influence caps.
- Anger is treated as a controlled runtime signal family, not as a diagnosis, personality verdict, or moral label.
- Partner-facing anger surfaces must be translated into safe relational language. Raw anger telemetry is not directly exposed to another user.
- This specification is meant to sit inside the official build set beside Drift, Shadow Activation, Sanctuary, and Human Emotion Genome companion contracts.

1. Role of Anger Hex inside the system

Anger Hex exists because relationship risk is not only about what a person says in an assessment. It is also about how pressure enters the body, how quickly boundary threat escalates, whether repair appears after activation, whether shutdown follows overload, and whether a person repeatedly routes distress into fire, ice, armor, sarcasm, or solution mode. The platform needs a deterministic way to capture that runtime behavior without letting a single bad day rewrite a person's blueprint.

- Primary jobs of Anger Hex: detect pressure and threat states; classify current band; track output style; measure recovery and repair latency; create safe write-backs; feed companion systems under governance.
- Upstream sources: onboarding intake, scored assessment answers, polarity answers, messaging telemetry, pair events, Sanctuary sessions, ExsAnonymous reflections, game telemetry where validated, and manual structured check-ins.

- Downstream consumers: readiness family, pair forecast family, cluster engine, drift system, shadow activation runtime, Human Emotion Genome features, Sanctuary prompts, AI reflection layer, and partner-safe summaries.

2. Canonical Anger Hex taxonomy

The production taxonomy uses six primary bands. Each band has a canonical code, a named phase, a safe interpretation lane, and a standard action posture. The system stores both the numeric band and the normalized label. The normalized label is the code-table value; any product copy shown to users is derived from that code through a copy library rather than hard-coded in the event stream.

Code	Band name	Operational meaning	Typical internal sentence	Default risk posture	Default partner-safe stance
R1	#FFB36A / Disruption	Small flow break or annoyance signal. Pressure is real but still light.	Why now?	Low	Minor tension, still recoverable
R2	#FF894D / Violation	Boundary, standard, or expectation feels crossed.	This should not be happening.	Low-moderate	Friction is rising; clarity matters
R3	#FF5A36 / Threat	Order, identity, fairness, or safety feels compromised.	I should not have to deal with this.	Moderate	Stress is active; repair matters soon
R4	#E43117 / Defense	Anger becomes protective posture and tone hardens.	Fix this now.	Moderate-high	Relational temperature is high; caution
R5	#B10F0F / Fire	Overload discharges outward through sharp expression or rupture behavior.	I'm done.	High	Rupture risk; cooling and repair required
R6	#5A0B0B / Collapse	Anger implodes into shutdown, disappear, or emotional freeze.	I don't care anymore.	High	Withdrawal / disconnection risk; re-entry needed

2.1 Secondary dimensions captured beside the band

- Trigger family: fairness, disrespect, overwhelm, identity threat, uncertainty, betrayal, repeated nonresponse, control loss, boundary breach, environmental disruption.
- Output style: fire, ice, armor, redirect, leadership, repair.

- Context lane: solo, pair-chat, live call, Sanctuary, ExsAnonymous, games, post-conflict reflection, archived retrospective.
- Repair posture: none, attempted, partial, completed, refused, blocked by partner response, blocked by platform safety.
- Recovery arc: fast, normal, slow, stalled, looping, collapsed.
- Confidence lane: inferred-light, inferred-moderate, confirmed-structured, multi-source-confirmed.

3. Core data objects

Object	Purpose	Cardinality	Notes
anger_event	Atomic log of one anger signal, transition, repair action, or decay update.	Many per user / pair	Append-only event store with event labels.
anger_episode	Session-level object that groups related anger events into one bounded episode.	Many per user / pair	Used for latency, peak band, and repair calculations.
anger_profile_rollup	Rolling summary of observed anger tendencies by user.	One active per user + historical versions	Feeds partner-safe summaries and Genome features.
pair_anger_rollup	Shared summary of tension/repair pattern inside a pair.	One active per pair + versions	Forecast consumer; never shown raw to partner.
anger_display_summary	Safe translation layer for user-facing or partner-safe text.	One per visible surface revision	Derived object, not source-of-truth.
anger_registry_code	Canonical lookup table row for codes, labels, event names, decay classes, and display policies.	Controlled code table	Must be versioned.

4. Production event catalog

The missing piece flagged in the audit was not the idea of anger, but the production event schema. The catalog below locks the event labels that developers should implement. Each event must carry an event_label, actor_scope, source_lane, evidence_band, anger_band_before, anger_band_after, and payload_version. Optional fields are attached by event family.

Event label	When emitted	Required fields	Primary write targets
anger.spark_detected	Initial pressure cue appears before full band classification.	trigger_family, cue_source, body_signal_count	episode, event store

Event label	When emitted	Required fields	Primary write targets
anger.band_entered	System enters a new anger band R1-R6.	band_after, source_reason, confidence_band	episode, rollup candidate
anger.band_escalated	Band increases within same episode.	band_before, band_after, escalation_reason	episode, cluster candidate
anger.band_deescalated	Band decreases within same episode.	band_before, band_after, cooling_reason	episode
anger.output_style_detected	Behavioral output style inferred or confirmed.	output_style, evidence_type	episode, Genome feature queue
anger.repair_attempted	User attempts to reconnect, clarify, apologize, or resolve.	repair_type, target_scope	episode, readiness candidate
anger.repair_received	Partner or system receives repair attempt.	repair_type, accepted_state	pair episode
anger.repair_completed	Repair loop resolves enough to close episode.	repair_duration_sec, closure_quality	rollup, forecast support
anger.shutdown_entered	Episode routes to silence / disappear / collapse pattern.	reentry_expected, collapse_marker	episode, cluster candidate
anger.trigger_repeated	Same trigger family repeats in active window.	trigger_family, repeat_count, window_hours	cluster candidate, drift support
anger.decay_tick	Scheduled decayer lowers active weight without deleting history.	decay_class, previous_weight, new_weight	rollup refresh
anger.episode_closed	No longer active; episode sealed.	peak_band, duration_sec, repair_outcome	rollup, archive
anger.partner_safe_summary_updated	Derived summary for visible surfaces changed.	display_band, safety_policy_id	display summary
anger.visibility_suppressed	Raw or translated anger display withheld due to policy.	suppression_reason	display summary, audit
anger.manual_review_flagged	Episode escalated for support / admin review.	reason_code, review_scope	audit queue

4.1 Universal anger_event payload

```
{
  "event_id": "uuid",
  "payload_version": "anger_v1",
  "event_label": "anger.band_entered",
  "event_time_utc": "ISO-8601",
  "actor_scope": "user | pair | system | ai | admin",
  "actor_user_id": "nullable",
  "pair_id": "nullable",
```

```
  "episode_id": "nullable until grouped",
  "source_lane": "assessment | polarity | onboarding | messaging | sanctuary | exs | game | checkin | admin",
  "source_ref": "question_id / message_id / session_id / game_event_id / etc",
  "evidence_band": "EB1 | EB2 | EB3 | EB4",
  "confidence_score": 0.0,
  "anger_band_before": "R0-R6",
  "anger_band_after": "R0-R6",
  "trigger_family": "fairness | disrespect | overwhelm | identity | uncertainty | betrayal | control | boundary | env_disruption | mixed",
  "output_style": "fire | ice | armor | redirect | leadership | repair | unknown",
  "repair_posture": "none | attempted | partial | completed | refused | blocked",
  "visibility_policy": "internal_only | self_visible | partner_safe | suppressed",
  "writeback_class": "none | transient | light | standard | elevated",
  "notes": "optional structured text"
}
```

5. Source-to-feature capture map

Anger Hex does not rely on a single capture lane. The system gets anger data from different sources at different times for different reasons. The source lane matters because it changes confidence, decay speed, visibility, and downstream authority.

Source lane	Where data comes from	What we get	When we get it	Why we keep it
Onboarding	Intake prompts, baseline temperament questions, conflict history, self-described anger style.	Baseline self-report: trigger families, default output style, known cooling habits, prior rupture themes.	At account creation and structured updates.	Creates initial hypotheses and safe defaults before runtime evidence exists.
Scored assessment	Core readiness questions and shadow/behavior items.	Trait-level indicators: fairness sensitivity, control threat, impatience, repair posture, shutdown tendency, pressure logic.	During assessment completion and retakes.	Supports readiness interpretation and cluster priors.
Polarity assessment	Initiation, receptivity, surrender safety, tension attraction, pursuit, pace questions.	Pattern around activation under mismatch, pressure around pursuit/withdrawal, calm-vs-charge preferences.	When polarity section is completed.	Explains why certain relational dynamics activate anger or collapse.
Messaging telemetry	Response timing, escalation markers, punctuation bursts, cold-short sequences, rupture language, silence after heat.	Live runtime anger cues and repair latency.	During active pair communication.	Most valuable operational source for episodes, decay, and pair forecast.
Sanctuary	Emotional Mirror, Whisper Window, Archive, Last Light,	Self-reported and AI-assisted reflective evidence about anger,	When user enters Sanctuary modules.	Adds reflective context and soft confirmation under governance.

Source lane	Where data comes from	What we get	When we get it	Why we keep it
	ExsAnonymous, DM reflection.	overwhelm, hurt, and repair meaning.		
ExsAnonymous	Narratives about prior ruptures, blame loops, repeated trigger stories, repair failures.	Historical anger patterns and residue themes.	Inside Sanctuary-Exs module flow.	Helps detect recurrent trigger families and legacy scripts.
Game telemetry	Validated cooperative/competitive game events and post-game reactions.	Pressure under stakes, fairness response, loss tolerance, cooperative repair, blame routing.	When games are active and event coverage is validated.	Adds stress-behavior evidence in controlled micro-contexts.
Manual check-ins	Structured user reports after a conflict or stressful moment.	Self-labeled band, trigger, repair attempt, and outcome.	When prompted or user-initiated.	Useful for grounding AI and filling gaps without overclaiming.

6. State machine

Anger Hex should be implemented as an episode-aware state machine rather than a free-floating score. The state machine governs how the system knows whether a signal is fresh, active, cooling, closed, or decayed. This is important because readiness, forecast, clusters, and partner-safe summaries should not consume unbounded noise.

State	Meaning	Entry event	Exit event(s)
idle	No active anger episode.	none / episode closed	spark_detected
sparked	Pressure has appeared but band not stabilized yet.	anger.spark_detected	band_entered or idle timeout
active	Episode is live in R1-R6.	band_entered / escalated	repair_attempted, deescalated to R0, shutdown_entered
cooling	Band is decreasing and system is evaluating whether conflict is resolving.	band_deescalated	episode_closed, re-escalation
repairing	Explicit repair action is underway.	repair_attempted	repair_completed, blocked, re-escalation
collapsed	Shutdown/withdrawal path is active.	shutdown_entered	repair_attempted, episode_closed after timeout
closed	Episode has ended but retains fresh influence weight.	episode_closed	decay_tick sequence
decayed	Episode remains historical	scheduled decay threshold	historical only

State	Meaning	Entry event	Exit event(s)
	but no longer has fresh runtime authority.	met	

7. Write-back law

Anger events do not all deserve the same downstream weight. The write-back class tells the system how far an event is allowed to travel. A small annoyance seen once should not behave like repeated rupture without repair. The table below locks the five write-back classes.

Class	When used	Where it may write	What it may not do
none	Noise, ambiguous signals, single low-confidence cue.	event store only	Cannot affect readiness, forecast, or clusters.
transient	Low-confidence but usable short-lived cue.	active episode, ephemeral UI hints	Cannot persist to long-term rollups beyond transient window.
light	Confirmed small event such as R1/R2 with no repetition.	user rollup, Genome low-weight feature	Cannot materially drag forecast alone.
standard	Confirmed moderate episode, repeated trigger, or meaningful repair event.	user rollup, pair rollup, readiness/forecast candidates, cluster engine	Cannot independently rewrite archetype or shadow.
elevated	R4-R6 episode, recurring ruptures, shutdown loops, or repeated failure to repair.	all standard targets plus safety queues and stronger suppression logic	Still cannot act as diagnosis or permanent identity label.

7.1 Specific write-back targets

- Self rollup: stores baseline trigger distribution, output distribution, average peak band, recovery time bands, repair completion rate, and recent active risk flags.
- Pair rollup: stores pair-specific trigger patterns, tension frequency, escalation arc, mutual repair latency, shutdown occurrence, and trend direction.
- Readiness family: anger contributes under controlled caps to relational regulation, repair reliability, and conflict sustainability interpretations.
- Forecast family: anger contributes as a pair-state condition, especially when peak bands are high, repair is absent, or shutdown repeats.
- Cluster engine: anger can seed or confirm Control Avoidance, Push-Pull, Repair Gap, Slow Drift, Readiness Overstatement, and other governed clusters.

- Shadow activation: anger event labels can satisfy dependency thresholds for Stonewaller, Scorekeeper, Self-Saboteur, Love Bomber crash, Manipulator pressure responses, and Chameleon collapse patterns only where the runtime contract allows.
- Sanctuary and AI: anger can trigger reflection prompts, cooling rituals, or repair suggestions, but only through safe text and never as a moral indictment.

8. Decay rules

Decay is mandatory. Without decay, the system becomes cruel and inaccurate by letting old heat over-govern the present. Decay does not erase history; it lowers operational weight over time when no confirming evidence appears. The decay class is assigned per episode.

Decay class	Default examples	Fresh influence window	Half-life intent	Historical posture
D1 fast	R1-R2 one-off irritation resolved quickly.	24-72 hours	Very fast	Reference only after first week.
D2 normal	Moderate but repaired episode.	7-14 days	Moderate	Still visible in trend stats, low operational weight after 30 days.
D3 slow	R4+ episode, or repeated trigger with incomplete repair.	21-45 days	Slow	Remains meaningful in pair trend and forecast until healthy counter-evidence appears.
D4 persistent	R5-R6, repeated shutdown, or recurrent rupture pattern across window.	45-90 days	Very slow	Persists in historical trend and may keep suppression or caution labels active until multiple good counter-signals arrive.

8.1 Decay algorithm stance

- The active_weight of an anger episode decreases on scheduled decay ticks; the event history itself remains immutable.
- Repair quality may accelerate decay. Failed or blocked repair may slow decay.
- Counter-signals such as calm clarification, timely re-entry, and repeated completed repair can reduce future operational weight but do not rewrite the original event.
- Decay never changes the original peak band. It changes only present operational influence.
- Episodes tied to safety reviews may keep an administrative hold even if statistical weight decays.

9. Partner-safe display logic

This is one of the most important sections. Another user should never be shown raw anger telemetry such as 'R5, high fairness-triggered defense, 3 escalation bursts, 480-minute repair

latency.' That is technically interesting and socially radioactive. Partner-safe display logic translates internal evidence into bounded, non-creepy relational language.

Policy lane	Who can see it	Allowed content	Never show
internal_only	System services, audit, governed admin tools.	Raw events, bands, trigger families, timings, dependency labels.	Anything public or direct partner surface.
self_visible	The user who generated the episode.	Band summary, trigger translation, repair suggestions, cooling guidance, trend label.	Sensitive moderator notes or partner-internal data.
partner_safe	Matched partner or pair-shared insight surface.	Soft relational translation such as 'can get sharp under repeated pressure' or 'may need a little time before repair.'	Raw band codes, exact trigger histories, exact timestamps, medicalized or accusatory wording.
suppressed	No end-user surface.	Nothing visible beyond generic notice if needed.	Any specific anger detail.

9.1 Partner-safe translation rules

- Translate behavior into relational safety language, not pathology. Example: 'tends to need cooling time after conflict' instead of 'shutdown risk R6.'
- Use trend bands rather than exact counts where possible. Example: 'repair usually returns within a day' rather than 'median repair latency = 13.4 hours.'
- Never reveal private Sanctuary, ExsAnonymous, or game-specific raw context to a partner.
- Do not surface a partner-safe anger insight unless minimum confidence and minimum recency rules are satisfied.
- If signals are mixed, favor neutral descriptions such as 'stress responses vary by context' rather than hard labels.
- If safety suppression is active, show either nothing or a generic statement controlled by copy policy.

Internal fact pattern	Partner-safe translation
Repeated R3-R4 escalation when messages feel dismissive; repair returns within hours.	Can get tense when they feel brushed off, but usually tries to repair fairly soon.
R5 outburst pattern with low repair completion.	Under heavy pressure, conflict can become sharper than ideal; caution is warranted.
R6 shutdown after repeated fairness trigger; re-entry often delayed.	May pull back when conflict feels unresolved and can need time before reconnecting.
Strong repair completion despite R3 spikes.	Stress shows up, but follow-through after conflict is usually a strength.

10. Calibration with the current Genesis OS system

- Readiness calibration: Anger Hex should influence regulation, repair reliability, and sustainability dimensions, not dominate the final readiness score by itself.
- Forecast calibration: Pair forecast should weight repeated anger + absent repair much more than isolated anger + completed repair.
- Drift calibration: Slow drift is strengthened by repeated cooling failures, shutdown entry, or unresolved low-grade pressure. Drift should not be triggered by one sealed episode alone.
- Shadow calibration: Shadow activation dependencies must rely on event labels and rule packs, not on hand-wavy interpretation. Example: repeated shutdown_entered + blocked repair may satisfy part of Stonewaller activation, but only if the shadow runtime contract says so.
- Genome calibration: The Human Emotion Genome may extract trigger_family distributions, output_style frequencies, and recovery arcs as feature families, but it cannot back-edit workbook scoring law.
- Sanctuary calibration: Sanctuary can use anger summaries to route the user into cooling rituals, reflection prompts, or repair preparation flows. ExsAnonymous may add historic trigger residues but cannot by itself create a high-risk partner label.

11. API surfaces and code-table dependencies

- Code tables required: anger_band_registry, anger_event_registry, anger_trigger_family_registry, anger_output_style_registry, anger_decay_policy_registry, anger_display_policy_registry, anger_writeback_class_registry.
- Service endpoints recommended: POST /anger/events, POST /anger/episodes/close, GET /anger/profile/{user_id}, GET /anger/pair-summary/{pair_id}, POST /anger/manual-checkin, POST /anger/repair-update.
- All visible text must route through copy IDs and display policies, not through raw event labels embedded in the client.
- Every downstream consumer should subscribe to normalized summaries rather than reading raw event tables whenever possible.

12. QA scenarios required before production

- Single R1 irritation in chat, immediate clarification, no repeat. Expected: light write-back only, fast decay, no partner caution surface.
- R3 escalation, two-hour silence, successful repair same day. Expected: standard episode, moderate operational weight, positive repair support preserved.
- R4 to R6 pattern with repeated shutdown and no repair across window. Expected: elevated write-back, slower decay, forecast caution, possible partner-safe suppression rather than raw insight.

- ExsAnonymous narrative suggests chronic anger issues but runtime evidence is thin. Expected: reflective context only; no high-confidence partner-facing summary.
- Game event shows competitive spike but chat and Sanctuary show healthy repair. Expected: game source enriches profile without over-punishing final outputs.
- Multiple mild R2 episodes over weeks with repeated fairness trigger. Expected: cluster candidate for recurring trigger or repair gap depending on repair data; drift support possible.

13. Implementation order

- Phase 1: publish code tables and event registry; implement anger_event and anger_episode objects; wire the state machine.
- Phase 2: connect messaging, Sanctuary, and manual check-in inputs; enable repair and decay events.
- Phase 3: enable rollups, Genome feature extraction, cluster dependencies, and drift/shadow interlocks.
- Phase 4: ship self-visible summaries and strictly bounded partner-safe translations.
- Phase 5: tune thresholds with observed data while preserving audit history and registry versioning.

14. Final contract statement

Anger Hex is now promoted to a formal companion-system contract. In production, it must behave like a disciplined evidence engine: observable, versioned, decaying, bounded, and translation-safe. It must be strong enough to help Genesis OS understand rupture, regulation, and repair, while restrained enough that one hot moment never becomes a permanent sentence.

15. Trigger family registry

Trigger family is the system's explanation for why the anger event exists. Production code should not store free-text triggers when a normalized family can be selected. Free text may be kept as a supporting note, but downstream services consume the normalized trigger family.

Trigger family	Definition	Typical upstream signs	Common downstream concern
fairness	User perceives something as unfair, lopsided, or not honoring a shared standard.	Scorekeeping language, 'that's not right', repeated comparison, rule breach references.	Resentment accumulation, scorekeeper shadow risk, recurring friction.
disrespect	User experiences tone, dismissal, humiliation, or status injury.	Sharp reactions to being brushed off, talked over, mocked, or minimized.	Armor mode, coldness, sharpness escalation.
overwhelm	Too much input, noise, demand, or emotional load arrives at once.	Fast fatigue, short replies, withdrawal after overload, sensory strain.	Collapse or shutdown route, low repair capacity in the moment.
identity	Anger arises because the	Strong statements about code,	Threat/defense pattern; higher

Trigger family	Definition	Typical upstream signs	Common downstream concern
	situation violates how the person sees themselves or wants their environment ordered.	standards, dignity, who they are, or how life should run.	stakes than event size may imply.
uncertainty	Ambiguity, unpredictability, or unstable signals drive irritation or threat.	Repeated clarification requests, anxious control attempts, future pressure.	Control behavior, chasing certainty, escalation under vague signals.
betrayal	Trust injury or perceived disloyalty activates anger.	Cheating narratives, withholding, broken promises, secrecy themes.	High forecast drag if repeated, deep memory persistence.
control_loss	Something removes perceived agency or order.	Reactivity to disruption, forced changes, helplessness cues.	Spike from R1/R2 to R3/R4 under external pressure.
boundary	A line the user believes matters gets crossed.	'I said no', repeated asks, intrusions, access violations.	Repair sensitivity and relational safety questions.
env_disruption	Environmental nuisance or interruption causes activation.	Noise, clutter, interruptions, scheduling chaos, broken flow.	Usually decays fast unless it chains into identity or overwhelm.
mixed	Multiple validated trigger families are active in the same episode.	Conflicting cues with several clear contributors.	Requires careful translation and may slow decay.

16. Output style registry

The band tells the system how hot the anger is; the output style tells the system what the anger does. Both are needed. An R3 event routed through leadership is not the same as an R3 event routed through fire or ice.

Output style	Meaning	What the system looks for	Why it matters
fire	Outward discharge: yelling, snapping, venting, harsh burst.	Escalation wording, intensity spikes, repeated force markers, strong rupture actions.	Highest rupture visibility; may justify stronger suppression or caution.
ice	Silent withdrawal, disengagement, disappearance, refusal to answer.	Long silence after heat, abrupt drop-off, no re-entry without prompt.	Important for drift, shutdown loops, and partner uncertainty.
armor	Controlled but sharpened protection; cold precision, sarcasm, hardness.	Short clipped tone, protective language, defensive structure without full outburst.	Common in people who remain functional but relationally sharp.
redirect	Deflection, topic switch, pseudo-calm, 'I’m fine' masking.	Mismatch between content and tension indicators, evasive move away from rupture.	Can hide unresolved heat and create cluster noise if not tracked.
leadership	Firm but regulated conflict handling.	Clear boundary plus calm request, paced tone, problem-solving posture.	Protective factor; important counterweight for readiness.
repair	Movement toward	Explicit ownership, reconnect	Major positive modifier for

Output style	Meaning	What the system looks for	Why it matters
	reconnection, apology, clarification, or solution.	attempt, future-facing conflict repair.	decay and forecast interpretation.

17. Source-to-anger mapping by intake family

Below is the engineering view of how anger evidence enters from the major question families. This does not override workbook scoring; it tells developers what Anger Hex is allowed to infer, seed, or confirm from each source lane.

Question family	Example evidence type	Anger Hex result
Onboarding temperament	Self-report such as 'I get quiet when I'm mad' or 'I hate disrespect.'	Seed default output_style and likely trigger_family hypotheses at low confidence.
Onboarding history	Past conflict stories, prior rupture pattern, repair self-description.	Create historical residue hints and preferred cooling-path defaults.
Core readiness assessment	Answers showing control threat, impatience, repair follow-through, fairness sensitivity.	Support regulatory and repair priors; may raise confidence on recurring styles.
Shadow questions	Manipulation, withholding, scorekeeping, self-sabotage patterns.	Can satisfy dependency hints for shadow runtime when paired with anger events.
Polarity questions	Pressure under pursuit, surrender safety, pace mismatch, initiation burden, attraction to tension.	Explain why certain pair dynamics trigger anger, collapse, or protest behavior.
Messaging telemetry	Observed runtime words, silence, escalation, re-entry, apology, clarification.	Primary source for production anger episodes and repair latency.
Sanctuary reflection	User names hurt, pressure, resentment, or cooling needs after conflict.	Contextual support and self-visible translation material; limited direct partner impact.
ExsAnonymous narratives	Historical ex patterns tied to disrespect, betrayal, chaos, or repeated unfairness.	Historic trigger map and residue tags; supportive but never sole basis for caution labels.

18. Episode closure, reopen, and merge rules

- Close an episode when the system sees either confirmed repair_completed, confirmed de-escalation to baseline with quiet stabilization, or timeout plus no continuing signal inside the active window.
- Reopen the same episode rather than creating a new one if renewed escalation clearly belongs to the same trigger chain and occurs inside the configured continuation window.

- Create a new episode if the trigger family changes materially, if the quiet period exceeds the continuation window, or if the context lane changes enough to make the event semantically different.
- Merge episodes only under controlled reconciliation logic. Example: a messaging spike followed immediately by a Sanctuary reflection about the same active conflict may enrich the existing episode; it should not duplicate it.
- Once sealed, an episode is immutable except for derived rollups, decay ticks, administrative annotations, and visibility translations.

19. Suppression and review policy registry

Reason code	When used	Visible outcome
SUP_LOW_CONF	Evidence too thin or conflicting for safe translation.	Show nothing or generic 'insufficient context' if needed.
SUP_PRIVATE_CONTEXT	Source depends on private Sanctuary / Exs context that cannot be translated safely.	No partner-facing output; self-visible only where allowed.
SUP_SAFETY_REVIEW	Episode or pattern has been queued for trust/safety or admin review.	Suppress partner-facing detail until review resolved.
SUP_MIXED_SIGNALS	Evidence is mixed enough that any summary would overclaim.	Use only neutral self-visible note or no display.
SUP_RECENCY_GUARD	Event too fresh to translate responsibly without seeing repair arc.	Delay partner-safe display until minimum observation window passes.
SUP_POLICY_HOLD	Policy or jurisdictional copy rules require withholding this surface.	No output until policy release.

20. Examples of partner-safe display packs

Scenario	Self-visible summary	Partner-safe summary	Why difference exists
One-off R2 irritation, fast clarification	You looked mildly activated by disruption, then regulated quickly.	No partner surface.	No need to externalize harmless noise.
Repeated R3 fairness-triggered tension, repair usually same day	Fairness and follow-through matter a lot for your nervous system; repair is usually possible when both people stay engaged.	Can get tense when things feel uneven, but often tries to work it out.	Partner gets relationship-relevant gist, not private mechanism detail.
R4 armor responses under perceived disrespect	You protect yourself with hardness when tone feels off.	May become sharper when they feel disrespected.	Translated away from clinical / internal phrasing.
R6 shutdown after overwhelm with delayed re-entry	When overloaded, you can disappear before you know how to re-engage; re-entry planning may help.	May need a little space before reconnecting after conflict.	Safe, useful, and not humiliating.

Scenario	Self-visible summary	Partner-safe summary	Why difference exists
Repeated blocked repair with high-heat episodes	High-intensity conflict plus unresolved repair is creating a meaningful strain pattern.	Conflict may take longer to settle than ideal; caution is warranted.	Shows caution without exposing the full telemetry stack.

21. Versioning and migration rules

- Any change to anger event labels, decay classes, display policies, or write-back class behavior must increment the anger registry version and be recorded in migration notes.
- Historical anger events remain stored under their original payload_version. Derived rollups may be recomputed, but the source event should not be rewritten.
- If a new display policy would have hidden older partner-safe summaries, new renders must respect current policy even if prior versions were less strict.
- When thresholds are tuned later, audit notes should record whether old episodes were grandfathered, recalculated, or left untouched.