GENESIS OS / RELATIONSHIP SCORES

Master Cluster Atlas & Rule-Pack Source Map - v2

Workbook-aligned redevelopment of the cluster pack using the corrected master blueprint and polarity workbook as the source frame.

Purpose	**Explain every major cluster in the current system, show where it comes from, when it is collected, why it matters, and how onboarding, assessment, polarity, drift, Sanctuary, messaging, shadows, and companion systems feed the rule engine.**
Authoritative references	Genesis_OS_Unified_Master_Blueprint_Corrected RS.docx and genesis_os_final_quantum_scoring_workbook_v10_polarity.xlsx
What this replaces	The shorter Batch 1 packet; this version is designed to act like the cluster atlas your developers can actually build from instead of a polite summary.
Important stance	Question answers and prompts can seed cluster families, but final activation still depends on the rules engine, severity ladder, corroboration logic, confidence rules, and write-back policy.
Scope	Onboarding intake, 50 core assessment questions, 10 shadow questions, 32 polarity prompts, governed clusters, support/watch clusters, companion-state clusters, and cross-system evidence flows.

Executive use note

This document is deliberately long. That is not an accident. The cluster layer is where score law, behavioral evidence, consent, and product surfaces meet. If this layer stays vague, different services will infer different truths and the whole engine starts speaking in dialects. This atlas is meant to stop that.

- The workbook remains the authority for answer-row metrics, pair quantum interactions, polarity axes, and burden fields.
- The corrected master blueprint remains the authority for build order, service boundaries, product surfaces, and object ownership.
- This atlas governs interpretation: what clusters exist, what can seed them, what confirms them, what softens them, what suppresses them, and what each system is allowed to say about them.

1. Cluster operating principles

Clusters are governed interpretations, not stickers.

A cluster is not a personality badge. It is a governed interpretation sitting on top of workbook vectors, burden fields, pair interactions, telemetry, and consent-bounded narrative evidence.

A seed is not a firing.

An onboarding field, assessment answer, polarity answer, or companion signal can nominate a cluster family without activating it. Activation requires threshold logic, evidence sufficiency, confidence, and write-back policy.

Lane discipline matters.

Slow-lane evidence comes from onboarding, assessment, shadow, and scheduled polarity refreshes. Medium-lane evidence comes from messaging rollups, DriftEligibility, couple tools, ExsAnonymous, and repeated Sanctuary behavior. Fast-lane evidence comes from anger events, acute conflicts, and same-day emotional spikes. Fast lane may warn; it may not casually rewrite identity.

The engine must always answer four questions.

What was observed, where did it come from, when was it captured, and why the system cared enough to keep it.

Companion systems may support, soften, or suppress.

Grief, anger, ex-history, and game data are real, but they are modifiers and context engines. They are not excuses to let one subsystem hijack baseline readiness or pair law.

Lane	Typical sources	What it can do	What it cannot do
Slow lane	Onboarding, assessment, shadow section, scheduled polarity refresh	Seed blueprint-level cluster families, baseline risk posture, and governed starting states	Rewrite itself off one dramatic conversation, one grief-heavy week, or one angry night
Medium lane	Messaging rollups, DriftEligibility, couple tools, ExsAnonymous, repeated Sanctuary usage	Confirm persistence, escalation, decay, pair effects, and forecast relevance	Override workbook pair law or bypass threshold gates
Fast lane	Single conflict windows, anger events, acute-state overlays, same-day drift support signals	Create temporary caution signals, hold states, and reflection prompts	Become a permanent readiness rewrite without corroboration

2. Master cluster catalog

The catalog below is the full cluster universe currently visible in the workbook-aligned cluster layer. The first six are governed v1 launch clusters. The rest are support/watch, shadow, or companion-state clusters that either feed the first six, translate evidence, or preserve context for readiness and forecast decisions.

ID	Cluster	Family	Status	Primary sources	When captured	Why the system keeps it
CL_VULN_LOOP	Vulnerability Loop	Governed contradiction	Governed v1	Assessment + messaging + Sanctuary	Initial assessment, then rechecked when closeness, disclosure, or retreat events occur	Captures the pattern where someone wants closeness but retracts once it becomes emotionally real; protects readiness from reading high just because insight language sounds mature.
CL_CONTROL_AVOID	Control Avoidance	Governed contradiction	Governed v1	Assessment + conflict events + messaging	Assessment pass, conflict windows, repair attempts, repeated ambiguity events	Flags the posture of preserving control by delaying clarity, staying vague, or making the other person work too hard for certainty.
CL_SLOW_DRIFT	Slow Drift	Governed behavioral	Governed v1	Assessment + DriftEligibility + messaging cadence + relationship rituals	Only after enough longitudinal evidence exists; never from one bad day alone	Separates temporary distance from a meaningful erosion of reciprocity, warmth, and relational upkeep.
CL_PUSH_PULL	Push-Pull	Governed behavioral	Governed v1	Assessment + polarity + messaging cadence	Assessment seeds it; behavior either confirms or clears it	Tracks the pursue-withdraw rhythm that inflates chemistry while reducing stability and relational safety.
CL_REPAIR_GAP	Repair Gap	Governed conflict	Governed v1	Assessment + conflict/repair telemetry + Sanctuary reflections	Triggered after rupture events, failed repairs, or repeated non-repair	Measures the distance between caring and actually repairing; a major forecast drag factor.
CL_READINESS_OVERSTATE	Readiness Overstatement	Governed readiness	Governed v1	Assessment + behavior deltas + forecast write-backs	After assessment and whenever behavior materially contradicts stated maturity	Prevents polished language, insight, or intention from overstating relational

						readiness when behavior does not hold.
CL_INTEGRITY_STRAIN	Integrity Strain	Core burden	Support/watch	Assessment + shadow items + event ledger	Assessment baseline and any integrity-relevant behavior events	Tracks mismatch between stated values and observed behavior, especially around honesty, accountability, and congruence.
CL_RECIPROCITY_IMBALANCE	Reciprocity Imbalance	Core burden	Support/watch	Assessment + messaging + couple mode	Assessment seeds, then recalculated through interaction data	Captures one-sided effort, uneven pursuit, or chronic emotional labor asymmetry.
CL_INTENT_MISMATCH	Intent Mismatch	Core burden	Support/watch	Onboarding profile + assessment + match packet	At onboarding, at assessment close, and at pair assembly	Prevents high chemistry from being mistaken for aligned direction when relationship intent is structurally mismatched.
CL_PACE_MISMATCH	Pace Mismatch	Core burden	Support/watch	Assessment + matching + drift	Assessment and pair assembly; rechecked during pace shocks	Captures misaligned speed for intimacy, commitment, and relationship progression.
CL_ATTACHMENT_SPIKE	Attachment Spike	Attachment strain	Support/watch	Assessment + messaging + distress windows	Assessment baseline plus sudden silence or uncertainty windows	Tracks abrupt activation under uncertainty: spiraling, freezing, over-functioning, or emotional collapse.
CL_VALIDATION_HUNGER	Validation Dependency	Chemistry-risk	Support/watch	Assessment + polarity + messaging	Assessment, attraction windows, approval-seeking sequences	Captures chemistry inflated by being chosen, mirrored, or reassured rather than by actual relational fit.
CL_SURRENDER_SAFETY_DEFICIT	Surrender Safety Deficit	Chemistry-risk	Support/watch	Assessment + polarity + Sanctuary	Assessment and intimacy escalation windows	Flags when closeness does not feel safe enough to relax into, even if

						attraction is present.
CL_TENSION_BONDING	Tension Bonding	Chemistry-risk	Support/watch	Polarity + messaging conflict cadence	During attraction ramp and conflict repetition	Captures chemistry that is being powered by friction, uncertainty, chaos, or emotional whiplash.
CL_STRUCTURE_COLLISION	Structure Collision	Chemistry-risk	Support/watch	Polarity + onboarding intent + couple mode	At polarity scoring and pair assembly	Flags mismatch in who carries direction, decision-making, and relational shape.
CL_FLEX_RIGIDITY	Flex-Rigidity Collision	Chemistry-risk	Support/watch	Polarity + conflict/repair telemetry	Assessment and conflict windows	Captures stalemates created when one or both people cannot adapt their style under stress.
CL_EROTIC_SAFETY_SPLIT	Erotic-Safety Split	Chemistry-risk	Support/watch	Polarity + relationship mode + user reflections	Assessment, chemistry report generation, and couple-state reviews	Tracks pairings where desire and safety run on separate rails, creating intensity without durable trust.
CL_BOUNDARY_FOG	Boundary Fog	Needs/ boundaries	Support/watch	Assessment + Sanctuary + ExsAnonymous	Assessment and reflective writing windows	Captures unclear limits, over-accommodation, hinting instead of asking, and confusion between love and self-erasure.
CL_FORECAST_DRAG	Forecast Drag	Forecast meta-cluster	Support/watch	Pair engine + drift + repair + outcome telemetry	After pair assembly and over time	Represents accumulated downward pressure on pair stability from multiple lighter burdens even when no single cluster is catastrophic.
CL_SH_MANIPULATOR	Shadow Activation: Manipulator	Shadow activation	Governed shadow	Shadow questions + messaging + moderation events	Assessment baseline, then behavior-confirmed only	Captures strategic guilt, baiting, leverage, or engineered emotional

						positioning.
CL_SH_STONEWALLER	Shadow Activation: Stonewaller	Shadow activation	Governed shadow	Shadow questions + messaging + conflict events	Assessment baseline, then confirmed in shutdown sequences	Captures hard withholding, closure refusal, emotional lockout, and punitive silence.
CL_SH_LOVEBOMBER	Shadow Activation: Love Bomber	Shadow activation	Governed shadow	Shadow questions + messaging ramp + drift aftermath	Fast attachment windows, reconciliation bursts, over-promise moments	Captures intensity spikes that exceed actual stability, pacing, or follow-through.
CL_SH_CHAMELEON	Shadow Activation: Chameleon	Shadow activation	Governed shadow	Shadow questions + messaging tone shifts + Sanctuary writing	Assessment baseline, then confirmed when identity presentation shifts for safety or approval	Captures adaptive self-erasure, mirroring, and identity drift to preserve connection.
CL_SH_SCOREKEEPER	Shadow Activation: Scorekeeper	Shadow activation	Governed shadow	Shadow questions + reciprocity events + couple conflict logs	Assessment baseline and repeated fairness/receipt moments	Captures ledger-style relating, hidden resentments, and conditional giving.
CL_SH_SELFSABOTEUR	Shadow Activation: Self-Saboteur	Shadow activation	Governed shadow	Shadow questions + retreat events + grief/attachment strain	Assessment baseline, then escalates under perceived rejection or hope	Captures pre-emptive withdrawal, self-blame, and collapse before reality is clarified.
CL_ANGER_ESCALATION	Anger Escalation	Companion state	Companion state	Anger Hex + event labeler + repair telemetry	Real-time to daily in high-conflict windows	Separates irritation, protest, shutdown, and retaliatory expression so the system knows whether safety and repair are improving or worsening.
CL_GRIEF_LOAD	Grief / Depletion Load	Companion state	Companion state	E_GRIEF + Sanctuary + Last Light + user check-ins	Daily to milestone windows; always decays with time and evidence	Prevents depletion, loss, and spiritual drag from being mistaken for identity truth or permanent compatibility failure.
CL_EX_TETHER	Ex Tether /	Companion	Companion	ExsAnonymous	Whenever ex-	Tracks whether

	Narrative Carryover	state	state	+ Sanctuary archive + reflections	processing activity is active or recurring	old relational loops are still shaping present trust, comparison, or emotional availability.
CL_GAME_COMPENSATION	Game Compensation / Bond Distortion	Companion state	Companion state	RS Game Ecosystem + couple sessions + telemetry	During game sessions and post-session analysis	Captures when games are healthy bonding versus a substitute for direct intimacy, clarity, or repair.

3. Zero-pass onboarding intake - what is captured before the scored questions

These are not scored like the 60-question assessment, but they still matter. They shape eligibility, visibility, observability, routing, and later cluster interpretation. Several do not create a cluster by themselves; instead they decide whether the right clusters can run later and whether the system has permission to narrate what it sees.

Field / prompt	System lane	Cluster or rule impact	When captured	Why kept
RS code / identity handle	Account creation	No cluster by itself; enables protected singles telemetry, game identity, and pair joins	Immediately at account creation	Keeps identity consistent across systems without exposing raw private IDs.
Dating status selector	Onboarding profile	CL_INTENT_MISMATCH; CL_PACE_MISMATCH	At onboarding and whenever status changes	Lets the engine separate attraction from actual relational direction.
Relationship state	Onboarding profile	CL_SLOW_DRIFT context; CL_FORECAST_DRAG context	At onboarding and every state change	Determines whether singles logic, couple logic, or recovery logic should run.
Gender + DOB	Onboarding profile	No direct cluster; supports zodiac / numerology keys and cohort analysis	At onboarding	Needed for overlay keys and future calibration, not for punitive clustering.
Sun sign / life path keys	Onboarding profile	No direct cluster; pair packet context only	At onboarding or after birth data derivation	Supports the 12% cosmic overlay without inventing behavior truth.
Location + search radius	Onboarding profile	Operational only; pair availability context	At onboarding and preference edits	Prevents candidate generation from reading like

				incompatibility when it is really search-scope friction.
Visibility settings	Onboarding profile	Eligibility / suppression control only	At onboarding and privacy edits	Controls whether the system may surface a user at all.
Ghost Mode flag	Onboarding profile	Suppresses public drift interpretations while user is hidden	At activation/deactivation	Stops hidden self-work from being misread as relational fading.
AI consent	Governance intake	No cluster; controls whether AI may explain or enrich evidence	At onboarding and consent edits	Cluster evidence can still exist deterministically, but AI may not narrate or expand it without consent.
Chat-analysis consent	Governance intake	Enables behavior confirmation for CL_SLOW_DRIFT, CL_PUSH_PULL, CL_REPAIR_GAP, shadow activation	At onboarding and consent edits	Separates deterministic score law from opt-in behavioral enrichment.
Sanctuary visibility / use permissions	Governance intake	Enables CL_GRIEF_LOAD, CL_EX_TETHER, CL_BOUNDARY_FOG support signals	At onboarding and Sanctuary setup	Keeps private healing surfaces private while still letting the owner receive accurate support.
Intention / preference prompts	Profile setup	CL_INTENT_MISMATCH; CL_READINESS_OVERSTATE context	At onboarding and preference edits	Prevents the product from rewarding vague desire more than declared direction.

4. Full assessment source map - all 60 scored questions

The workbook gives the row-level answer vectors, burden fields, and dimensions. This section maps each question to the cluster families it can seed. Seeding is not the same as firing. The engine still decides whether the cluster becomes active, visible, penalized, or merely stored as supporting evidence.

Questions 1-8 | CommunicationHealth / EmotionalNeeds

Q#	Question source theme	Primary cluster families seeded	Workbook burden hooks	When captured	Why the system keeps it
1	Reading tone shifts and 'weird energy'	Vulnerability Loop, Control Avoidance, Attachment Spike	cluster 0.79; integrity 0.00; risk 0.33; forecast 1.00	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
2	What happens after your vulnerability is shrugged off	Vulnerability Loop, Repair Gap, Boundary Fog	cluster 1.00; integrity 0.00; risk 0.67; forecast 1.00	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence

3	How you show early interest safely	Vulnerability Loop, Push-Pull, Surrender Safety Deficit	cluster 0.50; integrity 0.00; risk 0.28; forecast 1.00	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
4	How you handle early red flags	Control Avoidance, Intent Mismatch, Readiness Overstatement	cluster 0.36; integrity 0.00; risk 0.67; forecast 1.00	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
5	Not feeling emotionally met after opening up	Vulnerability Loop, Boundary Fog, Repair Gap	cluster 0.93; integrity 0.00; risk 0.67; forecast 1.00	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
6	Reaction to hot-and-cold behavior	Slow Drift, Push-Pull, Attachment Spike	cluster 0.64; integrity 0.00; risk 0.67; forecast 1.00	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
7	Response to silence after closeness	Attachment Spike, Vulnerability Loop, Control Avoidance	cluster 0.36; integrity 0.00; risk 0.44; forecast 1.00	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
8	Detecting a fade in consistency	Slow Drift, Push-Pull, Attachment Spike	cluster 0.79; integrity 0.00; risk 0.33; forecast 0.60	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence

Questions 9-16 | ConflictRepair / CommunicationHealth

Q#	Question source theme	Primary cluster families seeded	Workbook burden hooks	When captured	Why the system keeps it
9	Reaction to feeling deprioritized	Reciprocity Imbalance, Control Avoidance, Boundary Fog	cluster 0.36; integrity 0.00; risk 0.33; forecast 0.20	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
10	Being leaned on without reciprocity	Reciprocity Imbalance, Boundary Fog, Readiness Overstatement	cluster 0.36; integrity 0.00; risk 0.00; forecast 0.60	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
11	When someone suddenly pulls back	Slow Drift, Push-Pull, Control Avoidance	cluster 0.43; integrity 0.00; risk 0.67; forecast 0.80	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
12	Sensing a change and testing for clarity	Control Avoidance, Attachment Spike, Slow Drift	cluster 0.93; integrity 0.00; risk 0.44; forecast 1.00	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
13	How interest becomes visible over time	Surrender Safety Deficit, Vulnerability Loop, Push-Pull	cluster 0.00; integrity 0.00; risk 0.33; forecast 0.80	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
14	Matching reduced effort and mixed signals	Reciprocity Imbalance, Reciprocity Imbalance, Push-Pull	cluster 0.36; integrity 0.00; risk 1.00; forecast 0.40	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
15	Receiving	Validation	cluster 0.36;	Initial assessment /	Seeds slow-lane

	unusually healthy interest	Dependency, Surrender Safety Deficit, Readiness Overstatement	integrity 0.00; risk 0.00; forecast 0.40	versioned retake	blueprint evidence
16	Second pass on hot-and-cold inconsistency	Push-Pull, Slow Drift, Attachment Spike	cluster 0.64; integrity 0.00; risk 0.67; forecast 1.00	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence

Questions 17-24 | EmotionalRegulation / AttachmentSafety

Q#	Question source theme	Primary cluster families seeded	Workbook burden hooks	When captured	Why the system keeps it
17	Partner goes quiet under emotion	Attachment Spike, Control Avoidance, Surrender Safety Deficit	cluster 0.36; integrity 0.00; risk 0.67; forecast 1.00	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
18	Response to pullback after feeling seen	Push-Pull, Attachment Spike, Slow Drift	cluster 0.36; integrity 0.00; risk 0.44; forecast 0.40	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
19	Curiosity imbalance and one-sided knowing	Reciprocity Imbalance, Boundary Fog, Control Avoidance	cluster 0.57; integrity 0.00; risk 0.11; forecast 1.00	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
20	Mismatch between chemistry and readiness	Intent Mismatch, Pace Mismatch, Readiness Overstatement	cluster 0.36; integrity 0.00; risk 0.67; forecast 1.00	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
21	Repeated hot-and-cold pattern test	Push-Pull, Slow Drift, Attachment Spike	cluster 0.64; integrity 0.00; risk 0.67; forecast 1.00	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
22	Repeated silence / ambiguity test	Attachment Spike, Control Avoidance, Surrender Safety Deficit	cluster 0.36; integrity 0.00; risk 0.67; forecast 1.00	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
23	Energy mismatch and overinvestment threshold	Reciprocity Imbalance, Pace Mismatch, Slow Drift	cluster 0.79; integrity 0.00; risk 0.67; forecast 0.80	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
24	Effort that goes unacknowledged	Reciprocity Imbalance, Shadow Activation: Scorekeeper, Repair Gap	cluster 0.36; integrity 0.00; risk 0.33; forecast 0.80	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence

Questions 25-32 | ValuesLifestyle / IntentAlignment

Q#	Question source theme	Primary cluster families seeded	Workbook burden hooks	When captured	Why the system keeps it
25	Unclear mutual feeling / who likes who	Intent Mismatch, Boundary Fog, Pace Mismatch	cluster 0.36; integrity 0.00; risk 0.33; forecast 0.80	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence

26	Unmet ask and follow-through reliability	Reciprocity Imbalance, Repair Gap, Boundary Fog	cluster 0.36; integrity 0.00; risk 0.33; forecast 0.60	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
27	Feeling awkward after emotional disclosure	Vulnerability Loop, Shadow Activation: Scorekeeper, Boundary Fog	cluster 0.36; integrity 0.00; risk 0.44; forecast 0.60	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
28	Suspected dishonesty and trust rupture	Integrity Strain, Control Avoidance, Attachment Spike	cluster 0.57; integrity 0.00; risk 0.67; forecast 1.00	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
29	Naming or denying drift	Slow Drift, Push-Pull, Control Avoidance	cluster 0.36; integrity 0.00; risk 0.33; forecast 1.00	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
30	Vulnerability landing badly or being mocked	Vulnerability Loop, Surrender Safety Deficit, Boundary Fog	cluster 0.36; integrity 0.00; risk 0.33; forecast 1.00	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
31	What you do after hurting someone	Repair Gap, Control Avoidance, Shadow Activation: Self-Saboteur	cluster 0.36; integrity 0.00; risk 0.33; forecast 1.00	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
32	One-sided initiation and outreach load	Reciprocity Imbalance, Slow Drift, Push-Pull	cluster 0.57; integrity 0.00; risk 0.67; forecast 1.00	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence

Questions 33-40 | CommitmentPace / IntentAlignment

Q#	Question source theme	Primary cluster families seeded	Workbook burden hooks	When captured	Why the system keeps it
33	How you leave when commitment is wrong	Slow Drift, Control Avoidance, Pace Mismatch	cluster 0.36; integrity 0.00; risk 0.44; forecast 0.40	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
34	Feeling used or conditionally wanted	Reciprocity Imbalance, Boundary Fog, Integrity Strain	cluster 0.36; integrity 0.00; risk 1.00; forecast 0.60	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
35	Receiving feedback that you caused pain	Repair Gap, Control Avoidance, Attachment Spike	cluster 0.36; integrity 0.00; risk 0.11; forecast 1.00	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
36	Being left for someone else / replaced	Shadow Activation: Self-Saboteur, Validation Dependency, Attachment Spike	cluster 0.79; integrity 0.00; risk 0.33; forecast 0.60	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
37	How you process internal upset before expression	Control Avoidance, Attachment Spike, Repair Gap	cluster 0.79; integrity 0.00; risk 0.33; forecast 0.60	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
38	What happens when someone sees you deeply	Surrender Safety Deficit, Vulnerability Loop,	cluster 0.36; integrity 0.00; risk 0.33; forecast 0.60	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence

		Validation Dependency			
39	How you respond to another person's overwhelm	Repair Gap, Boundary Fog, Control Avoidance	cluster 0.36; integrity 0.00; risk 0.33; forecast 0.80	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
40	The hidden reason you want partnership	Validation Dependency, Boundary Fog, Readiness Overstatement	cluster 0.43; integrity 0.00; risk 0.00; forecast 0.80	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence

Questions 41-50 | EmotionalNeeds / IntimacyThreshold

Q#	Question source theme	Primary cluster families seeded	Workbook burden hooks	When captured	Why the system keeps it
41	Who you are attracted to and why	Validation Dependency, Surrender Safety Deficit, Push-Pull	cluster 0.43; integrity 0.00; risk 0.11; forecast 0.80	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
42	Your default self-sabotage pattern in closeness	Shadow Activation: Self-Saboteur, Vulnerability Loop, Boundary Fog	cluster 0.36; integrity 0.00; risk 0.11; forecast 0.40	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
43	Primary fear once love becomes real	Surrender Safety Deficit, Shadow Activation: Self-Saboteur, Vulnerability Loop	cluster 0.36; integrity 0.00; risk 0.33; forecast 0.80	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
44	Why feelings stay unspoken	Control Avoidance, Boundary Fog, Vulnerability Loop	cluster 0.36; integrity 0.00; risk 0.67; forecast 0.80	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
45	What actually helps you reopen	Repair Gap, Surrender Safety Deficit, Vulnerability Loop	cluster 0.36; integrity 0.00; risk 0.33; forecast 1.00	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
46	How you metabolize heartbreak	Grief / Depletion Load, Shadow Activation: Self-Saboteur, Control Avoidance	cluster 0.36; integrity 0.00; risk 0.33; forecast 1.00	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
47	How healthy care feels in your body	Surrender Safety Deficit, Validation Dependency, Vulnerability Loop	cluster 0.36; integrity 0.00; risk 0.33; forecast 0.60	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
48	What destabilizes closeness once it becomes real	Surrender Safety Deficit, Attachment Spike, Push-Pull	cluster 0.36; integrity 0.00; risk 0.11; forecast 0.60	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
49	Which old wound still governs love	Ex Tether / Narrative Carryover, Vulnerability Loop,	cluster 0.36; integrity 0.00; risk 0.00; forecast 0.60	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence

		Shadow Activation: Self-Saboteur			
50	What scares you about healthy love	Surrender Safety Deficit, Grief / Depletion Load, Vulnerability Loop	cluster 0.36; integrity 0.00; risk 0.11; forecast 0.80	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence

Questions 100-109 | ShadowPattern / IntegrityRisk

Q#	Question source theme	Primary cluster families seeded	Workbook burden hooks	When captured	Why the system keeps it
100	Shadow reaction when attraction or respect slips	Shadow Activation: Manipulator, Shadow Activation: Stonewaller, Shadow Activation: Self-Saboteur	cluster 0.00; integrity 0.60; risk 0.67; forecast 0.60	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
101	Shadow response when your story stops matching	Integrity Strain, Shadow Activation: Manipulator, Shadow Activation: Chameleon	cluster 0.00; integrity 0.60; risk 0.33; forecast 0.60	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
102	Shadow move after being hurt or unseen	Shadow Activation: Scorekeeper, Shadow Activation: Stonewaller, Shadow Activation: Manipulator	cluster 0.00; integrity 0.60; risk 0.33; forecast 0.60	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
103	Shadow move when the other person wants space	Shadow Activation: Stonewaller, Shadow Activation: Chameleon, Shadow Activation: Self-Saboteur	cluster 0.00; integrity 0.60; risk 0.00; forecast 0.60	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
104	Shadow move after exposed vulnerability lands badly	Vulnerability Loop, Shadow Activation: Stonewaller, Shadow Activation: Chameleon	cluster 0.00; integrity 0.60; risk 0.00; forecast 0.60	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
105	Shadow move after forgetfulness / small hurt	Shadow Activation: Scorekeeper, Shadow Activation: Stonewaller, Shadow Activation: Self-Saboteur	cluster 0.00; integrity 0.60; risk 0.33; forecast 0.60	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
106	Shadow move when you feel them drifting	Slow Drift, Shadow Activation: Love Bomber, Shadow Activation: Chameleon	cluster 0.00; integrity 0.60; risk 0.33; forecast 0.60	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
107	Shadow apology style after causing harm	Repair Gap, Shadow Activation: Manipulator,	cluster 0.00; integrity 0.60; risk 0.00; forecast 0.60	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence

		Shadow Activation: Love Bomber			
108	Shadow move in active conflict escalation	Anger Escalation, Shadow Activation: Stonewaller, Shadow Activation: Scorekeeper	cluster 0.00; integrity 0.60; risk 0.00; forecast 0.60	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence
109	Shadow move when something emotional must be said	Boundary Fog, Shadow Activation: Scorekeeper, Shadow Activation: Self-Saboteur	cluster 0.00; integrity 0.60; risk 0.33; forecast 0.60	Initial assessment / versioned retake	Seeds slow-lane blueprint evidence

5. Full polarity source map - all 32 polarity prompts

The polarity package exists to translate attraction style into measurable chemistry-risk signals. In the workbook, 32 prompts are organized into four repeating eight-question blocks. The repeated note-pattern is intentional: each block deepens the same axis families under higher weighting and lets the engine measure consistency, strain, and drift exposure instead of relying on one dramatic answer.

Polarity Q1-Q8

Q#	Polarity prompt note	Primary axes	Secondary axes	Likely cluster families seeded	When captured
Q1	Plan shifts / adaptation	structure_preference, flexibility	initiation, leadership	Push-Pull, Structure Collision, Control Avoidance, Flex-Rigidity Collision	At polarity completion and any future polarity version refresh
Q2	Silence and pause	calm_receptivity, surrender_safety	flexibility, structure_preference	Surrender Safety Deficit, Vulnerability Loop, Structure Collision, Flex-Rigidity Collision	At polarity completion and any future polarity version refresh
Q3	Choice under uncertainty	initiation, leadership	flexibility, validation_dependency	Push-Pull, Structure Collision, Control Avoidance, Flex-Rigidity Collision	At polarity completion and any future polarity version refresh
Q4	Early friction / unresolved tension	tension_attraction, calm_receptivity	surrender_safety, structure_preference	Surrender Safety Deficit, Vulnerability Loop, Tension Bonding, Structure Collision	At polarity completion and any future polarity version refresh
Q5	Time pressure	initiation, structure_preference	flexibility, repair_endurance	Push-Pull, Structure Collision, Control Avoidance, Flex-Rigidity Collision	At polarity completion and any future polarity version refresh
Q6	Emotional disclosure	surrender_safety, erotic_safety_alignment	calm_receptivity, validation_dependency	Surrender Safety Deficit, Vulnerability Loop, Validation	At polarity completion and any future polarity version refresh

				Dependency, Erotic-Safety Split	
Q7	Conflicting preferences	leadership, structure_preference	tension_attraction, calm_receptivity	Push-Pull, Structure Collision, Control Avoidance, Surrender Safety Deficit	At polarity completion and any future polarity version refresh
Q8	Being observed / exposure	validation_dependency, flexibility	surrender_safety, structure_preference	Surrender Safety Deficit, Vulnerability Loop, Structure Collision, Flex-Rigidity Collision	At polarity completion and any future polarity version refresh

Polarity Q9-Q16

Q#	Polarity prompt note	Primary axes	Secondary axes	Likely cluster families seeded	When captured
Q9	Plan shifts / adaptation	structure_preference, flexibility	initiation, leadership	Push-Pull, Structure Collision, Control Avoidance, Flex-Rigidity Collision	At polarity completion and any future polarity version refresh
Q10	Silence and pause	calm_receptivity, surrender_safety	flexibility, structure_preference	Surrender Safety Deficit, Vulnerability Loop, Structure Collision, Flex-Rigidity Collision	At polarity completion and any future polarity version refresh
Q11	Choice under uncertainty	initiation, leadership	flexibility, validation_dependency	Push-Pull, Structure Collision, Control Avoidance, Flex-Rigidity Collision	At polarity completion and any future polarity version refresh
Q12	Early friction / unresolved tension	tension_attraction, calm_receptivity	surrender_safety, structure_preference	Surrender Safety Deficit, Vulnerability Loop, Tension Bonding, Structure Collision	At polarity completion and any future polarity version refresh
Q13	Time pressure	initiation, structure_preference	flexibility, repair_endurance	Push-Pull, Structure Collision, Control Avoidance, Flex-Rigidity Collision	At polarity completion and any future polarity version refresh
Q14	Emotional disclosure	surrender_safety, erotic_safety_alignment	calm_receptivity, validation_dependency	Surrender Safety Deficit, Vulnerability Loop, Validation Dependency, Erotic-Safety Split	At polarity completion and any future polarity version refresh
Q15	Conflicting preferences	leadership, structure_preference	tension_attraction, calm_receptivity	Push-Pull, Structure Collision, Control Avoidance, Surrender Safety Deficit	At polarity completion and any future polarity version refresh

Q16	Being observed / exposure	validation_dependency, flexibility	surrender_safety, structure_preference	Surrender Safety Deficit, Vulnerability Loop, Structure Collision, Flex-Rigidity Collision	At polarity completion and any future polarity version refresh

Polarity Q17-Q24

Q#	Polarity prompt note	Primary axes	Secondary axes	Likely cluster families seeded	When captured
Q17	Plan shifts / adaptation	structure_preference, flexibility	initiation, leadership	Push-Pull, Structure Collision, Control Avoidance, Flex-Rigidity Collision	At polarity completion and any future polarity version refresh
Q18	Silence and pause	calm_receptivity, surrender_safety	flexibility, structure_preference	Surrender Safety Deficit, Vulnerability Loop, Structure Collision, Flex-Rigidity Collision	At polarity completion and any future polarity version refresh
Q19	Choice under uncertainty	initiation, leadership	flexibility, validation_dependency	Push-Pull, Structure Collision, Control Avoidance, Flex-Rigidity Collision	At polarity completion and any future polarity version refresh
Q20	Early friction / unresolved tension	tension_attraction, calm_receptivity	surrender_safety, structure_preference	Surrender Safety Deficit, Vulnerability Loop, Tension Bonding, Structure Collision	At polarity completion and any future polarity version refresh
Q21	Time pressure	initiation, structure_preference	flexibility, repair_endurance	Push-Pull, Structure Collision, Control Avoidance, Flex-Rigidity Collision	At polarity completion and any future polarity version refresh
Q22	Emotional disclosure	surrender_safety, erotic_safety_alignment	calm_receptivity, validation_dependency	Surrender Safety Deficit, Vulnerability Loop, Validation Dependency, Erotic-Safety Split	At polarity completion and any future polarity version refresh
Q23	Conflicting preferences	leadership, structure_preference	tension_attraction, calm_receptivity	Push-Pull, Structure Collision, Control Avoidance, Surrender Safety Deficit	At polarity completion and any future polarity version refresh
Q24	Being observed / exposure	validation_dependency, flexibility	surrender_safety, structure_preference	Surrender Safety Deficit, Vulnerability Loop, Structure Collision, Flex-Rigidity Collision	At polarity completion and any future polarity version refresh

Polarity Q25-Q32

Q#	Polarity prompt note	Primary axes	Secondary axes	Likely cluster families seeded	When captured
Q25	Plan shifts / adaptation	structure_preference, flexibility	initiation, leadership	Push-Pull, Structure Collision, Control Avoidance, Flex-Rigidity Collision	At polarity completion and any future polarity version refresh
Q26	Silence and pause	calm_receptivity, surrender_safety	flexibility, structure_preference	Surrender Safety Deficit, Vulnerability Loop, Structure Collision, Flex-Rigidity Collision	At polarity completion and any future polarity version refresh
Q27	Choice under uncertainty	initiation, leadership	flexibility, validation_dependency	Push-Pull, Structure Collision, Control Avoidance, Flex-Rigidity Collision	At polarity completion and any future polarity version refresh
Q28	Early friction / unresolved tension	tension_attraction, calm_receptivity	surrender_safety, structure_preference	Surrender Safety Deficit, Vulnerability Loop, Tension Bonding, Structure Collision	At polarity completion and any future polarity version refresh
Q29	Time pressure	initiation, structure_preference	flexibility, repair_endurance	Push-Pull, Structure Collision, Control Avoidance, Flex-Rigidity Collision	At polarity completion and any future polarity version refresh
Q30	Emotional disclosure	surrender_safety, erotic_safety_alignment	calm_receptivity, validation_dependency	Surrender Safety Deficit, Vulnerability Loop, Validation Dependency, Erotic-Safety Split	At polarity completion and any future polarity version refresh
Q31	Conflicting preferences	leadership, structure_preference	tension_attraction, calm_receptivity	Push-Pull, Structure Collision, Control Avoidance, Surrender Safety Deficit	At polarity completion and any future polarity version refresh
Q32	Being observed / exposure	validation_dependency, flexibility	surrender_safety, structure_preference	Surrender Safety Deficit, Vulnerability Loop, Structure Collision, Flex-Rigidity Collision	At polarity completion and any future polarity version refresh

6. Every system that can create, support, confirm, soften, or suppress clusters

This is the cross-system map. Some modules only seed. Some confirm. Some create companion-state context. Some suppress escalation so the engine does not confuse grief, acute anger, or fantasy chemistry with durable relational truth.

System / module	Cluster work it performs	When it captures evidence	Why that evidence matters
Assessment engine (50 core)	Primary seed source for governed contradiction, needs, pacing, repair, regulation, and readiness clusters	Single completed assessment plus versioned rescoring windows	This is the slow-lane truth layer; it seeds clusters before behavior confirms or clears them.
Shadow questions (10)	Primary seed source for six shadow-activation clusters and integrity strain context	Assessment completion and scheduled retakes	These answers tell the engine which unhealthy move someone rationalizes when pressure rises.
Polarity engine (32 questions)	Primary seed source for chemistry-risk clusters like CL_SURRENDER_SAFETY_DEFICIT, CL_TENSION_BONDING, CL_STRUCTURE_COLLISION, CL_EROTIC_SAFETY_SPLIT	Polarity test completion and versioned recompute	Chemistry can feel incredible while still being unstable. These questions separate magnetism from durability.
Matching service / pair packet	Generates pair-side burden reads: CL_INTENT_MISMATCH, CL_PACE_MISMATCH, CL_FORECAST_DRAG	Every pair compute	Some clusters are not visible until two blueprints touch each other.
Messaging telemetry	Confirms or clears CL_SLOW_DRIFT, CL_PUSH_PULL, CL_RECIPROCITY_IMBALANCE, CL_REPAIR_GAP, shadow activations	Event-driven, real time to daily rollups	Behavior either backs up assessment posture or exposes the gap between self-story and lived pattern.
Knock system	Supports intent and reciprocity evidence; reveals threshold behavior under limited access	At every Knock send / accept / ignore event	How someone handles partial access is often more diagnostic than how they behave once access is guaranteed.
Sanctuary: Stillness / Emotional Mirror / Whisper Window	Supports CL_GRIEF_LOAD, CL_BOUNDARY_FOG, CL_VULN_LOOP, CL_EX_TETHER	On every opted-in reflection session	Private self-work often reveals the inner loop long before public behavior does.
Archive of Self / Alignment Compass	Medium-lane support for trend persistence and forecast drag	Daily to milestone reviews	Gives the system memory without letting one dramatic day rewrite identity.
DriftEligibility	Owns CL_SLOW_DRIFT state transitions and drift reason mix	Daily / threshold-based review only after enough evidence exists	Separates true fading from grief, overload, privacy, or one-off life noise.
Shadow Activation service	Confirms when a seeded shadow becomes behaviorally active	Conflict windows, abandonment windows, apology windows, reconciliation bursts	A baseline shadow type is not the same as active expression. This service prevents lazy over-labeling.
Emotional Genome engine	Produces signal families that support multiple clusters	Fast to medium lane depending on source	Turns tone, phrasing, ritual use, and behavior into

	with confidence and freshness		governed features instead of vibes.
Anger Hex / Anger map	Supports CL_ANGER_ESCALATION, CL_REPAIR_GAP, CL_PUSH_PULL	Real time to daily in conflict periods	Lets the engine distinguish protest, cold withdrawal, retaliatory expression, and actual repair potential.
E_GRIEF / Last Light	Supports CL_GRIEF_LOAD and protects against false identity conclusions	Daily to milestone cadence with decay	Loss and depletion change behavior temporarily; this system stops the engine from mistaking pain for permanent character.
ExsAnonymous	Supports CL_EX_TETHER, CL_BOUNDARY_FOG, CL_READINESS_OVERSTATE context	Whenever ex-processing is active	Old attachments often keep writing the script for the next relationship unless the system tracks carryover cleanly.
RS Game Ecosystem	Supports CL_GAME_COMPENSATION, CL_RECIPROCITY_IMBALANCE, rhythm / bond telemetry	Per session and post-session summary	Games can expose healthy play, pursuit rhythm, or avoidance disguised as bonding.
Relationship mode / couple tools	Confirms pair-side repair, reciprocity, ritual upkeep, and forecast drag	After pair confirmation / relationship mode activation	This is where theory meets maintenance. Stable pairs must keep showing the work.
Outcome telemetry / FeedbackData	Supports CL_FORECAST_DRAG calibration and future threshold tuning	7-day, 30-day, 90-day outcome windows	Without outcome truth, forecasting becomes expensive fan fiction.

7. Cluster capture lifecycle

Lifecycle point	What becomes available	Why it matters
Account creation	Identity, intent, consent, eligibility context	No direct score mutation; determines which later clusters are observable and whether AI may narrate them.
Assessment completion	Baseline slow-lane cluster seeding from 50 core + 10 shadow questions	This is the first durable cluster snapshot.
Polarity completion	Chemistry-risk cluster seeding	Adds the attraction-risk layer that the core assessment cannot fully see.
First pair compute	Pair-only burdens such as intent mismatch, pace mismatch, forecast drag context	Some clusters only become meaningful when two people are compared.
Early messaging window	Push-pull, reciprocity imbalance, repair gap, drift support, shadow activation support	Behavior begins to confirm or soften what the assessment predicted.
Conflict / rupture window	Repair gap, anger escalation, stonewalling, scorekeeping, control	High-value moments because coping style gets exposed under pressure.

	avoidance	
Sanctuary / reflection usage	Grief load, ex tether, boundary fog, vulnerability support signals	Private data often explains why public behavior is off without turning that pain into permanent identity.
Relationship mode / couple tools	Forecast drag, ritual upkeep, bond quality, pair repair health	This is where the system can finally distinguish chemistry from maintenance.
Outcome windows (7/30/90 days)	Forecast calibration and threshold refinement	Without outcome truth, clusters never mature into reliable forecast inputs.

8. Detailed cluster profiles

Each profile below is designed to answer the five questions your devs actually need answered: what the cluster is, where the system gets it from, when evidence enters the engine, why the system keeps it, and how the cluster should be treated in the wider architecture.

CL_CONTROL_AVOID - Control Avoidance

Control Avoidance sits in the governed contradiction family and is currently classified as governed v1. Its primary evidence path is assessment + conflict events + messaging. The base timing rule is: assessment pass, conflict windows, repair attempts, repeated ambiguity events. The reason the engine preserves it is straightforward: Flags the posture of preserving control by delaying clarity, staying vague, or making the other person work too hard for certainty.

What it is	**A governed or support-state reading of behavior, posture, or context. It should never be treated as a personality sentence; it is an operational signal that can strengthen, soften, suppress, or redirect readiness and forecast interpretation depending on corroboration.**
Where we get it	Assessment + conflict events + messaging
When we get it	Assessment pass, conflict windows, repair attempts, repeated ambiguity events
Why we keep it	Flags the posture of preserving control by delaying clarity, staying vague, or making the other person work too hard for certainty.
Direct onboarding hooks	None directly seeded at onboarding; onboarding may still create context, eligibility, or permission conditions.
Direct system emitters	No dedicated emitter line item in the high-level system table; rely on the primary-source path above and the rule engine.

Assessment seeding. The slow-lane assessment currently seeds this cluster from 18 scored questions, including Q1: Reading tone shifts and 'weird energy', Q4: How you handle early red flags, Q7: Response to silence after closeness, Q9: Reaction to feeling deprioritized, Q11: When someone suddenly pulls back, Q12: Sensing a change and testing for clarity, Q17: Partner goes quiet under emotion, Q19: Curiosity imbalance and one-sided knowing, Q22: Repeated silence / ambiguity test, Q28: Suspected dishonesty and trust rupture, Q29: Naming or denying drift, Q31: What you do after hurting someone

and 6 more seeded assessment links. These do not activate the cluster by themselves; they provide blueprint-level evidence that later behavior either confirms, contradicts, or cools down.

Polarity seeding. The polarity engine points at this cluster in 16 prompt slots, including Q1 (Plan shifts / adaptation), Q3 (Choice under uncertainty), Q5 (Time pressure), Q7 (Conflicting preferences), Q9 (Plan shifts / adaptation), Q11 (Choice under uncertainty), Q13 (Time pressure), Q15 (Conflicting preferences), Q17 (Plan shifts / adaptation), Q19 (Choice under uncertainty) and 6 more polarity links. In practice this means the cluster often rides on attraction-style misalignment, structure-flow mismatch, surrender safety strain, or chemistry that outruns stability.

Onboarding context. There is no direct onboarding trigger mapped here. Even so, onboarding still affects observability through identity, consent, relationship state, visibility, and intention routing.

Implementation guardrails:

- Do not let one isolated event, one private reflection, or one consent-bounded surface become a permanent identity rewrite.

Source family	Where it appears now	When it matters most	What developers should do with it
Primary sources	Assessment + conflict events + messaging	Assessment pass, conflict windows, repair attempts, repeated ambiguity events	Wire these as allowed emitters and keep source hashes or event references.
Assessment map	18 mapped questions	Assessment completion, rescoring windows	Use as seed evidence, not instant activation.
Polarity map	16 mapped prompts	Polarity completion and polarity refresh	Use as chemistry-risk context or direct seed only when the map explicitly names it.
Onboarding context	0 direct fields	Account creation and profile edits	Use as routing, eligibility, context, and permission logic.
Cross-system emitters	See primary-source path	During lived behavior windows	Use for confirmation, suppression, or explanation depending on lane discipline.

CL_PUSH_PULL - Push-Pull

Push-Pull sits in the governed behavioral family and is currently classified as governed v1. Its primary evidence path is assessment + polarity + messaging cadence. The base timing rule is: assessment seeds it; behavior either confirms or clears it. The reason the engine preserves it is straightforward: Tracks the pursue-withdraw rhythm that inflates chemistry while reducing stability and relational safety.

What it is	A governed or support-state reading of behavior, posture, or context. It should never be treated as a personality sentence; it is an operational signal that can strengthen, soften, suppress, or redirect readiness and forecast interpretation depending on corroboration.
Where we get it	Assessment + polarity + messaging cadence
When we get it	Assessment seeds it; behavior either confirms or clears it

Why we keep it	Tracks the pursue-withdraw rhythm that inflates chemistry while reducing stability and relational safety.
Direct onboarding hooks	None directly seeded at onboarding; onboarding may still create context, eligibility, or permission conditions.
Direct system emitters	No dedicated emitter line item in the high-level system table; rely on the primary-source path above and the rule engine.

Assessment seeding. The slow-lane assessment currently seeds this cluster from 13 scored questions, including Q3: How you show early interest safely, Q6: Reaction to hot-and-cold behavior, Q8: Detecting a fade in consistency, Q11: When someone suddenly pulls back, Q13: How interest becomes visible over time, Q14: Matching reduced effort and mixed signals, Q16: Second pass on hot-and-cold inconsistency, Q18: Response to pullback after feeling seen, Q21: Repeated hot-and-cold pattern test, Q29: Naming or denying drift, Q32: One-sided initiation and outreach load, Q41: Who you are attracted to and why and 1 more seeded assessment links. These do not activate the cluster by themselves; they provide blueprint-level evidence that later behavior either confirms, contradicts, or cools down.

Polarity seeding. The polarity engine points at this cluster in 16 prompt slots, including Q1 (Plan shifts / adaptation), Q3 (Choice under uncertainty), Q5 (Time pressure), Q7 (Conflicting preferences), Q9 (Plan shifts / adaptation), Q11 (Choice under uncertainty), Q13 (Time pressure), Q15 (Conflicting preferences), Q17 (Plan shifts / adaptation), Q19 (Choice under uncertainty) and 6 more polarity links. In practice this means the cluster often rides on attraction-style misalignment, structure-flow mismatch, surrender safety strain, or chemistry that outruns stability.

Onboarding context. There is no direct onboarding trigger mapped here. Even so, onboarding still affects observability through identity, consent, relationship state, visibility, and intention routing.

Implementation guardrails:

- Do not let one isolated event, one private reflection, or one consent-bounded surface become a permanent identity rewrite.

Source family	Where it appears now	When it matters most	What developers should do with it
Primary sources	Assessment + polarity + messaging cadence	Assessment seeds it; behavior either confirms or clears it	Wire these as allowed emitters and keep source hashes or event references.
Assessment map	13 mapped questions	Assessment completion, rescoring windows	Use as seed evidence, not instant activation.
Polarity map	16 mapped prompts	Polarity completion and polarity refresh	Use as chemistry-risk context or direct seed only when the map explicitly names it.
Onboarding context	0 direct fields	Account creation and profile edits	Use as routing, eligibility, context, and permission logic.
Cross-system emitters	See primary-source path	During lived behavior windows	Use for confirmation, suppression, or explanation depending on lane discipline.

CL_READINESS_OVERSTATE - Readiness Overstatement

Readiness Overstatement sits in the governed readiness family and is currently classified as governed v1. Its primary evidence path is assessment + behavior deltas + forecast write-backs. The base timing rule is: after assessment and whenever behavior materially contradicts stated maturity. The reason the engine preserves it is straightforward: Prevents polished language, insight, or intention from overstating relational readiness when behavior does not hold.

What it is	A governed or support-state reading of behavior, posture, or context. It should never be treated as a personality sentence; it is an operational signal that can strengthen, soften, suppress, or redirect readiness and forecast interpretation depending on corroboration.
Where we get it	Assessment + behavior deltas + forecast write-backs
When we get it	After assessment and whenever behavior materially contradicts stated maturity
Why we keep it	Prevents polished language, insight, or intention from overstating relational readiness when behavior does not hold.
Direct onboarding hooks	None directly seeded at onboarding; onboarding may still create context, eligibility, or permission conditions.
Direct system emitters	No dedicated emitter line item in the high-level system table; rely on the primary-source path above and the rule engine.

Assessment seeding. The slow-lane assessment currently seeds this cluster from 5 scored questions, including Q4: How you handle early red flags, Q10: Being leaned on without reciprocity, Q15: Receiving unusually healthy interest, Q20: Mismatch between chemistry and readiness, Q40: The hidden reason you want partnership. These do not activate the cluster by themselves; they provide blueprint-level evidence that later behavior either confirms, contradicts, or cools down.

Polarity seeding. The current polarity map does not explicitly name this cluster. That usually means polarity acts as background chemistry context rather than as a direct seed source.

Onboarding context. There is no direct onboarding trigger mapped here. Even so, onboarding still affects observability through identity, consent, relationship state, visibility, and intention routing.

Implementation guardrails:

- Use this to protect the public readiness surface from fake certainty; do not weaponize it as punishment for optimism.

Source family	Where it appears now	When it matters most	What developers should do with it
Primary sources	Assessment + behavior deltas + forecast write-backs	After assessment and whenever behavior materially contradicts stated maturity	Wire these as allowed emitters and keep source hashes or event references.
Assessment map	5 mapped questions	Assessment completion, rescoring windows	Use as seed evidence, not instant activation.
Polarity map	0 mapped prompts	Polarity completion and	Use as chemistry-risk context

		polarity refresh	or direct seed only when the map explicitly names it.
Onboarding context	0 direct fields	Account creation and profile edits	Use as routing, eligibility, context, and permission logic.
Cross-system emitters	See primary-source path	During lived behavior windows	Use for confirmation, suppression, or explanation depending on lane discipline.

CL_REPAIR_GAP - Repair Gap

Repair Gap sits in the governed conflict family and is currently classified as governed v1. Its primary evidence path is assessment + conflict/repair telemetry + sanctuary reflections. The base timing rule is: triggered after rupture events, failed repairs, or repeated non-repair. The reason the engine preserves it is straightforward: Measures the distance between caring and actually repairing; a major forecast drag factor.

What it is	A governed or support-state reading of behavior, posture, or context. It should never be treated as a personality sentence; it is an operational signal that can strengthen, soften, suppress, or redirect readiness and forecast interpretation depending on corroboration.
Where we get it	Assessment + conflict/repair telemetry + Sanctuary reflections
When we get it	Triggered after rupture events, failed repairs, or repeated non-repair
Why we keep it	Measures the distance between caring and actually repairing; a major forecast drag factor.
Direct onboarding hooks	None directly seeded at onboarding; onboarding may still create context, eligibility, or permission conditions.
Direct system emitters	No dedicated emitter line item in the high-level system table; rely on the primary-source path above and the rule engine.

Assessment seeding. The slow-lane assessment currently seeds this cluster from 10 scored questions, including Q2: What happens after your vulnerability is shrugged off, Q5: Not feeling emotionally met after opening up, Q24: Effort that goes unacknowledged, Q26: Unmet ask and follow-through reliability, Q31: What you do after hurting someone, Q35: Receiving feedback that you caused pain, Q37: How you process internal upset before expression, Q39: How you respond to another person's overwhelm, Q45: What actually helps you reopen, Q107: Shadow apology style after causing harm. These do not activate the cluster by themselves; they provide blueprint-level evidence that later behavior either confirms, contradicts, or cools down.

Polarity seeding. The current polarity map does not explicitly name this cluster. That usually means polarity acts as background chemistry context rather than as a direct seed source.

Onboarding context. There is no direct onboarding trigger mapped here. Even so, onboarding still affects observability through identity, consent, relationship state, visibility, and intention routing.

Implementation guardrails:

- Do not let one isolated event, one private reflection, or one consent-bounded surface become a permanent identity rewrite.

Source family	Where it appears now	When it matters most	What developers should do with it
Primary sources	Assessment + conflict/repair telemetry + Sanctuary reflections	Triggered after rupture events, failed repairs, or repeated non-repair	Wire these as allowed emitters and keep source hashes or event references.
Assessment map	10 mapped questions	Assessment completion, rescoring windows	Use as seed evidence, not instant activation.
Polarity map	0 mapped prompts	Polarity completion and polarity refresh	Use as chemistry-risk context or direct seed only when the map explicitly names it.
Onboarding context	0 direct fields	Account creation and profile edits	Use as routing, eligibility, context, and permission logic.
Cross-system emitters	See primary-source path	During lived behavior windows	Use for confirmation, suppression, or explanation depending on lane discipline.

CL_SLOW_DRIFT - Slow Drift

Slow Drift sits in the governed behavioral family and is currently classified as governed v1. Its primary evidence path is assessment + drifteligibility + messaging cadence + relationship rituals. The base timing rule is: only after enough longitudinal evidence exists; never from one bad day alone. The reason the engine preserves it is straightforward: Separates temporary distance from a meaningful erosion of reciprocity, warmth, and relational upkeep.

What it is	A governed or support-state reading of behavior, posture, or context. It should never be treated as a personality sentence; it is an operational signal that can strengthen, soften, suppress, or redirect readiness and forecast interpretation depending on corroboration.
Where we get it	Assessment + DriftEligibility + messaging cadence + relationship rituals
When we get it	Only after enough longitudinal evidence exists; never from one bad day alone
Why we keep it	Separates temporary distance from a meaningful erosion of reciprocity, warmth, and relational upkeep.
Direct onboarding hooks	None directly seeded at onboarding; onboarding may still create context, eligibility, or permission conditions.
Direct system emitters	No dedicated emitter line item in the high-level system table; rely on the primary-source path above and the rule engine.

Assessment seeding. The slow-lane assessment currently seeds this cluster from 12 scored questions, including Q6: Reaction to hot-and-cold behavior, Q8: Detecting a fade in consistency, Q11: When someone suddenly pulls back, Q12: Sensing a change and testing for clarity, Q16: Second pass on hot-

and-cold inconsistency, Q18: Response to pullback after feeling seen, Q21: Repeated hot-and-cold pattern test, Q23: Energy mismatch and overinvestment threshold, Q29: Naming or denying drift, Q32: One-sided initiation and outreach load, Q33: How you leave when commitment is wrong, Q106: Shadow move when you feel them drifting. These do not activate the cluster by themselves; they provide blueprint-level evidence that later behavior either confirms, contradicts, or cools down.

Polarity seeding. The current polarity map does not explicitly name this cluster. That usually means polarity acts as background chemistry context rather than as a direct seed source.

Onboarding context. There is no direct onboarding trigger mapped here. Even so, onboarding still affects observability through identity, consent, relationship state, visibility, and intention routing.

Implementation guardrails:

- Never infer this off one argument, one quiet day, or one overloaded week; DriftEligibility and rollup windows must be satisfied.

Source family	Where it appears now	When it matters most	What developers should do with it
Primary sources	Assessment + DriftEligibility + messaging cadence + relationship rituals	Only after enough longitudinal evidence exists; never from one bad day alone	Wire these as allowed emitters and keep source hashes or event references.
Assessment map	12 mapped questions	Assessment completion, rescoring windows	Use as seed evidence, not instant activation.
Polarity map	0 mapped prompts	Polarity completion and polarity refresh	Use as chemistry-risk context or direct seed only when the map explicitly names it.
Onboarding context	0 direct fields	Account creation and profile edits	Use as routing, eligibility, context, and permission logic.
Cross-system emitters	See primary-source path	During lived behavior windows	Use for confirmation, suppression, or explanation depending on lane discipline.

CL_VULN_LOOP - Vulnerability Loop

Vulnerability Loop sits in the governed contradiction family and is currently classified as governed v1. Its primary evidence path is assessment + messaging + sanctuary. The base timing rule is: initial assessment, then rechecked when closeness, disclosure, or retreat events occur. The reason the engine preserves it is straightforward: Captures the pattern where someone wants closeness but retracts once it becomes emotionally real; protects readiness from reading high just because insight language sounds mature.

What it is	A governed or support-state reading of behavior, posture, or context. It should never be treated as a personality sentence; it is an operational signal that can strengthen, soften, suppress, or redirect readiness and forecast interpretation depending on corroboration.
Where we get it	Assessment + messaging + Sanctuary
When we get it	Initial assessment, then rechecked when closeness,

	disclosure, or retreat events occur
Why we keep it	Captures the pattern where someone wants closeness but retracts once it becomes emotionally real; protects readiness from reading high just because insight language sounds mature.
Direct onboarding hooks	None directly seeded at onboarding; onboarding may still create context, eligibility, or permission conditions.
Direct system emitters	No dedicated emitter line item in the high-level system table; rely on the primary-source path above and the rule engine.

Assessment seeding. The slow-lane assessment currently seeds this cluster from 17 scored questions, including Q1: Reading tone shifts and 'weird energy', Q2: What happens after your vulnerability is shrugged off, Q3: How you show early interest safely, Q5: Not feeling emotionally met after opening up, Q7: Response to silence after closeness, Q13: How interest becomes visible over time, Q27: Feeling awkward after emotional disclosure, Q30: Vulnerability landing badly or being mocked, Q38: What happens when someone sees you deeply, Q42: Your default self-sabotage pattern in closeness, Q43: Primary fear once love becomes real, Q44: Why feelings stay unspoken and 5 more seeded assessment links. These do not activate the cluster by themselves; they provide blueprint-level evidence that later behavior either confirms, contradicts, or cools down.

Polarity seeding. The polarity engine points at this cluster in 16 prompt slots, including Q2 (Silence and pause), Q4 (Early friction / unresolved tension), Q6 (Emotional disclosure), Q8 (Being observed / exposure), Q10 (Silence and pause), Q12 (Early friction / unresolved tension), Q14 (Emotional disclosure), Q16 (Being observed / exposure), Q18 (Silence and pause), Q20 (Early friction / unresolved tension) and 6 more polarity links. In practice this means the cluster often rides on attraction-style misalignment, structure-flow mismatch, surrender safety strain, or chemistry that outruns stability.

Onboarding context. There is no direct onboarding trigger mapped here. Even so, onboarding still affects observability through identity, consent, relationship state, visibility, and intention routing.

Implementation guardrails:

- Do not collapse every healthy pause into this cluster; the engine must distinguish regulated pacing from retreat after closeness.

Source family	Where it appears now	When it matters most	What developers should do with it
Primary sources	Assessment + messaging + Sanctuary	Initial assessment, then rechecked when closeness, disclosure, or retreat events occur	Wire these as allowed emitters and keep source hashes or event references.
Assessment map	17 mapped questions	Assessment completion, rescoring windows	Use as seed evidence, not instant activation.
Polarity map	16 mapped prompts	Polarity completion and polarity refresh	Use as chemistry-risk context or direct seed only when the map explicitly names it.
Onboarding context	0 direct fields	Account creation and profile edits	Use as routing, eligibility, context, and permission logic.

Cross-system emitters	See primary-source path	During lived behavior windows	Use for confirmation, suppression, or explanation depending on lane discipline.

CL_SH_CHAMELEON - Shadow Activation: Chameleon

Shadow Activation: Chameleon sits in the shadow activation family and is currently classified as governed shadow. Its primary evidence path is shadow questions + messaging tone shifts + sanctuary writing. The base timing rule is: assessment baseline, then confirmed when identity presentation shifts for safety or approval. The reason the engine preserves it is straightforward: Captures adaptive self-erasure, mirroring, and identity drift to preserve connection.

What it is	A governed or support-state reading of behavior, posture, or context. It should never be treated as a personality sentence; it is an operational signal that can strengthen, soften, suppress, or redirect readiness and forecast interpretation depending on corroboration.
Where we get it	Shadow questions + messaging tone shifts + Sanctuary writing
When we get it	Assessment baseline, then confirmed when identity presentation shifts for safety or approval
Why we keep it	Captures adaptive self-erasure, mirroring, and identity drift to preserve connection.
Direct onboarding hooks	None directly seeded at onboarding; onboarding may still create context, eligibility, or permission conditions.
Direct system emitters	No dedicated emitter line item in the high-level system table; rely on the primary-source path above and the rule engine.

Assessment seeding. The slow-lane assessment currently seeds this cluster from 4 scored questions, including Q101: Shadow response when your story stops matching, Q103: Shadow move when the other person wants space, Q104: Shadow move after exposed vulnerability lands badly, Q106: Shadow move when you feel them drifting. These do not activate the cluster by themselves; they provide blueprint-level evidence that later behavior either confirms, contradicts, or cools down.

Polarity seeding. The current polarity map does not explicitly name this cluster. That usually means polarity acts as background chemistry context rather than as a direct seed source.

Onboarding context. There is no direct onboarding trigger mapped here. Even so, onboarding still affects observability through identity, consent, relationship state, visibility, and intention routing.

Implementation guardrails:

- A shadow type in the questionnaire is not the same as active shadow pressure. Activation requires pressure, corroboration, and current-state evidence.

Source family	Where it appears now	When it matters most	What developers should do with it
Primary sources	Shadow questions + messaging tone shifts + Sanctuary writing	Assessment baseline, then confirmed when identity presentation shifts for safety	Wire these as allowed emitters and keep source hashes or

		or approval	event references.
Assessment map	4 mapped questions	Assessment completion, rescoring windows	Use as seed evidence, not instant activation.
Polarity map	0 mapped prompts	Polarity completion and polarity refresh	Use as chemistry-risk context or direct seed only when the map explicitly names it.
Onboarding context	0 direct fields	Account creation and profile edits	Use as routing, eligibility, context, and permission logic.
Cross-system emitters	See primary-source path	During lived behavior windows	Use for confirmation, suppression, or explanation depending on lane discipline.

CL_SH_LOVEBOMBER - Shadow Activation: Love Bomber

Shadow Activation: Love Bomber sits in the shadow activation family and is currently classified as governed shadow. Its primary evidence path is shadow questions + messaging ramp + drift aftermath. The base timing rule is: fast attachment windows, reconciliation bursts, over-promise moments. The reason the engine preserves it is straightforward: Captures intensity spikes that exceed actual stability, pacing, or follow-through.

What it is	A governed or support-state reading of behavior, posture, or context. It should never be treated as a personality sentence; it is an operational signal that can strengthen, soften, suppress, or redirect readiness and forecast interpretation depending on corroboration.
Where we get it	Shadow questions + messaging ramp + drift aftermath
When we get it	Fast attachment windows, reconciliation bursts, over-promise moments
Why we keep it	Captures intensity spikes that exceed actual stability, pacing, or follow-through.
Direct onboarding hooks	None directly seeded at onboarding; onboarding may still create context, eligibility, or permission conditions.
Direct system emitters	No dedicated emitter line item in the high-level system table; rely on the primary-source path above and the rule engine.

Assessment seeding. The slow-lane assessment currently seeds this cluster from 2 scored questions, including Q106: Shadow move when you feel them drifting, Q107: Shadow apology style after causing harm. These do not activate the cluster by themselves; they provide blueprint-level evidence that later behavior either confirms, contradicts, or cools down.

Polarity seeding. The current polarity map does not explicitly name this cluster. That usually means polarity acts as background chemistry context rather than as a direct seed source.

Onboarding context. There is no direct onboarding trigger mapped here. Even so, onboarding still affects observability through identity, consent, relationship state, visibility, and intention routing.

Implementation guardrails:

- A shadow type in the questionnaire is not the same as active shadow pressure. Activation requires pressure, corroboration, and current-state evidence.

Source family	Where it appears now	When it matters most	What developers should do with it
Primary sources	Shadow questions + messaging ramp + drift aftermath	Fast attachment windows, reconciliation bursts, over-promise moments	Wire these as allowed emitters and keep source hashes or event references.
Assessment map	2 mapped questions	Assessment completion, rescoring windows	Use as seed evidence, not instant activation.
Polarity map	0 mapped prompts	Polarity completion and polarity refresh	Use as chemistry-risk context or direct seed only when the map explicitly names it.
Onboarding context	0 direct fields	Account creation and profile edits	Use as routing, eligibility, context, and permission logic.
Cross-system emitters	See primary-source path	During lived behavior windows	Use for confirmation, suppression, or explanation depending on lane discipline.

CL_SH_MANIPULATOR - Shadow Activation: Manipulator

Shadow Activation: Manipulator sits in the shadow activation family and is currently classified as governed shadow. Its primary evidence path is shadow questions + messaging + moderation events. The base timing rule is: assessment baseline, then behavior-confirmed only. The reason the engine preserves it is straightforward: Captures strategic guilt, baiting, leverage, or engineered emotional positioning.

What it is	A governed or support-state reading of behavior, posture, or context. It should never be treated as a personality sentence; it is an operational signal that can strengthen, soften, suppress, or redirect readiness and forecast interpretation depending on corroboration.
Where we get it	Shadow questions + messaging + moderation events
When we get it	Assessment baseline, then behavior-confirmed only
Why we keep it	Captures strategic guilt, baiting, leverage, or engineered emotional positioning.
Direct onboarding hooks	None directly seeded at onboarding; onboarding may still create context, eligibility, or permission conditions.
Direct system emitters	No dedicated emitter line item in the high-level system table; rely on the primary-source path above and the rule engine.

Assessment seeding. The slow-lane assessment currently seeds this cluster from 4 scored questions, including Q100: Shadow reaction when attraction or respect slips, Q101: Shadow response when your story stops matching, Q102: Shadow move after being hurt or unseen, Q107: Shadow apology style after causing harm. These do not activate the cluster by themselves; they provide blueprint-level evidence that later behavior either confirms, contradicts, or cools down.

Polarity seeding. The current polarity map does not explicitly name this cluster. That usually means polarity acts as background chemistry context rather than as a direct seed source.

Onboarding context. There is no direct onboarding trigger mapped here. Even so, onboarding still affects observability through identity, consent, relationship state, visibility, and intention routing.

Implementation guardrails:

- A shadow type in the questionnaire is not the same as active shadow pressure. Activation requires pressure, corroboration, and current-state evidence.

Source family	Where it appears now	When it matters most	What developers should do with it
Primary sources	Shadow questions + messaging + moderation events	Assessment baseline, then behavior-confirmed only	Wire these as allowed emitters and keep source hashes or event references.
Assessment map	4 mapped questions	Assessment completion, rescoring windows	Use as seed evidence, not instant activation.
Polarity map	0 mapped prompts	Polarity completion and polarity refresh	Use as chemistry-risk context or direct seed only when the map explicitly names it.
Onboarding context	0 direct fields	Account creation and profile edits	Use as routing, eligibility, context, and permission logic.
Cross-system emitters	See primary-source path	During lived behavior windows	Use for confirmation, suppression, or explanation depending on lane discipline.

CL_SH_SCOREKEEPER - Shadow Activation: Scorekeeper

Shadow Activation: Scorekeeper sits in the shadow activation family and is currently classified as governed shadow. Its primary evidence path is shadow questions + reciprocity events + couple conflict logs. The base timing rule is: assessment baseline and repeated fairness/receipt moments. The reason the engine preserves it is straightforward: Captures ledger-style relating, hidden resentments, and conditional giving.

What it is	A governed or support-state reading of behavior, posture, or context. It should never be treated as a personality sentence; it is an operational signal that can strengthen, soften, suppress, or redirect readiness and forecast interpretation depending on corroboration.
Where we get it	Shadow questions + reciprocity events + couple conflict logs
When we get it	Assessment baseline and repeated fairness/receipt moments
Why we keep it	Captures ledger-style relating, hidden resentments, and conditional giving.
Direct onboarding hooks	None directly seeded at onboarding; onboarding may still create context, eligibility, or permission conditions.
Direct system emitters	No dedicated emitter line item in the high-level system table; rely on the primary-source path above and the rule engine.

Assessment seeding. The slow-lane assessment currently seeds this cluster from 6 scored questions, including Q24: Effort that goes unacknowledged, Q27: Feeling awkward after emotional disclosure, Q102: Shadow move after being hurt or unseen, Q105: Shadow move after forgetfulness / small hurt, Q108: Shadow move in active conflict escalation, Q109: Shadow move when something emotional must be said. These do not activate the cluster by themselves; they provide blueprint-level evidence that later behavior either confirms, contradicts, or cools down.

Polarity seeding. The current polarity map does not explicitly name this cluster. That usually means polarity acts as background chemistry context rather than as a direct seed source.

Onboarding context. There is no direct onboarding trigger mapped here. Even so, onboarding still affects observability through identity, consent, relationship state, visibility, and intention routing.

Implementation guardrails:

- A shadow type in the questionnaire is not the same as active shadow pressure. Activation requires pressure, corroboration, and current-state evidence.

Source family	Where it appears now	When it matters most	What developers should do with it
Primary sources	Shadow questions + reciprocity events + couple conflict logs	Assessment baseline and repeated fairness/receipt moments	Wire these as allowed emitters and keep source hashes or event references.
Assessment map	6 mapped questions	Assessment completion, rescoring windows	Use as seed evidence, not instant activation.
Polarity map	0 mapped prompts	Polarity completion and polarity refresh	Use as chemistry-risk context or direct seed only when the map explicitly names it.
Onboarding context	0 direct fields	Account creation and profile edits	Use as routing, eligibility, context, and permission logic.
Cross-system emitters	See primary-source path	During lived behavior windows	Use for confirmation, suppression, or explanation depending on lane discipline.

CL_SH_SELFSABOTEUR - Shadow Activation: Self-Saboteur

Shadow Activation: Self-Saboteur sits in the shadow activation family and is currently classified as governed shadow. Its primary evidence path is shadow questions + retreat events + grief/attachment strain. The base timing rule is: assessment baseline, then escalates under perceived rejection or hope. The reason the engine preserves it is straightforward: Captures pre-emptive withdrawal, self-blame, and collapse before reality is clarified.

What it is	A governed or support-state reading of behavior, posture, or context. It should never be treated as a personality sentence; it is an operational signal that can strengthen, soften, suppress, or redirect readiness and forecast interpretation depending on corroboration.
Where we get it	Shadow questions + retreat events + grief/attachment strain
When we get it	Assessment baseline, then escalates under perceived

	rejection or hope
Why we keep it	Captures pre-emptive withdrawal, self-blame, and collapse before reality is clarified.
Direct onboarding hooks	None directly seeded at onboarding; onboarding may still create context, eligibility, or permission conditions.
Direct system emitters	No dedicated emitter line item in the high-level system table; rely on the primary-source path above and the rule engine.

Assessment seeding. The slow-lane assessment currently seeds this cluster from 10 scored questions, including Q31: What you do after hurting someone, Q36: Being left for someone else / replaced, Q42: Your default self-sabotage pattern in closeness, Q43: Primary fear once love becomes real, Q46: How you metabolize heartbreak, Q49: Which old wound still governs love, Q100: Shadow reaction when attraction or respect slips, Q103: Shadow move when the other person wants space, Q105: Shadow move after forgetfulness / small hurt, Q109: Shadow move when something emotional must be said. These do not activate the cluster by themselves; they provide blueprint-level evidence that later behavior either confirms, contradicts, or cools down.

Polarity seeding. The current polarity map does not explicitly name this cluster. That usually means polarity acts as background chemistry context rather than as a direct seed source.

Onboarding context. There is no direct onboarding trigger mapped here. Even so, onboarding still affects observability through identity, consent, relationship state, visibility, and intention routing.

Implementation guardrails:

- A shadow type in the questionnaire is not the same as active shadow pressure. Activation requires pressure, corroboration, and current-state evidence.

Source family	Where it appears now	When it matters most	What developers should do with it
Primary sources	Shadow questions + retreat events + grief/attachment strain	Assessment baseline, then escalates under perceived rejection or hope	Wire these as allowed emitters and keep source hashes or event references.
Assessment map	10 mapped questions	Assessment completion, rescoring windows	Use as seed evidence, not instant activation.
Polarity map	0 mapped prompts	Polarity completion and polarity refresh	Use as chemistry-risk context or direct seed only when the map explicitly names it.
Onboarding context	0 direct fields	Account creation and profile edits	Use as routing, eligibility, context, and permission logic.
Cross-system emitters	See primary-source path	During lived behavior windows	Use for confirmation, suppression, or explanation depending on lane discipline.

CL_SH_STONEWALLER - Shadow Activation: Stonewaller

Shadow Activation: Stonewaller sits in the shadow activation family and is currently classified as governed shadow. Its primary evidence path is shadow questions + messaging + conflict events. The

base timing rule is: assessment baseline, then confirmed in shutdown sequences. The reason the engine preserves it is straightforward: Captures hard withholding, closure refusal, emotional lockout, and punitive silence.

What it is	A governed or support-state reading of behavior, posture, or context. It should never be treated as a personality sentence; it is an operational signal that can strengthen, soften, suppress, or redirect readiness and forecast interpretation depending on corroboration.
Where we get it	Shadow questions + messaging + conflict events
When we get it	Assessment baseline, then confirmed in shutdown sequences
Why we keep it	Captures hard withholding, closure refusal, emotional lockout, and punitive silence.
Direct onboarding hooks	None directly seeded at onboarding; onboarding may still create context, eligibility, or permission conditions.
Direct system emitters	No dedicated emitter line item in the high-level system table; rely on the primary-source path above and the rule engine.

Assessment seeding. The slow-lane assessment currently seeds this cluster from 6 scored questions, including Q100: Shadow reaction when attraction or respect slips, Q102: Shadow move after being hurt or unseen, Q103: Shadow move when the other person wants space, Q104: Shadow move after exposed vulnerability lands badly, Q105: Shadow move after forgetfulness / small hurt, Q108: Shadow move in active conflict escalation. These do not activate the cluster by themselves; they provide blueprint-level evidence that later behavior either confirms, contradicts, or cools down.

Polarity seeding. The current polarity map does not explicitly name this cluster. That usually means polarity acts as background chemistry context rather than as a direct seed source.

Onboarding context. There is no direct onboarding trigger mapped here. Even so, onboarding still affects observability through identity, consent, relationship state, visibility, and intention routing.

Implementation guardrails:

- A shadow type in the questionnaire is not the same as active shadow pressure. Activation requires pressure, corroboration, and current-state evidence.

Source family	Where it appears now	When it matters most	What developers should do with it
Primary sources	Shadow questions + messaging + conflict events	Assessment baseline, then confirmed in shutdown sequences	Wire these as allowed emitters and keep source hashes or event references.
Assessment map	6 mapped questions	Assessment completion, rescoring windows	Use as seed evidence, not instant activation.
Polarity map	0 mapped prompts	Polarity completion and polarity refresh	Use as chemistry-risk context or direct seed only when the map explicitly names it.
Onboarding context	0 direct fields	Account creation and profile edits	Use as routing, eligibility, context, and permission logic.

Cross-system emitters	See primary-source path	During lived behavior windows	Use for confirmation, suppression, or explanation depending on lane discipline.

CL_ATTACHMENT_SPIKE - Attachment Spike

Attachment Spike sits in the attachment strain family and is currently classified as support/watch. Its primary evidence path is assessment + messaging + distress windows. The base timing rule is: assessment baseline plus sudden silence or uncertainty windows. The reason the engine preserves it is straightforward: Tracks abrupt activation under uncertainty: spiraling, freezing, over-functioning, or emotional collapse.

What it is	A governed or support-state reading of behavior, posture, or context. It should never be treated as a personality sentence; it is an operational signal that can strengthen, soften, suppress, or redirect readiness and forecast interpretation depending on corroboration.
Where we get it	Assessment + messaging + distress windows
When we get it	Assessment baseline plus sudden silence or uncertainty windows
Why we keep it	Tracks abrupt activation under uncertainty: spiraling, freezing, over-functioning, or emotional collapse.
Direct onboarding hooks	None directly seeded at onboarding; onboarding may still create context, eligibility, or permission conditions.
Direct system emitters	No dedicated emitter line item in the high-level system table; rely on the primary-source path above and the rule engine.

Assessment seeding. The slow-lane assessment currently seeds this cluster from 15 scored questions, including Q1: Reading tone shifts and 'weird energy', Q6: Reaction to hot-and-cold behavior, Q7: Response to silence after closeness, Q8: Detecting a fade in consistency, Q12: Sensing a change and testing for clarity, Q16: Second pass on hot-and-cold inconsistency, Q17: Partner goes quiet under emotion, Q18: Response to pullback after feeling seen, Q21: Repeated hot-and-cold pattern test, Q22: Repeated silence / ambiguity test, Q28: Suspected dishonesty and trust rupture, Q35: Receiving feedback that you caused pain and 3 more seeded assessment links. These do not activate the cluster by themselves; they provide blueprint-level evidence that later behavior either confirms, contradicts, or cools down.

Polarity seeding. The current polarity map does not explicitly name this cluster. That usually means polarity acts as background chemistry context rather than as a direct seed source.

Onboarding context. There is no direct onboarding trigger mapped here. Even so, onboarding still affects observability through identity, consent, relationship state, visibility, and intention routing.

Implementation guardrails:

- Treat this as a support, watch, or translation layer unless another governed cluster or confidence threshold elevates it.

Source family	Where it appears now	When it matters most	What developers should do

			with it
Primary sources	Assessment + messaging + distress windows	Assessment baseline plus sudden silence or uncertainty windows	Wire these as allowed emitters and keep source hashes or event references.
Assessment map	15 mapped questions	Assessment completion, rescoring windows	Use as seed evidence, not instant activation.
Polarity map	0 mapped prompts	Polarity completion and polarity refresh	Use as chemistry-risk context or direct seed only when the map explicitly names it.
Onboarding context	0 direct fields	Account creation and profile edits	Use as routing, eligibility, context, and permission logic.
Cross-system emitters	See primary-source path	During lived behavior windows	Use for confirmation, suppression, or explanation depending on lane discipline.

CL_BOUNDARY_FOG - Boundary Fog

Boundary Fog sits in the needs/boundaries family and is currently classified as support/watch. Its primary evidence path is assessment + sanctuary + exsanonymous. The base timing rule is: assessment and reflective writing windows. The reason the engine preserves it is straightforward: Captures unclear limits, over-accommodation, hinting instead of asking, and confusion between love and self-erasure.

What it is	A governed or support-state reading of behavior, posture, or context. It should never be treated as a personality sentence; it is an operational signal that can strengthen, soften, suppress, or redirect readiness and forecast interpretation depending on corroboration.
Where we get it	Assessment + Sanctuary + ExsAnonymous
When we get it	Assessment and reflective writing windows
Why we keep it	Captures unclear limits, over-accommodation, hinting instead of asking, and confusion between love and self-erasure.
Direct onboarding hooks	None directly seeded at onboarding; onboarding may still create context, eligibility, or permission conditions.
Direct system emitters	No dedicated emitter line item in the high-level system table; rely on the primary-source path above and the rule engine.

Assessment seeding. The slow-lane assessment currently seeds this cluster from 15 scored questions, including Q2: What happens after your vulnerability is shrugged off, Q5: Not feeling emotionally met after opening up, Q9: Reaction to feeling deprioritized, Q10: Being leaned on without reciprocity, Q19: Curiosity imbalance and one-sided knowing, Q25: Unclear mutual feeling / who likes who, Q26: Unmet ask and follow-through reliability, Q27: Feeling awkward after emotional disclosure, Q30: Vulnerability landing badly or being mocked, Q34: Feeling used or conditionally wanted, Q39: How you respond to another person's overwhelm, Q40: The hidden reason you want partnership and 3 more seeded

assessment links. These do not activate the cluster by themselves; they provide blueprint-level evidence that later behavior either confirms, contradicts, or cools down.

Polarity seeding. The current polarity map does not explicitly name this cluster. That usually means polarity acts as background chemistry context rather than as a direct seed source.

Onboarding context. There is no direct onboarding trigger mapped here. Even so, onboarding still affects observability through identity, consent, relationship state, visibility, and intention routing.

Implementation guardrails:

- Treat this as a support, watch, or translation layer unless another governed cluster or confidence threshold elevates it.

Source family	Where it appears now	When it matters most	What developers should do with it
Primary sources	Assessment + Sanctuary + ExsAnonymous	Assessment and reflective writing windows	Wire these as allowed emitters and keep source hashes or event references.
Assessment map	15 mapped questions	Assessment completion, rescoring windows	Use as seed evidence, not instant activation.
Polarity map	0 mapped prompts	Polarity completion and polarity refresh	Use as chemistry-risk context or direct seed only when the map explicitly names it.
Onboarding context	0 direct fields	Account creation and profile edits	Use as routing, eligibility, context, and permission logic.
Cross-system emitters	See primary-source path	During lived behavior windows	Use for confirmation, suppression, or explanation depending on lane discipline.

CL_EROTIC_SAFETY_SPLIT - Erotic-Safety Split

Erotic-Safety Split sits in the chemistry-risk family and is currently classified as support/watch. Its primary evidence path is polarity + relationship mode + user reflections. The base timing rule is: assessment, chemistry report generation, and couple-state reviews. The reason the engine preserves it is straightforward: Tracks pairings where desire and safety run on separate rails, creating intensity without durable trust.

What it is	A governed or support-state reading of behavior, posture, or context. It should never be treated as a personality sentence; it is an operational signal that can strengthen, soften, suppress, or redirect readiness and forecast interpretation depending on corroboration.
Where we get it	Polarity + relationship mode + user reflections
When we get it	Assessment, chemistry report generation, and couple-state reviews
Why we keep it	Tracks pairings where desire and safety run on separate rails, creating intensity without durable trust.

Direct onboarding hooks	None directly seeded at onboarding; onboarding may still create context, eligibility, or permission conditions.
Direct system emitters	No dedicated emitter line item in the high-level system table; rely on the primary-source path above and the rule engine.

Assessment seeding. This cluster is not directly named in the current 60-question assessment map. That does not make it irrelevant; it simply means it is usually inferred from pair law, polarity, live telemetry, or companion-state systems rather than from a first-pass blueprint answer.

Polarity seeding. The polarity engine points at this cluster in 4 prompt slots, including Q6 (Emotional disclosure), Q14 (Emotional disclosure), Q22 (Emotional disclosure), Q30 (Emotional disclosure). In practice this means the cluster often rides on attraction-style misalignment, structure-flow mismatch, surrender safety strain, or chemistry that outruns stability.

Onboarding context. There is no direct onboarding trigger mapped here. Even so, onboarding still affects observability through identity, consent, relationship state, visibility, and intention routing.

Implementation guardrails:

- Treat this as a support, watch, or translation layer unless another governed cluster or confidence threshold elevates it.

Source family	Where it appears now	When it matters most	What developers should do with it
Primary sources	Polarity + relationship mode + user reflections	Assessment, chemistry report generation, and couple-state reviews	Wire these as allowed emitters and keep source hashes or event references.
Assessment map	0 mapped questions	Assessment completion, rescoring windows	Use as seed evidence, not instant activation.
Polarity map	4 mapped prompts	Polarity completion and polarity refresh	Use as chemistry-risk context or direct seed only when the map explicitly names it.
Onboarding context	0 direct fields	Account creation and profile edits	Use as routing, eligibility, context, and permission logic.
Cross-system emitters	See primary-source path	During lived behavior windows	Use for confirmation, suppression, or explanation depending on lane discipline.

CL_FLEX_RIGIDITY - Flex-Rigidity Collision

Flex-Rigidity Collision sits in the chemistry-risk family and is currently classified as support/watch. Its primary evidence path is polarity + conflict/repair telemetry. The base timing rule is: assessment and conflict windows. The reason the engine preserves it is straightforward: Captures stalemates created when one or both people cannot adapt their style under stress.

What it is	A governed or support-state reading of behavior, posture, or context. It should never be treated as a personality sentence; it is an operational signal that can strengthen, soften, suppress, or redirect readiness and forecast interpretation depending on corroboration.

Where we get it	Polarity + conflict/repair telemetry
When we get it	Assessment and conflict windows
Why we keep it	Captures stalemates created when one or both people cannot adapt their style under stress.
Direct onboarding hooks	None directly seeded at onboarding; onboarding may still create context, eligibility, or permission conditions.
Direct system emitters	No dedicated emitter line item in the high-level system table; rely on the primary-source path above and the rule engine.

Assessment seeding. This cluster is not directly named in the current 60-question assessment map. That does not make it irrelevant; it simply means it is usually inferred from pair law, polarity, live telemetry, or companion-state systems rather than from a first-pass blueprint answer.

Polarity seeding. The polarity engine points at this cluster in 20 prompt slots, including Q1 (Plan shifts / adaptation), Q2 (Silence and pause), Q3 (Choice under uncertainty), Q5 (Time pressure), Q8 (Being observed / exposure), Q9 (Plan shifts / adaptation), Q10 (Silence and pause), Q11 (Choice under uncertainty), Q13 (Time pressure), Q16 (Being observed / exposure) and 10 more polarity links. In practice this means the cluster often rides on attraction-style misalignment, structure-flow mismatch, surrender safety strain, or chemistry that outruns stability.

Onboarding context. There is no direct onboarding trigger mapped here. Even so, onboarding still affects observability through identity, consent, relationship state, visibility, and intention routing.

Implementation guardrails:

- Treat this as a support, watch, or translation layer unless another governed cluster or confidence threshold elevates it.

Source family	Where it appears now	When it matters most	What developers should do with it
Primary sources	Polarity + conflict/repair telemetry	Assessment and conflict windows	Wire these as allowed emitters and keep source hashes or event references.
Assessment map	0 mapped questions	Assessment completion, rescoring windows	Use as seed evidence, not instant activation.
Polarity map	20 mapped prompts	Polarity completion and polarity refresh	Use as chemistry-risk context or direct seed only when the map explicitly names it.
Onboarding context	0 direct fields	Account creation and profile edits	Use as routing, eligibility, context, and permission logic.
Cross-system emitters	See primary-source path	During lived behavior windows	Use for confirmation, suppression, or explanation depending on lane discipline.

CL_FORECAST_DRAG - Forecast Drag

Forecast Drag sits in the forecast meta-cluster family and is currently classified as support/watch. Its primary evidence path is pair engine + drift + repair + outcome telemetry. The base timing rule is: after

pair assembly and over time. The reason the engine preserves it is straightforward: Represents accumulated downward pressure on pair stability from multiple lighter burdens even when no single cluster is catastrophic.

What it is	A governed or support-state reading of behavior, posture, or context. It should never be treated as a personality sentence; it is an operational signal that can strengthen, soften, suppress, or redirect readiness and forecast interpretation depending on corroboration.
Where we get it	Pair engine + drift + repair + outcome telemetry
When we get it	After pair assembly and over time
Why we keep it	Represents accumulated downward pressure on pair stability from multiple lighter burdens even when no single cluster is catastrophic.
Direct onboarding hooks	None directly seeded at onboarding; onboarding may still create context, eligibility, or permission conditions.
Direct system emitters	Archive of Self / Alignment Compass; Relationship mode / couple tools

Assessment seeding. This cluster is not directly named in the current 60-question assessment map. That does not make it irrelevant; it simply means it is usually inferred from pair law, polarity, live telemetry, or companion-state systems rather than from a first-pass blueprint answer.

Polarity seeding. The current polarity map does not explicitly name this cluster. That usually means polarity acts as background chemistry context rather than as a direct seed source.

Onboarding context. There is no direct onboarding trigger mapped here. Even so, onboarding still affects observability through identity, consent, relationship state, visibility, and intention routing.

Implementation guardrails:

- Treat this as a support, watch, or translation layer unless another governed cluster or confidence threshold elevates it.

Source family	Where it appears now	When it matters most	What developers should do with it
Primary sources	Pair engine + drift + repair + outcome telemetry	After pair assembly and over time	Wire these as allowed emitters and keep source hashes or event references.
Assessment map	0 mapped questions	Assessment completion, rescoring windows	Use as seed evidence, not instant activation.
Polarity map	0 mapped prompts	Polarity completion and polarity refresh	Use as chemistry-risk context or direct seed only when the map explicitly names it.
Onboarding context	0 direct fields	Account creation and profile edits	Use as routing, eligibility, context, and permission logic.
Cross-system emitters	Archive of Self / Alignment Compass; Relationship mode / couple tools	During lived behavior windows	Use for confirmation, suppression, or explanation depending on lane discipline.

CL_INTEGRITY_STRAIN - Integrity Strain

Integrity Strain sits in the core burden family and is currently classified as support/watch. Its primary evidence path is assessment + shadow items + event ledger. The base timing rule is: assessment baseline and any integrity-relevant behavior events. The reason the engine preserves it is straightforward: Tracks mismatch between stated values and observed behavior, especially around honesty, accountability, and congruence.

What it is	A governed or support-state reading of behavior, posture, or context. It should never be treated as a personality sentence; it is an operational signal that can strengthen, soften, suppress, or redirect readiness and forecast interpretation depending on corroboration.
Where we get it	Assessment + shadow items + event ledger
When we get it	Assessment baseline and any integrity-relevant behavior events
Why we keep it	Tracks mismatch between stated values and observed behavior, especially around honesty, accountability, and congruence.
Direct onboarding hooks	None directly seeded at onboarding; onboarding may still create context, eligibility, or permission conditions.
Direct system emitters	Shadow questions (10)

Assessment seeding. The slow-lane assessment currently seeds this cluster from 3 scored questions, including Q28: Suspected dishonesty and trust rupture, Q34: Feeling used or conditionally wanted, Q101: Shadow response when your story stops matching. These do not activate the cluster by themselves; they provide blueprint-level evidence that later behavior either confirms, contradicts, or cools down.

Polarity seeding. The current polarity map does not explicitly name this cluster. That usually means polarity acts as background chemistry context rather than as a direct seed source.

Onboarding context. There is no direct onboarding trigger mapped here. Even so, onboarding still affects observability through identity, consent, relationship state, visibility, and intention routing.

Implementation guardrails:

- Treat this as a support, watch, or translation layer unless another governed cluster or confidence threshold elevates it.

Source family	Where it appears now	When it matters most	What developers should do with it
Primary sources	Assessment + shadow items + event ledger	Assessment baseline and any integrity-relevant behavior events	Wire these as allowed emitters and keep source hashes or event references.
Assessment map	3 mapped questions	Assessment completion, rescoring windows	Use as seed evidence, not instant activation.
Polarity map	0 mapped prompts	Polarity completion and polarity refresh	Use as chemistry-risk context or direct seed only when the

			map explicitly names it.
Onboarding context	0 direct fields	Account creation and profile edits	Use as routing, eligibility, context, and permission logic.
Cross-system emitters	Shadow questions (10)	During lived behavior windows	Use for confirmation, suppression, or explanation depending on lane discipline.

CL_INTENT_MISMATCH - Intent Mismatch

Intent Mismatch sits in the core burden family and is currently classified as support/watch. Its primary evidence path is onboarding profile + assessment + match packet. The base timing rule is: at onboarding, at assessment close, and at pair assembly. The reason the engine preserves it is straightforward: Prevents high chemistry from being mistaken for aligned direction when relationship intent is structurally mismatched.

What it is	A governed or support-state reading of behavior, posture, or context. It should never be treated as a personality sentence; it is an operational signal that can strengthen, soften, suppress, or redirect readiness and forecast interpretation depending on corroboration.
Where we get it	Onboarding profile + assessment + match packet
When we get it	At onboarding, at assessment close, and at pair assembly
Why we keep it	Prevents high chemistry from being mistaken for aligned direction when relationship intent is structurally mismatched.
Direct onboarding hooks	None directly seeded at onboarding; onboarding may still create context, eligibility, or permission conditions.
Direct system emitters	No dedicated emitter line item in the high-level system table; rely on the primary-source path above and the rule engine.

Assessment seeding. The slow-lane assessment currently seeds this cluster from 3 scored questions, including Q4: How you handle early red flags, Q20: Mismatch between chemistry and readiness, Q25: Unclear mutual feeling / who likes who. These do not activate the cluster by themselves; they provide blueprint-level evidence that later behavior either confirms, contradicts, or cools down.

Polarity seeding. The current polarity map does not explicitly name this cluster. That usually means polarity acts as background chemistry context rather than as a direct seed source.

Onboarding context. There is no direct onboarding trigger mapped here. Even so, onboarding still affects observability through identity, consent, relationship state, visibility, and intention routing.

Implementation guardrails:

- Treat this as a support, watch, or translation layer unless another governed cluster or confidence threshold elevates it.

Source family	Where it appears now	When it matters most	What developers should do with it

Primary sources	Onboarding profile + assessment + match packet	At onboarding, at assessment close, and at pair assembly	Wire these as allowed emitters and keep source hashes or event references.
Assessment map	3 mapped questions	Assessment completion, rescoring windows	Use as seed evidence, not instant activation.
Polarity map	0 mapped prompts	Polarity completion and polarity refresh	Use as chemistry-risk context or direct seed only when the map explicitly names it.
Onboarding context	0 direct fields	Account creation and profile edits	Use as routing, eligibility, context, and permission logic.
Cross-system emitters	See primary-source path	During lived behavior windows	Use for confirmation, suppression, or explanation depending on lane discipline.

CL_PACE_MISMATCH - Pace Mismatch

Pace Mismatch sits in the core burden family and is currently classified as support/watch. Its primary evidence path is assessment + matching + drift. The base timing rule is: assessment and pair assembly; rechecked during pace shocks. The reason the engine preserves it is straightforward: Captures misaligned speed for intimacy, commitment, and relationship progression.

What it is	**A governed or support-state reading of behavior, posture, or context. It should never be treated as a personality sentence; it is an operational signal that can strengthen, soften, suppress, or redirect readiness and forecast interpretation depending on corroboration.**
Where we get it	Assessment + matching + drift
When we get it	Assessment and pair assembly; rechecked during pace shocks
Why we keep it	Captures misaligned speed for intimacy, commitment, and relationship progression.
Direct onboarding hooks	None directly seeded at onboarding; onboarding may still create context, eligibility, or permission conditions.
Direct system emitters	No dedicated emitter line item in the high-level system table; rely on the primary-source path above and the rule engine.

Assessment seeding. The slow-lane assessment currently seeds this cluster from 4 scored questions, including Q20: Mismatch between chemistry and readiness, Q23: Energy mismatch and overinvestment threshold, Q25: Unclear mutual feeling / who likes who, Q33: How you leave when commitment is wrong. These do not activate the cluster by themselves; they provide blueprint-level evidence that later behavior either confirms, contradicts, or cools down.

Polarity seeding. The current polarity map does not explicitly name this cluster. That usually means polarity acts as background chemistry context rather than as a direct seed source.

Onboarding context. There is no direct onboarding trigger mapped here. Even so, onboarding still affects observability through identity, consent, relationship state, visibility, and intention routing.

Implementation guardrails:

- Treat this as a support, watch, or translation layer unless another governed cluster or confidence threshold elevates it.

Source family	Where it appears now	When it matters most	What developers should do with it
Primary sources	Assessment + matching + drift	Assessment and pair assembly; rechecked during pace shocks	Wire these as allowed emitters and keep source hashes or event references.
Assessment map	4 mapped questions	Assessment completion, rescoring windows	Use as seed evidence, not instant activation.
Polarity map	0 mapped prompts	Polarity completion and polarity refresh	Use as chemistry-risk context or direct seed only when the map explicitly names it.
Onboarding context	0 direct fields	Account creation and profile edits	Use as routing, eligibility, context, and permission logic.
Cross-system emitters	See primary-source path	During lived behavior windows	Use for confirmation, suppression, or explanation depending on lane discipline.

CL_RECIPROCITY_IMBALANCE - Reciprocity Imbalance

Reciprocity Imbalance sits in the core burden family and is currently classified as support/watch. Its primary evidence path is assessment + messaging + couple mode. The base timing rule is: assessment seeds, then recalculated through interaction data. The reason the engine preserves it is straightforward: Captures one-sided effort, uneven pursuit, or chronic emotional labor asymmetry.

What it is	A governed or support-state reading of behavior, posture, or context. It should never be treated as a personality sentence; it is an operational signal that can strengthen, soften, suppress, or redirect readiness and forecast interpretation depending on corroboration.
Where we get it	Assessment + messaging + couple mode
When we get it	Assessment seeds, then recalculated through interaction data
Why we keep it	Captures one-sided effort, uneven pursuit, or chronic emotional labor asymmetry.
Direct onboarding hooks	None directly seeded at onboarding; onboarding may still create context, eligibility, or permission conditions.
Direct system emitters	No dedicated emitter line item in the high-level system table; rely on the primary-source path above and the rule engine.

Assessment seeding. The slow-lane assessment currently seeds this cluster from 9 scored questions, including Q9: Reaction to feeling deprioritized, Q10: Being leaned on without reciprocity, Q14: Matching reduced effort and mixed signals, Q19: Curiosity imbalance and one-sided knowing, Q23: Energy mismatch and overinvestment threshold, Q24: Effort that goes unacknowledged, Q26: Unmet ask and follow-through reliability, Q32: One-sided initiation and outreach load, Q34: Feeling used or

conditionally wanted. These do not activate the cluster by themselves; they provide blueprint-level evidence that later behavior either confirms, contradicts, or cools down.

Polarity seeding. The current polarity map does not explicitly name this cluster. That usually means polarity acts as background chemistry context rather than as a direct seed source.

Onboarding context. There is no direct onboarding trigger mapped here. Even so, onboarding still affects observability through identity, consent, relationship state, visibility, and intention routing.

Implementation guardrails:

- Treat this as a support, watch, or translation layer unless another governed cluster or confidence threshold elevates it.

Source family	Where it appears now	When it matters most	What developers should do with it
Primary sources	Assessment + messaging + couple mode	Assessment seeds, then recalculated through interaction data	Wire these as allowed emitters and keep source hashes or event references.
Assessment map	9 mapped questions	Assessment completion, rescoring windows	Use as seed evidence, not instant activation.
Polarity map	0 mapped prompts	Polarity completion and polarity refresh	Use as chemistry-risk context or direct seed only when the map explicitly names it.
Onboarding context	0 direct fields	Account creation and profile edits	Use as routing, eligibility, context, and permission logic.
Cross-system emitters	See primary-source path	During lived behavior windows	Use for confirmation, suppression, or explanation depending on lane discipline.

CL_STRUCTURE_COLLISION - Structure Collision

Structure Collision sits in the chemistry-risk family and is currently classified as support/watch. Its primary evidence path is polarity + onboarding intent + couple mode. The base timing rule is: at polarity scoring and pair assembly. The reason the engine preserves it is straightforward: Flags mismatch in who carries direction, decision-making, and relational shape.

What it is	A governed or support-state reading of behavior, posture, or context. It should never be treated as a personality sentence; it is an operational signal that can strengthen, soften, suppress, or redirect readiness and forecast interpretation depending on corroboration.
Where we get it	Polarity + onboarding intent + couple mode
When we get it	At polarity scoring and pair assembly
Why we keep it	Flags mismatch in who carries direction, decision-making, and relational shape.
Direct onboarding hooks	None directly seeded at onboarding; onboarding may still create context, eligibility, or permission conditions.

Direct system emitters	No dedicated emitter line item in the high-level system table; rely on the primary-source path above and the rule engine.

Assessment seeding. This cluster is not directly named in the current 60-question assessment map. That does not make it irrelevant; it simply means it is usually inferred from pair law, polarity, live telemetry, or companion-state systems rather than from a first-pass blueprint answer.

Polarity seeding. The polarity engine points at this cluster in 28 prompt slots, including Q1 (Plan shifts / adaptation), Q2 (Silence and pause), Q3 (Choice under uncertainty), Q4 (Early friction / unresolved tension), Q5 (Time pressure), Q7 (Conflicting preferences), Q8 (Being observed / exposure), Q9 (Plan shifts / adaptation), Q10 (Silence and pause), Q11 (Choice under uncertainty) and 18 more polarity links. In practice this means the cluster often rides on attraction-style misalignment, structure-flow mismatch, surrender safety strain, or chemistry that outruns stability.

Onboarding context. There is no direct onboarding trigger mapped here. Even so, onboarding still affects observability through identity, consent, relationship state, visibility, and intention routing.

Implementation guardrails:

- Treat this as a support, watch, or translation layer unless another governed cluster or confidence threshold elevates it.

Source family	Where it appears now	When it matters most	What developers should do with it
Primary sources	Polarity + onboarding intent + couple mode	At polarity scoring and pair assembly	Wire these as allowed emitters and keep source hashes or event references.
Assessment map	0 mapped questions	Assessment completion, rescoring windows	Use as seed evidence, not instant activation.
Polarity map	28 mapped prompts	Polarity completion and polarity refresh	Use as chemistry-risk context or direct seed only when the map explicitly names it.
Onboarding context	0 direct fields	Account creation and profile edits	Use as routing, eligibility, context, and permission logic.
Cross-system emitters	See primary-source path	During lived behavior windows	Use for confirmation, suppression, or explanation depending on lane discipline.

CL_SURRENDER_SAFETY_DEFICIT - Surrender Safety Deficit

Surrender Safety Deficit sits in the chemistry-risk family and is currently classified as support/watch. Its primary evidence path is assessment + polarity + sanctuary. The base timing rule is: assessment and intimacy escalation windows. The reason the engine preserves it is straightforward: Flags when closeness does not feel safe enough to relax into, even if attraction is present.

What it is	A governed or support-state reading of behavior, posture, or context. It should never be treated as a personality sentence; it is an operational signal that can strengthen, soften, suppress, or redirect readiness and forecast interpretation depending on corroboration.

Where we get it	Assessment + polarity + Sanctuary
When we get it	Assessment and intimacy escalation windows
Why we keep it	Flags when closeness does not feel safe enough to relax into, even if attraction is present.
Direct onboarding hooks	None directly seeded at onboarding; onboarding may still create context, eligibility, or permission conditions.
Direct system emitters	No dedicated emitter line item in the high-level system table; rely on the primary-source path above and the rule engine.

Assessment seeding. The slow-lane assessment currently seeds this cluster from 13 scored questions, including Q3: How you show early interest safely, Q13: How interest becomes visible over time, Q15: Receiving unusually healthy interest, Q17: Partner goes quiet under emotion, Q22: Repeated silence / ambiguity test, Q30: Vulnerability landing badly or being mocked, Q38: What happens when someone sees you deeply, Q41: Who you are attracted to and why, Q43: Primary fear once love becomes real, Q45: What actually helps you reopen, Q47: How healthy care feels in your body, Q48: What destabilizes closeness once it becomes real and 1 more seeded assessment links. These do not activate the cluster by themselves; they provide blueprint-level evidence that later behavior either confirms, contradicts, or cools down.

Polarity seeding. The polarity engine points at this cluster in 20 prompt slots, including Q2 (Silence and pause), Q4 (Early friction / unresolved tension), Q6 (Emotional disclosure), Q7 (Conflicting preferences), Q8 (Being observed / exposure), Q10 (Silence and pause), Q12 (Early friction / unresolved tension), Q14 (Emotional disclosure), Q15 (Conflicting preferences), Q16 (Being observed / exposure) and 10 more polarity links. In practice this means the cluster often rides on attraction-style misalignment, structure-flow mismatch, surrender safety strain, or chemistry that outruns stability.

Onboarding context. There is no direct onboarding trigger mapped here. Even so, onboarding still affects observability through identity, consent, relationship state, visibility, and intention routing.

Implementation guardrails:

- Treat this as a support, watch, or translation layer unless another governed cluster or confidence threshold elevates it.

Source family	Where it appears now	When it matters most	What developers should do with it
Primary sources	Assessment + polarity + Sanctuary	Assessment and intimacy escalation windows	Wire these as allowed emitters and keep source hashes or event references.
Assessment map	13 mapped questions	Assessment completion, rescoring windows	Use as seed evidence, not instant activation.
Polarity map	20 mapped prompts	Polarity completion and polarity refresh	Use as chemistry-risk context or direct seed only when the map explicitly names it.
Onboarding context	0 direct fields	Account creation and profile edits	Use as routing, eligibility, context, and permission logic.

Cross-system emitters	See primary-source path	During lived behavior windows	Use for confirmation, suppression, or explanation depending on lane discipline.

CL_TENSION_BONDING - Tension Bonding

Tension Bonding sits in the chemistry-risk family and is currently classified as support/watch. Its primary evidence path is polarity + messaging conflict cadence. The base timing rule is: during attraction ramp and conflict repetition. The reason the engine preserves it is straightforward: Captures chemistry that is being powered by friction, uncertainty, chaos, or emotional whiplash.

What it is	A governed or support-state reading of behavior, posture, or context. It should never be treated as a personality sentence; it is an operational signal that can strengthen, soften, suppress, or redirect readiness and forecast interpretation depending on corroboration.
Where we get it	Polarity + messaging conflict cadence
When we get it	During attraction ramp and conflict repetition
Why we keep it	Captures chemistry that is being powered by friction, uncertainty, chaos, or emotional whiplash.
Direct onboarding hooks	None directly seeded at onboarding; onboarding may still create context, eligibility, or permission conditions.
Direct system emitters	No dedicated emitter line item in the high-level system table; rely on the primary-source path above and the rule engine.

Assessment seeding. This cluster is not directly named in the current 60-question assessment map. That does not make it irrelevant; it simply means it is usually inferred from pair law, polarity, live telemetry, or companion-state systems rather than from a first-pass blueprint answer.

Polarity seeding. The polarity engine points at this cluster in 4 prompt slots, including Q4 (Early friction / unresolved tension), Q12 (Early friction / unresolved tension), Q20 (Early friction / unresolved tension), Q28 (Early friction / unresolved tension). In practice this means the cluster often rides on attraction-style misalignment, structure-flow mismatch, surrender safety strain, or chemistry that outruns stability.

Onboarding context. There is no direct onboarding trigger mapped here. Even so, onboarding still affects observability through identity, consent, relationship state, visibility, and intention routing.

Implementation guardrails:

- Treat this as a support, watch, or translation layer unless another governed cluster or confidence threshold elevates it.

Source family	Where it appears now	When it matters most	What developers should do with it
Primary sources	Polarity + messaging conflict cadence	During attraction ramp and conflict repetition	Wire these as allowed emitters and keep source hashes or event references.
Assessment map	0 mapped questions	Assessment completion,	Use as seed evidence, not

		rescoring windows	instant activation.
Polarity map	4 mapped prompts	Polarity completion and polarity refresh	Use as chemistry-risk context or direct seed only when the map explicitly names it.
Onboarding context	0 direct fields	Account creation and profile edits	Use as routing, eligibility, context, and permission logic.
Cross-system emitters	See primary-source path	During lived behavior windows	Use for confirmation, suppression, or explanation depending on lane discipline.

CL_VALIDATION_HUNGER - Validation Dependency

Validation Dependency sits in the chemistry-risk family and is currently classified as support/watch. Its primary evidence path is assessment + polarity + messaging. The base timing rule is: assessment, attraction windows, approval-seeking sequences. The reason the engine preserves it is straightforward: Captures chemistry inflated by being chosen, mirrored, or reassured rather than by actual relational fit.

What it is	A governed or support-state reading of behavior, posture, or context. It should never be treated as a personality sentence; it is an operational signal that can strengthen, soften, suppress, or redirect readiness and forecast interpretation depending on corroboration.
Where we get it	Assessment + polarity + messaging
When we get it	Assessment, attraction windows, approval-seeking sequences
Why we keep it	Captures chemistry inflated by being chosen, mirrored, or reassured rather than by actual relational fit.
Direct onboarding hooks	None directly seeded at onboarding; onboarding may still create context, eligibility, or permission conditions.
Direct system emitters	No dedicated emitter line item in the high-level system table; rely on the primary-source path above and the rule engine.

Assessment seeding. The slow-lane assessment currently seeds this cluster from 6 scored questions, including Q15: Receiving unusually healthy interest, Q36: Being left for someone else / replaced, Q38: What happens when someone sees you deeply, Q40: The hidden reason you want partnership, Q41: Who you are attracted to and why, Q47: How healthy care feels in your body. These do not activate the cluster by themselves; they provide blueprint-level evidence that later behavior either confirms, contradicts, or cools down.

Polarity seeding. The polarity engine points at this cluster in 4 prompt slots, including Q6 (Emotional disclosure), Q14 (Emotional disclosure), Q22 (Emotional disclosure), Q30 (Emotional disclosure). In practice this means the cluster often rides on attraction-style misalignment, structure-flow mismatch, surrender safety strain, or chemistry that outruns stability.

Onboarding context. There is no direct onboarding trigger mapped here. Even so, onboarding still affects observability through identity, consent, relationship state, visibility, and intention routing.

Implementation guardrails:

- Treat this as a support, watch, or translation layer unless another governed cluster or confidence threshold elevates it.

Source family	Where it appears now	When it matters most	What developers should do with it
Primary sources	Assessment + polarity + messaging	Assessment, attraction windows, approval-seeking sequences	Wire these as allowed emitters and keep source hashes or event references.
Assessment map	6 mapped questions	Assessment completion, rescoring windows	Use as seed evidence, not instant activation.
Polarity map	4 mapped prompts	Polarity completion and polarity refresh	Use as chemistry-risk context or direct seed only when the map explicitly names it.
Onboarding context	0 direct fields	Account creation and profile edits	Use as routing, eligibility, context, and permission logic.
Cross-system emitters	See primary-source path	During lived behavior windows	Use for confirmation, suppression, or explanation depending on lane discipline.

CL_ANGER_ESCALATION - Anger Escalation

Anger Escalation sits in the companion state family and is currently classified as companion state. Its primary evidence path is anger hex + event labeler + repair telemetry. The base timing rule is: real-time to daily in high-conflict windows. The reason the engine preserves it is straightforward: Separates irritation, protest, shutdown, and retaliatory expression so the system knows whether safety and repair are improving or worsening.

What it is	A governed or support-state reading of behavior, posture, or context. It should never be treated as a personality sentence; it is an operational signal that can strengthen, soften, suppress, or redirect readiness and forecast interpretation depending on corroboration.
Where we get it	Anger Hex + event labeler + repair telemetry
When we get it	Real-time to daily in high-conflict windows
Why we keep it	Separates irritation, protest, shutdown, and retaliatory expression so the system knows whether safety and repair are improving or worsening.
Direct onboarding hooks	None directly seeded at onboarding; onboarding may still create context, eligibility, or permission conditions.
Direct system emitters	No dedicated emitter line item in the high-level system table; rely on the primary-source path above and the rule engine.

Assessment seeding. The slow-lane assessment currently seeds this cluster from 1 scored questions, including Q108: Shadow move in active conflict escalation. These do not activate the cluster by themselves; they provide blueprint-level evidence that later behavior either confirms, contradicts, or cools down.

Polarity seeding. The current polarity map does not explicitly name this cluster. That usually means polarity acts as background chemistry context rather than as a direct seed source.

Onboarding context. There is no direct onboarding trigger mapped here. Even so, onboarding still affects observability through identity, consent, relationship state, visibility, and intention routing.

Implementation guardrails:

- Do not let this companion-state signal rewrite baseline archetype, baseline shadow, or pair compatibility law.

Source family	Where it appears now	When it matters most	What developers should do with it
Primary sources	Anger Hex + event labeler + repair telemetry	Real-time to daily in high-conflict windows	Wire these as allowed emitters and keep source hashes or event references.
Assessment map	1 mapped questions	Assessment completion, rescoring windows	Use as seed evidence, not instant activation.
Polarity map	0 mapped prompts	Polarity completion and polarity refresh	Use as chemistry-risk context or direct seed only when the map explicitly names it.
Onboarding context	0 direct fields	Account creation and profile edits	Use as routing, eligibility, context, and permission logic.
Cross-system emitters	See primary-source path	During lived behavior windows	Use for confirmation, suppression, or explanation depending on lane discipline.

CL_EX_TETHER - Ex Tether / Narrative Carryover

Ex Tether / Narrative Carryover sits in the companion state family and is currently classified as companion state. Its primary evidence path is exsanonymous + sanctuary archive + reflections. The base timing rule is: whenever ex-processing activity is active or recurring. The reason the engine preserves it is straightforward: Tracks whether old relational loops are still shaping present trust, comparison, or emotional availability.

What it is	A governed or support-state reading of behavior, posture, or context. It should never be treated as a personality sentence; it is an operational signal that can strengthen, soften, suppress, or redirect readiness and forecast interpretation depending on corroboration.
Where we get it	ExsAnonymous + Sanctuary archive + reflections
When we get it	Whenever ex-processing activity is active or recurring
Why we keep it	Tracks whether old relational loops are still shaping present trust, comparison, or emotional availability.
Direct onboarding hooks	None directly seeded at onboarding; onboarding may still create context, eligibility, or permission conditions.
Direct system emitters	No dedicated emitter line item in the high-level system table; rely on the primary-source path above and the rule engine.

Assessment seeding. The slow-lane assessment currently seeds this cluster from 1 scored questions, including Q49: Which old wound still governs love. These do not activate the cluster by themselves; they provide blueprint-level evidence that later behavior either confirms, contradicts, or cools down.

Polarity seeding. The current polarity map does not explicitly name this cluster. That usually means polarity acts as background chemistry context rather than as a direct seed source.

Onboarding context. There is no direct onboarding trigger mapped here. Even so, onboarding still affects observability through identity, consent, relationship state, visibility, and intention routing.

Implementation guardrails:

- Do not let this companion-state signal rewrite baseline archetype, baseline shadow, or pair compatibility law.

Source family	Where it appears now	When it matters most	What developers should do with it
Primary sources	ExsAnonymous + Sanctuary archive + reflections	Whenever ex-processing activity is active or recurring	Wire these as allowed emitters and keep source hashes or event references.
Assessment map	1 mapped questions	Assessment completion, rescoring windows	Use as seed evidence, not instant activation.
Polarity map	0 mapped prompts	Polarity completion and polarity refresh	Use as chemistry-risk context or direct seed only when the map explicitly names it.
Onboarding context	0 direct fields	Account creation and profile edits	Use as routing, eligibility, context, and permission logic.
Cross-system emitters	See primary-source path	During lived behavior windows	Use for confirmation, suppression, or explanation depending on lane discipline.

CL_GAME_COMPENSATION - Game Compensation / Bond Distortion

Game Compensation / Bond Distortion sits in the companion state family and is currently classified as companion state. Its primary evidence path is rs game ecosystem + couple sessions + telemetry. The base timing rule is: during game sessions and post-session analysis. The reason the engine preserves it is straightforward: Captures when games are healthy bonding versus a substitute for direct intimacy, clarity, or repair.

What it is	A governed or support-state reading of behavior, posture, or context. It should never be treated as a personality sentence; it is an operational signal that can strengthen, soften, suppress, or redirect readiness and forecast interpretation depending on corroboration.
Where we get it	RS Game Ecosystem + couple sessions + telemetry
When we get it	During game sessions and post-session analysis
Why we keep it	Captures when games are healthy bonding versus a substitute for direct intimacy, clarity, or repair.

Direct onboarding hooks	None directly seeded at onboarding; onboarding may still create context, eligibility, or permission conditions.
Direct system emitters	No dedicated emitter line item in the high-level system table; rely on the primary-source path above and the rule engine.

Assessment seeding. This cluster is not directly named in the current 60-question assessment map. That does not make it irrelevant; it simply means it is usually inferred from pair law, polarity, live telemetry, or companion-state systems rather than from a first-pass blueprint answer.

Polarity seeding. The current polarity map does not explicitly name this cluster. That usually means polarity acts as background chemistry context rather than as a direct seed source.

Onboarding context. There is no direct onboarding trigger mapped here. Even so, onboarding still affects observability through identity, consent, relationship state, visibility, and intention routing.

Implementation guardrails:

- Do not let this companion-state signal rewrite baseline archetype, baseline shadow, or pair compatibility law.

Source family	Where it appears now	When it matters most	What developers should do with it
Primary sources	RS Game Ecosystem + couple sessions + telemetry	During game sessions and post-session analysis	Wire these as allowed emitters and keep source hashes or event references.
Assessment map	0 mapped questions	Assessment completion, rescoring windows	Use as seed evidence, not instant activation.
Polarity map	0 mapped prompts	Polarity completion and polarity refresh	Use as chemistry-risk context or direct seed only when the map explicitly names it.
Onboarding context	0 direct fields	Account creation and profile edits	Use as routing, eligibility, context, and permission logic.
Cross-system emitters	See primary-source path	During lived behavior windows	Use for confirmation, suppression, or explanation depending on lane discipline.

CL_GRIEF_LOAD - Grief / Depletion Load

Grief / Depletion Load sits in the companion state family and is currently classified as companion state. Its primary evidence path is e_grief + sanctuary + last light + user check-ins. The base timing rule is: daily to milestone windows; always decays with time and evidence. The reason the engine preserves it is straightforward: Prevents depletion, loss, and spiritual drag from being mistaken for identity truth or permanent compatibility failure.

What it is	A governed or support-state reading of behavior, posture, or context. It should never be treated as a personality sentence; it is an operational signal that can strengthen, soften, suppress, or redirect readiness and forecast interpretation depending on corroboration.
Where we get it	E_GRIEF + Sanctuary + Last Light + user check-ins

When we get it	Daily to milestone windows; always decays with time and evidence
Why we keep it	Prevents depletion, loss, and spiritual drag from being mistaken for identity truth or permanent compatibility failure.
Direct onboarding hooks	None directly seeded at onboarding; onboarding may still create context, eligibility, or permission conditions.
Direct system emitters	No dedicated emitter line item in the high-level system table; rely on the primary-source path above and the rule engine.

Assessment seeding. The slow-lane assessment currently seeds this cluster from 2 scored questions, including Q46: How you metabolize heartbreak, Q50: What scares you about healthy love. These do not activate the cluster by themselves; they provide blueprint-level evidence that later behavior either confirms, contradicts, or cools down.

Polarity seeding. The current polarity map does not explicitly name this cluster. That usually means polarity acts as background chemistry context rather than as a direct seed source.

Onboarding context. There is no direct onboarding trigger mapped here. Even so, onboarding still affects observability through identity, consent, relationship state, visibility, and intention routing.

Implementation guardrails:

- Do not let this companion-state signal rewrite baseline archetype, baseline shadow, or pair compatibility law.

Source family	Where it appears now	When it matters most	What developers should do with it
Primary sources	E_GRIEF + Sanctuary + Last Light + user check-ins	Daily to milestone windows; always decays with time and evidence	Wire these as allowed emitters and keep source hashes or event references.
Assessment map	2 mapped questions	Assessment completion, rescoring windows	Use as seed evidence, not instant activation.
Polarity map	0 mapped prompts	Polarity completion and polarity refresh	Use as chemistry-risk context or direct seed only when the map explicitly names it.
Onboarding context	0 direct fields	Account creation and profile edits	Use as routing, eligibility, context, and permission logic.
Cross-system emitters	See primary-source path	During lived behavior windows	Use for confirmation, suppression, or explanation depending on lane discipline.

9. Developer write-back and storage guardrails

- Store source references. Every cluster event should preserve the answer_id, question number, event_id, rollup window, or companion-state source that nominated it.
- Keep seed evidence distinct from persisted state. A cluster may appear in memory before it qualifies for write-back.

- Do not let pair-state traces leak pair-private detail onto public individual surfaces. Mirrored user traces should store softened downstream meaning, not raw pair secrets.
- Version every threshold, every rule-pack change, and every write-back schema change. Clusters without migration history eventually become folklore with IDs.
- Severity, confidence, visibility, decay, and clear conditions must stay outside UI code. UI should consume allowed surfaces, not invent them.

10. Closing note

This atlas is meant to be the long-form answer to the question, 'what exactly are our clusters and where do they come from?' The answer is now explicit: from onboarding context, from slow-lane scored questions, from polarity chemistry logic, from pair-state rollups, from messaging behavior, from repair attempts, from Sanctuary usage, from shadow pressure, from grief and anger context, from ex-history, and from game telemetry - but only through governed lanes, thresholds, and write-back rules.

That is the difference between a system and a mood board.