Phase 1 Schema & Data Contract

Build-governing schema for MVP launch implementation

Purpose	Define the exact Phase 1 database objects, import tables, storage contracts, keys, and payloads so engineering can build the launch system without guessing.
Use with	1) Genesis_OS_Phase_1_MVP_Build_Spec_v1 2) genesis_os_final_quantum_scoring_workbook_v10_polarity.xlsx
Authoritative launch stance	Workbook-governed scoring law first; backend-only score computation; no simplified MVP formula; no client-side recreation.
Not build sources	Simulation workbook, legacy merged blueprints, and any older weighted compatibility formulas are non-authoritative for Phase 1 implementation.

Why this document exists

- Earlier blueprint generations mixed legacy logic with the corrected workbook-governed pair engine.
- For launch, the scoring system had to be corrected and locked so engineering builds the right system the first time.
- This schema contract exists to remove implementation ambiguity: what must be imported, what must be stored, what is computed, what is returned, and what is explicitly out of Phase 1.

Version 1.0 | Phase 1 launch schema | Postgres-friendly types shown; equivalent backend types are acceptable if behavior matches

1. Executive Summary

This document is the schema companion to the Phase 1 MVP build spec. It does not replace the workbook. The workbook remains the score law; this document tells engineering how to store the data, how to import the workbook artifacts, what launch tables must exist, how pair computations should be persisted, and how to keep the build deterministic and auditable.

Phase 1 is intentionally narrower than the full Genesis OS vision. The UI may be simple at launch, but the engine underneath must still be the real engine. That means AnswerPairQuantum lookup, question weighting, gate logic, rule penalties, and backend-only score computation all ship in Phase 1. Testing later means parity, validation, and calibration - not postponing or simplifying the actual score law.

Phase 1 handoff package

Artifact	Role in build	Rule
Phase 1 MVP Build Spec	Primary implementation scope	Defines what to build now, what to defer, and the required build order.
Phase 1 Schema & Data Contract (this doc)	Primary storage contract	Defines tables, fields, relationships, import metadata, keys, indexes, and payload shapes.
Quantum workbook	Primary score source of truth	Defines scoring law, imported row fields, pair interactions, weights, matrices, overlays, and tracked metrics.

Launch correction statement

- The build package changed because it needed to be correct for launch.
- The corrected launch package removes ambiguity between old weighted formulas and the workbook-governed pair engine.
- Developers are implementing a fixed scoring system, not inventing one and not deciding later whether the core logic works.

2. Source-of-Truth Order for Phase 1

Priority	Artifact	What it governs
1	genesis_os_final_quantum_scoring_workbook_v10_polarity.xlsx	Canonical score law, imported fields, pair quantum rows, matrices, overlays, and polarity package.
2	Phase 1 MVP Build Spec	Launch scope, implementation sequence, exclusions, and MVP acceptance criteria.
3	This schema contract	Storage design, table contract, import metadata, payload shapes, indexes, and persistence rules.
4	Everything else	Useful later for enrichment or future phases, but not authoritative for Phase 1 launch implementation.

- If a field, weight, or formula in code conflicts with the workbook, the workbook wins.
- If a legacy Word or PDF fragment conflicts with the Phase 1 build spec or this schema contract, the Phase 1 package wins.
- The simulation workbook may be used for later validation, but it is not a build source for Phase 1 implementation.

3. Phase 1 Scope Boundary

Build now	Defer
Onboarding identity and consent	Full Sanctuary module family
60-question assessment persistence	Full Human Emotion Genome feature store
Workbook imports and versioning	Advanced drift, shadow activation, and couple tools
AnswerPairQuantum scoring service	Quantum Mirror full premium memory stack
Match compute + tier + access gating	Autonomous model tuning or score reweighting
Minimal Knock + conversation shell	Broader social, badge, or economy surfaces
Event ledger + feedback capture foundation	Deep shared-relationship mode and Symphonic Match

- Phase 1 may return score, norms, tier, and safe translated drivers.
- Phase 1 must not simplify the score engine into a loose weighted approximation.
- Phase 1 may create reserved columns or empty tables for later modules, but those modules are not required for launch unless explicitly marked below.

4. Non-Negotiable Schema Rules

- Use UUID primary keys for platform records and stable text keys for workbook row identifiers (for example answer_id and pair_key).
- Use snake_case field names consistently across database tables, service payloads, and event objects.
- Use timestamptz / UTC timestamps for all persisted system events and score records.
- Use NUMERIC(6,4) for 0-1 norms, fits, burdens, penalties, and confidence values unless a stricter type is required.
- Use JSONB for diagnostics, payloads, reason codes, packet blocks, and visibility masks where structured variable-length content is needed.
- Imported workbook rows are immutable per import job. Never hand-edit imported score rows in production.
- No score math belongs in the client. Frontend surfaces render outputs only.
- Both answer pairs and user pairs must be order-insensitive. Always sort before creating pair keys.
- Do not destructively overwrite historical score records. Store recomputes as new rows or with versioned snapshots.
- Do not use the simulation workbook, old blueprints, or older formulas to fill missing schema decisions.

5. Workbook Import Contract

Every workbook import must be stored as a versioned import job. Sheet rows include the header row. Data rows exclude the header row. No code should hard-code disputed row counts when the imported sheet registry can store the truth.

Sheet	Sheet rows	Data rows / usable dims	Phase 1 status	Why it exists
AnswerBank	311	310	Required	Atomic answer vectors and metadata used for all assessment joins and aggregates.
AnswerPairQuantum	48,206	48,205	Required	Canonical pair lookup artifact; the core engine table.
ArchetypeMatrix	11	10x10	Required	Archetype fit lookup values.
ShadowMatrix	7	6x6	Required	Shadow interaction stability/risk lookup values.
ZodiacLookup	577	576	Required	7% cosmic overlay lookup.
NumerologyLookup	67	66	Required	5% cosmic overlay lookup.
MatchHarness	84	83 data rows + examples	Required	Parity and QA reference for deterministic score assembly.
Polarity_Axes_v10	11	10	Required import	Official chemistry dimensions and weights.
Polarity_RoleSeeds_v10	6	5	Required import	Scoring/content anchors for chemistry package auditability.
Polarity_QuestionMap_v10	33	32	Required import	Question-to-axis mapping for chemistry package traceability.
Polarity_AnswerBank_v10	161	160	Required import	Per-answer chemistry vectors and seed values.
Polarity_PairEngine_v10	22	21 data rows + formulas	Required import	Reference implementation of chemistry fit formulas.

Sheet	Sheet rows	Data rows / usable dims	Phase 1 status	Why it exists
FeedbackData	18	17	Required launch capture	Outcome telemetry schema for post-launch validation and calibration.
README / README_v10 / Rewrite_Rules_v10	Context sheets	Reference	Reference only	Useful for human interpretation, not direct production score math.

Required import metadata

Field	Type	Req	Notes
import_job_id	uuid	Yes	Primary import batch id for all workbook loads.
workbook_version	text	Yes	Human-readable workbook release tag or file version.
score_contract_version	text	Yes	Version of the launch scoring contract used with this import.
source_file_name	text	Yes	Original workbook file name persisted for auditability.
source_checksum	text	Yes	Checksum of the source workbook file; required for parity and audit.
sheet_name	text	Yes	Workbook sheet name as imported.
sheet_rows	int	Yes	Physical sheet row count including header row.
data_rows	int	Yes	Importable record count excluding header row.
imported_at	timestamptz	Yes	UTC time import completed.
imported_by	text	Yes	Actor, service, or environment that executed the import.
active_flag	bool	Yes	Marks the currently active import set for launch runtime.

6. Core Phase 1 Table Map

Table	Purpose	Phase 1
users	Launch identity, state, demographic keys, and basic profile routing.	Required
consent_state	Opt-ins, visibility, and policy scope required before AI or behavioral enrichment.	Required
assessment_sessions	Tracks one assessment submission cycle and version tags.	Required
user_answers	Stores selected answer rows for the 60-question assessment.	Required
workbook_import_jobs	Versioned import batches for workbook data.	Required
workbook_sheet_registry	Per-sheet row counts and active-sheet audit data.	Required
answer_bank_rows	Imported AnswerBank rows from workbook.	Required
answer_pair_quantum_rows	Imported AnswerPairQuantum rows from workbook.	Required
archetype_matrix_rows	Imported archetype fit lookup rows.	Required
shadow_matrix_rows	Imported shadow fit lookup rows.	Required
zodiac_lookup_rows	Imported zodiac overlay rows.	Required
numerology_lookup_rows	Imported numerology overlay rows.	Required
polarity_* import tables	Chemistry package reference imports.	Required import
user_profile_aggregates	Per-user aggregate blueprint and component summary rows.	Required
pair_computations	Persisted CompatibilityScore outputs and norms.	Required
pair_diagnostics	Top drivers, friction notes, and internal diagnostics.	Required
access_decisions	Direct-message vs Knock decision records.	Required
knocks	Restricted-access intent requests.	Required
conversations	Conversation shell once access opens.	Required minimal
messages	Minimal message records for live chat shell and telemetry foundation.	Required minimal
event_ledger	Append-only event record for scoring, matching, access, messaging, and reflection actions.	Required
feedback_outcomes	Outcome and checkpoint capture aligned to FeedbackData sheet.	Required launch capture
feature_store / genome_features / qm_*	Later enrichment, AI, and ML modules.	Deferred / optional skeleton only

6.1 users

Stores the platform identity record used by onboarding, matching, and access routing. Bubble or app-layer profile objects may exist, but this contract assumes a backend canonical user row.

Field	Type	Req	Notes
user_id	uuid	Yes	Primary user key used across all Phase 1 tables.

Field	Type	Req	Notes
rs_code	text(6)	Yes	Searchable identity handle; unique.
account_state	text/enum	Yes	assessment_incomplete, eligible_in_pool, knock_pending, conversation_active, in_relationship, post_breakup_recovery, ghost_mode.
relationship_state	text/enum	Yes	single, talking, matched, in_relationship, ended, paused.
date_of_birth	date	No	Used for age checks and astrology/numerology derivation if permitted.
gender	text	No	Used for configured overlay keys where product logic requires it.
location_city	text	No	Store city/metro, not exact address.
search_radius_miles	int	No	Candidate retrieval only.
dating_status	text/enum	No	Talking, Friends, Dating, Maybe long term, Marriage.
ghost_mode_flag	bool	Yes	Suppresses public surfacing while preserving self tools.
created_at	timestamptz	Yes	UTC record creation time.
updated_at	timestamptz	Yes	UTC record update time.

Implementation note

- users.rs_code must be unique. Do not let vanity, premium, or cosmetic identity affect score logic or ranking.

6.2 consent_state

Stores the policy gates required before AI reads content, before chat-derived features are saved, and before pair-level or shared content is exposed.

Field	Type	Req	Notes
consent_id	uuid	Yes	Primary row id.
user_id	uuid	Yes	FK to users.
ai_consent	bool	Yes	Master opt-in for AI interpretation features.
chat_analysis_consent	bool	Yes	Required before message content is used for derived features.
sanctuary_analysis_consent	bool	Yes	Required before Sanctuary free text is used by AI.
data_use_consent	bool	Yes	Legal/analytics consent distinct from feature consent.
match_visibility_mode	text/enum	Yes	public_pool, hidden, invite_only, ghost.
consent_scope_json	jsonb	Yes	Resolved policy payload for packet builders and AI services.

Field	Type	Req	Notes
policy_version	text	Yes	Version tag for the active policy rules.
updated_at	timestamptz	Yes	UTC update time.

6.3 assessment_sessions

Tracks a single assessment run or retake window. This table is the audit wrapper around user_answers.

Field	Type	Req	Notes
assessment_session_id	uuid	Yes	Primary assessment session id.
user_id	uuid	Yes	FK to users.
question_version	text	Yes	Question-bank version used for this session.
taxonomy_version	text	Yes	Canonical taxonomy version used for labels and clusters.
scoring_version	text	Yes	Scoring contract version in force at submission time.
status	text/enum	Yes	in_progress, completed, abandoned.
total_answers_expected	int	Yes	Default 60 for Phase 1 launch.
total_answers_received	int	Yes	Must equal 60 before the session is completed.
started_at	timestamptz	Yes	UTC start time.
submitted_at	timestamptz	No	UTC completion time; null until submitted.
created_at	timestamptz	Yes	Audit timestamp.

Implementation note

- Assessment completion should not surface the user into the candidate pool until required answers are complete and the score service has produced initial outputs.

6.4 user_answers

Stores the selected answer rows for each assessment question. This table should link to imported workbook rows by answer_id, not by raw answer text.

Field	Type	Req	Notes
user_answer_id	uuid	Yes	Primary selected-answer record.
assessment_session_id	uuid	Yes	FK to assessment_sessions.
user_id	uuid	Yes	FK to users.
question_number	int	Yes	Stable question key; 1-50 core and 100-109 shadow if mirroring workbook numbering.

Field	Type	Req	Notes
answer_id	text	Yes	Canonical workbook answer id such as Q01_A.
answer_letter	text	No	Convenience field for UI / QA traceability.
phase	text/enum	Yes	Core or Shadow.
question_weight_v3	numeric(4,2)	Yes	Copied from imported AnswerBank for auditability.
answer_text_snapshot	text	No	Optional display snapshot for historical audit only.
submitted_at	timestamptz	Yes	UTC answer submission time.
scoring_version	text	Yes	Version tag for parity and rescore audits.

Implementation note

- Enforce uniqueness on (assessment_session_id, question_number). Never let the client invent answer ids.

6.5 workbook_import_jobs

Tracks every launch import batch of workbook-owned data.

Field	Type	Req	Notes
import_job_id	uuid	Yes	Primary import batch id.
workbook_version	text	Yes	Human-readable workbook release tag.
score_contract_version	text	Yes	Schema/build contract version paired to this import.
source_file_name	text	Yes	Original workbook filename.
source_checksum	text	Yes	Checksum of workbook file used for this batch.
imported_by	text	Yes	Service, actor, or environment executing import.
imported_at	timestamptz	Yes	UTC time import finished.
active_flag	bool	Yes	Exactly one launch batch should be active at a time.
notes	text	No	Operational import notes.

6.6 workbook_sheet_registry

Stores one row per imported sheet so row counts, sheet-level checksum behavior, and active-state decisions can be audited cleanly.

Field	Type	Req	Notes
sheet_registry_id	uuid	Yes	Primary row id.
import_job_id	uuid	Yes	FK to workbook_import_jobs.

Field	Type	Req	Notes
sheet_name	text	Yes	Workbook sheet name.
sheet_rows	int	Yes	Physical sheet row count including header.
data_rows	int	Yes	Importable row count excluding header.
active_flag	bool	Yes	Whether this sheet import is the active runtime artifact.
created_at	timestamptz	Yes	UTC audit timestamp.

6.7 answer_bank_rows

Imported AnswerBank sheet. This table is the atomic answer registry and should be loaded from the workbook, not written by hand. Full field list is in Appendix A.

Field	Type	Req	Notes
answer_id	text	Yes	Primary workbook answer key such as Q01_A.
qa_key	text	Yes	Question/answer join key from workbook.
question_number_int	int	Yes	Stable question number.
phase	text	Yes	Core or Shadow.
question_weight_v3	numeric(4,2)	Yes	Official question weight.
answer_letter	text	Yes	Rendered answer slot identifier.
answer_text	text	Yes	Human-facing answer text.
archetype	text	No	Primary archetype classification anchor.
shadow_type	text	No	Primary shadow classification anchor when present.
dimension_primary	text	Yes	Workbook primary dimension family.
dimension_secondary	text	No	Workbook secondary dimension family.
import_job_id	uuid	Yes	FK to workbook import batch.
active_flag	bool	Yes	Runtime-active imported row.

Implementation note

- All numeric norm, risk, band, and behavior-label fields from the sheet must be preserved exactly. Do not collapse the answer bank into a reduced table.

6.8 answer_pair_quantum_rows

Imported AnswerPairQuantum sheet. This is the canonical pair-lookup artifact and the core of launch scoring. Full field list is in Appendix B.

Field	Type	Req	Notes
pair_key	text	Yes	Sorted answer pair key in the form ans_low|ans_high.
ans_a_id	text	Yes	First answer id from workbook row.
ans_b_id	text	Yes	Second answer id from workbook row.
q_a	int	Yes	Question number for answer A.
q_b	int	Yes	Question number for answer B.
phase_a	text	Yes	Core or Shadow.
phase_b	text	Yes	Core or Shadow.
risk_avg	numeric(6,4)	Yes	Average pair risk input used by gate logic.
risk_max	numeric(6,4)	Yes	Maximum pair risk input used by gate logic.
gate	numeric(6,4)	Yes	Workbook-defined gate value.
rule_penalty	numeric(6,4)	Yes	Workbook-defined rule penalty.
pair_norm_0_1	numeric(6,4)	Yes	Canonical normalized pair result after gate and penalty.
import_job_id	uuid	Yes	FK to workbook import batch.
active_flag	bool	Yes	Runtime-active imported row.

Implementation note

- Never recreate or simplify this table in client code. MVP scoring must look up these rows from backend-owned storage.

6.9 archetype_matrix_rows

Normalized storage of the 10x10 archetype matrix. Teams may also cache this in service config, but the import should remain auditable.

Field	Type	Req	Notes
archetype_matrix_row_id	uuid	Yes	Primary row id.
archetype_a	text	Yes	Left-hand archetype key.
archetype_b	text	Yes	Right-hand archetype key.
fit_value	numeric(6,4)	Yes	Imported matrix fit value.
import_job_id	uuid	Yes	FK to workbook import batch.

6.10 shadow_matrix_rows

Normalized storage of the 6x6 shadow interaction matrix.

Field	Type	Req	Notes
shadow_matrix_row_id	uuid	Yes	Primary row id.

Field	Type	Req	Notes
shadow_a	text	Yes	Left-hand shadow key.
shadow_b	text	Yes	Right-hand shadow key.
fit_value	numeric(6,4)	Yes	Imported matrix fit value.
import_job_id	uuid	Yes	FK to workbook import batch.

6.11 zodiac_lookup_rows

Zodiac overlay lookup imported directly from workbook.

Field	Type	Req	Notes
z_key	text	Yes	Primary zodiac lookup key.
gender_a	text	Yes	Overlay key component.
gender_b	text	Yes	Overlay key component.
sun_sign_1	text	Yes	Overlay key component.
sun_sign_2	text	Yes	Overlay key component.
score_1_10	numeric(4,2)	Yes	Raw workbook score.
z_norm	numeric(6,4)	Yes	Normalized overlay score.
import_job_id	uuid	Yes	FK to workbook import batch.

6.12 numerology_lookup_rows

Numerology overlay lookup imported directly from workbook.

Field	Type	Req	Notes
n_key	text	Yes	Primary numerology lookup key.
lp_a	int	Yes	Life path A.
lp_b	int	Yes	Life path B.
score_1_10	numeric(4,2)	Yes	Raw workbook score.
weight_points	numeric(4,2)	No	Optional raw weighting points from workbook.
n_norm	numeric(6,4)	Yes	Normalized overlay score.
import_job_id	uuid	Yes	FK to workbook import batch.

6.13 polarity package import tables

Phase 1 should import the chemistry package so the workbook-owned polarity subsystem is present for auditability and future chemistry surfaces. The score engine may rely on AnswerPairQuantum for launch outputs, but the polarity import package should still be versioned and available.

Table	Key fields to preserve	Why import now
polarity_axes_v10	Axis, Field, Fit Type, Inside Chemistry Weight, Definition, Why it matters	Locks the chemistry dimensions and weights used by the workbook package.
polarity_role_seeds_v10	Role, Role Label, axis targets, anchor labels	Supports chemistry auditability and future explanation surfaces.
polarity_question_map_v10	question_id, block_id, block_multiplier, question_weight, primary_axes, secondary_axes, notes	Provides traceability from chemistry prompts to axis logic.
polarity_answer_bank_v10	question_id, answer_slot, redesigned_role, axis payloads, polarity_norm_seed, attraction_pattern_norm_seed	Preserves per-answer chemistry vectors and seeds.
rewrite_rules_v10	Category, Guidance, Practical effect	Reference-only content rules; no runtime scoring logic.
polarity_pairengine_v10_reference	dimension, fit formula, fit score examples	Reference implementation and QA artifact for chemistry calculations.

6.14 user_profile_aggregates

Stores the most recent per-user aggregate blueprint computed from selected answer rows. This is a backend-owned derived object and should be recomputed deterministically from current assessment data.

Field	Type	Req	Notes
user_id	uuid	Yes	PK/FK to users.
latest_assessment_session_id	uuid	Yes	FK to assessment_sessions.
primary_archetype_id	text	No	Derived dominant archetype label.
secondary_archetype_id	text	No	Optional second archetype.
primary_shadow_type_id	text	No	Derived dominant shadow type; internal by default.
secondary_shadow_type_id	text	No	Optional secondary shadow type.
communication_norm	numeric(6,4)	Yes	Aggregate from AnswerBank norms.
regulation_norm	numeric(6,4)	Yes	Aggregate from AnswerBank norms.
love_norm	numeric(6,4)	Yes	Aggregate from AnswerBank norms.
readiness_norm	numeric(6,4)	Yes	Aggregate from AnswerBank norms.
pacing_norm	numeric(6,4)	Yes	Aggregate from AnswerBank norms.
attachment_risk_norm	numeric(6,4)	Yes	Aggregate from AnswerBank norms.
integrity_burden_norm	numeric(6,4)	Yes	Aggregate from AnswerBank penalty fields.

Field	Type	Req	Notes
cluster_burden_norm	numeric(6,4)	Yes	Aggregate from AnswerBank penalty fields.
forecast_signal_norm	numeric(6,4)	Yes	Aggregate forecast signal from AnswerBank.
baseline_polarity_profile_json	jsonb	No	Reserved for future chemistry surfaces and explanations.
profile_confidence_score	numeric(6,4)	No	Internal quality score for current blueprint stability.
scoring_version	text	Yes	Version tag for reproducibility.
updated_at	timestamptz	Yes	UTC recompute time.

6.15 pair_computations

Persists the canonical match output for a user pair. This table is what the app reads after backend score computation.

Field	Type	Req	Notes
pair_id	uuid	Yes	Primary pair computation id.
pair_user_key	text	Yes	Sorted user key: low_user_id|high_user_id.
user_a_id	uuid	Yes	User A id.
user_b_id	uuid	Yes	User B id.
source_assessment_session_a	uuid	Yes	Assessment snapshot used for A.
source_assessment_session_b	uuid	Yes	Assessment snapshot used for B.
compatibility_score	int	Yes	Public 350-850 score.
compatibility_tier	text/enum	Yes	Tier derived from score and policy.
final_norm	numeric(6,4)	Yes	Clamped final norm.
core_norm	numeric(6,4)	Yes	Weighted aggregate of pair_norm rows before cosmic overlay.
stability_avg	numeric(6,4)	Yes	Aggregate stability component for diagnostics.
chemistry_avg	numeric(6,4)	Yes	Aggregate chemistry component for diagnostics.
zodiac_norm	numeric(6,4)	No	Derived zodiac overlay.
numerology_norm	numeric(6,4)	No	Derived numerology overlay.
scoring_version	text	Yes	Required for parity and audit.
computed_at	timestamptz	Yes	UTC compute time.

Implementation note

- Enforce order-insensitive pair storage using pair_user_key. The same two users should not create duplicate active pair rows because of order.

6.16 pair_diagnostics

Stores translated or internal driver payloads linked to a pair computation. Diagnostics should be separate from the score row so visibility filters stay clean.

Field	Type	Req	Notes
pair_diagnostic_id	uuid	Yes	Primary row id.
pair_id	uuid	Yes	FK to pair_computations.
top_positive_drivers_json	jsonb	No	Translated strengths or major fit contributors.
top_friction_drivers_json	jsonb	No	Translated friction contributors.
active_pair_rules_json	jsonb	No	Internal list of pair rules or rule ids applied.
active_cluster_hits_json	jsonb	No	Internal cluster or contradiction hits; not match-visible raw.
chemistry_profile_json	jsonb	No	Chemistry summary payload if surfaced.
diagnostics_visibility_mask	jsonb	Yes	Controls what can be exposed to self, match, admin, or nobody.
created_at	timestamptz	Yes	UTC write time.

6.17 access_decisions

Stores the direct-message vs Knock gating decision that follows score computation.

Field	Type	Req	Notes
access_decision_id	uuid	Yes	Primary row id.
pair_id	uuid	Yes	FK to pair_computations.
compatibility_score	int	Yes	Copied for audit snapshot.
compatibility_tier	text/enum	Yes	Copied for audit snapshot.
direct_message_eligible	bool	Yes	Whether DM is immediately open.
knock_required_flag	bool	Yes	Whether Knock is required before messaging.
decision_reason_code	text	Yes	Policy reason used by backend.
policy_version	text	Yes	Policy version in force at decision time.
created_at	timestamptz	Yes	UTC decision time.

6.18 knocks

Minimal restricted-access request object for Phase 1. This is enough to support gated outreach before full relationship features exist.

Field	Type	Req	Notes
knock_id	uuid	Yes	Primary row id.
pair_id	uuid	Yes	FK to pair_computations.

Field	Type	Req	Notes
sender_user_id	uuid	Yes	User sending the Knock.
recipient_user_id	uuid	Yes	User receiving the Knock.
status	text/enum	Yes	pending, accepted, declined, expired.
message_text	text	No	Optional short note if product policy allows.
created_at	timestamptz	Yes	UTC send time.
responded_at	timestamptz	No	UTC response time.

6.19 conversations

Minimal conversation container for Phase 1 messaging.

Field	Type	Req	Notes
conversation_id	uuid	Yes	Primary conversation id.
pair_id	uuid	Yes	FK to pair_computations.
status	text/enum	Yes	active, closed, blocked, archived.
created_at	timestamptz	Yes	UTC create time.
last_message_at	timestamptz	No	UTC timestamp of latest message.

6.20 messages

Minimal message record for launch chat shell and telemetry foundation.

Field	Type	Req	Notes
message_id	uuid	Yes	Primary message id.
conversation_id	uuid	Yes	FK to conversations.
sender_user_id	uuid	Yes	Message sender.
body_text	text	Yes	Message body.
redaction_state	text/enum	No	raw, redacted, masked if moderation or policy requires.
created_at	timestamptz	Yes	UTC send time.
content_hash	text	No	Optional content hash/pointer for later telemetry and privacy-safe references.

6.21 event_ledger

Append-only event record that allows replay, audit, drift foundations, and post-launch validation. This table ships in Phase 1 even if advanced behavior models do not.

Field	Type	Req	Notes
event_id	uuid	Yes	Global unique event id.

Field	Type	Req	Notes
event_type	text/enum	Yes	answer_submitted, assessment_scored, match_computed, knock_sent, message_sent, consent_changed, etc.
module_name	text/enum	Yes	assessment, scoring, matching, messaging, system, policy, etc.
user_id_or_pair_id	text	Yes	Scope id depending on event subject.
actor_user_id	uuid	No	Optional acting user.
subject_user_id	uuid	No	Optional secondary user.
state_pre_json	jsonb	No	Before-state payload when relevant.
state_post_json	jsonb	No	After-state payload when relevant.
payload_json	jsonb	Yes	Typed event payload.
taxonomy_version	text	No	Version tag if taxonomy-bound.
scoring_version	text	No	Version tag if score-bound.
created_at	timestamptz	Yes	UTC write time.

Implementation note

- The event ledger is the launch foundation for later drift, forecasting, and calibration. It is not optional if the team wants a real system instead of a one-time quiz.

6.22 feedback_outcomes

Phase 1 launch-capture table aligned to FeedbackData. Even if tuning is later, the data capture structure should exist now.

Field	Type	Req	Notes
outcome_id	uuid	Yes	Primary row id.
user_id	uuid	Yes	Owner of the checkpoint record.
pair_id	uuid	No	Linked pair or match id.
match_score	int	No	CompatibilityScore snapshot at time of match.
tier_label	text	No	Tier label snapshot.
date_matched	date	No	Match creation date.
message_count_7d	int	No	7-day engagement proxy.
reply_latency_p50	text or numeric	No	Median reply latency metric.
date_first_meet	date	No	Offline conversion marker if captured.
relationship_status_30d	text	No	30-day outcome.
relationship_status_90d	text	No	90-day outcome.
breakup_reason	text	No	Coarse reason bucket if available.

Field	Type	Req	Notes
conflict_events_count	int	No	Observed conflict count.
repair_events_count	int	No	Observed repair count.
user_reported_satisfaction	int	No	Self-report label 1-10.
user_reported_safety	int	No	Self-report label 1-10.
user_reported_chemistry	int	No	Self-report label 1-10.
flag_harmful_behavior	bool	No	Serious safety label for later calibration.
collected_at	timestamptz	Yes	UTC collection time.

7. Minimal Packet Contracts for Phase 1

The launch system should move data between services using narrow, versioned packets rather than letting services read arbitrary tables. This keeps the implementation deterministic and reduces backend drift.

Packet	Required contents	Primary consumer
Self Packet	user_id, demographics keys, latest assessment session id, selected answer_ids, aggregate component norms, archetype/shadow baselines, visibility mask, scoring_version	Scoring, reporting, admin QA
Pair Packet	Two Self Packets + pair_user_key + policy mask + allowed overlay keys	Matching / scoring service
Access Packet	pair_id, score, tier, access policy version, reveal rules	Messaging / Knock gate
Training Capture Packet	pair_id, score snapshot, tier, feedback checkpoint fields, event references	Analytics later; capture now

8. Symmetry, Uniqueness, and Index Rules

Rule	Implementation requirement
Answer pair symmetry	pair_key must always be built from sorted answer ids before lookup; the same answer pair may not have duplicate active runtime keys.
User pair symmetry	pair_user_key must always be built from sorted user ids before score persistence or access decisions.
Assessment uniqueness	One user may have multiple assessment sessions over time, but only one active launch session should drive current match calculations.
Imported row immutability	Imported workbook rows should be append-only by import job. Switching imports should happen by active_flag, not by editing rows in place.
Current pair materialization	If the platform stores historical pair computations, mark one row as latest/current per pair_user_key and scoring_version.

- Recommended unique indexes: users(rs_code), user_answers(assessment_session_id, question_number), answer_bank_rows(import_job_id, answer_id), answer_pair_quantum_rows(import_job_id, pair_key), pair_computations(pair_user_key, computed_at), access_decisions(pair_id, created_at).
- Recommended query indexes: user_answers(user_id, submitted_at), pair_computations(user_a_id), pair_computations(user_b_id), knocks(recipient_user_id, status), messages(conversation_id, created_at), event_ledger(event_type, created_at).
- If the team is using Bubble-backed tables, the same logic still applies: sorted keys, immutable imports, and no page-level score computation.

9. What is explicitly not in Phase 1 schema law

- No client-side score fields, helper formulas, or page-workflow re-creations of workbook logic.
- No requirement to fully ship Quantum Mirror, Sanctuary module family, full genome extraction, or advanced drift state engines at launch.
- No monetization or cosmetic table may alter score, tier, rank, or access thresholds.

- No raw shadow, raw safety, or raw contradiction fields should be match-visible.
- No free-form schema additions that reinterpret the score package without an approved version update.

What 'test later' is allowed to mean

- Import parity checks against workbook row counts and checksums.
- Deterministic MatchHarness parity checks against known score outputs.
- Post-launch distribution review, outcome validation, and calibration using feedback captures.
- It does not mean using a fake MVP formula first and replacing it later.

10. Phase 1 Acceptance Criteria

- Engineers can import the workbook into versioned tables and verify row counts without hard-coded magic numbers.
- Assessment answers are stored by canonical answer_id and can be re-scored reproducibly.
- Matching service computes score from backend-owned AnswerPairQuantum lookups, question weights, gate logic, and penalties.
- The same two users receive the same deterministic score when the same assessment snapshots, imports, and version tags are used.
- Access decisions are stored separately from score rows and are driven by score/tier plus policy, not by UI interpretation.
- Event ledger records the minimum launch events needed for audit and later calibration.
- Feedback/outcome capture structure exists even if the first launch does not retrain any models.
- No old formula, simulation sheet, or frontend convenience layer is used as score law.

Appendix A. Full AnswerBank field contract

The table below mirrors the imported AnswerBank headers from the workbook. Suggested types are implementation guidance; the workbook remains the source of truth for field names and meaning.

Field	Suggested type	Purpose
answer_id	text	Canonical workbook answer row id.
qa_key	text	Question/answer join key.
question_number_int	int	Stable question number.
phase	text	Core or Shadow.
question_weight_v3	numeric(4,2)	Official question weight for pair aggregation.
dimension_primary	text	Primary workbook dimension family.
dimension_secondary	text	Secondary workbook dimension family.
answer_letter	text	Answer slot identifier.
answer_text	text	Human-facing answer text.
archetype	text	Primary archetype mapping.
shadow_type	text	Shadow mapping when present.
comm_health_norm	numeric(6,4)	Communication health norm.
reg_health_norm	numeric(6,4)	Regulation health norm.
love_relational_norm	numeric(6,4)	Love/relational norm.
readiness_norm	numeric(6,4)	Readiness norm component.
pacing_norm	numeric(6,4)	Pacing norm.
attachment_risk_norm	numeric(6,4)	Attachment risk norm.
centered_polarity_norm	numeric(6,4)	Legacy compatibility-friendly chemistry norm.
reciprocity_signal_norm	numeric(6,4)	Reciprocity signal norm.
integrity_penalty_norm	numeric(6,4)	Integrity penalty burden.
cluster_penalty_norm	numeric(6,4)	Cluster contradiction burden.
pattern_risk_norm	numeric(6,4)	Pattern risk burden.
forecast_norm	numeric(6,4)	Forecast signal norm.
pos_agg	numeric(6,4)	Positive aggregate band support.
chem_agg	numeric(6,4)	Chemistry aggregate band support.
risk_agg	numeric(6,4)	Risk aggregate band support.
risk_band	text	Risk band label.
pos_band	text	Positive band label.
chem_band	text	Chemistry band label.
access_style_label_v8	text	Access style narrative label.
conflict_style_label_v8	text	Conflict style narrative label.
repair_style_label_v8	text	Repair style narrative label.
attachment_style_label_v8	text	Attachment style narrative label.
nervous_system_bias_label_v8	text	Nervous-system bias label.
selection_gap_flag_v8	text/bool	Content QA or voice-balance flag.

Appendix B. Full AnswerPairQuantum field contract

The table below mirrors the imported AnswerPairQuantum headers from the workbook. These fields should be preserved exactly so the runtime engine, QA harness, and audit trail all speak the same language.

Field	Suggested type	Purpose
pair_key	text	Sorted answer pair key.
ans_a_id	text	Answer A id.
ans_b_id	text	Answer B id.
q_a	int	Question number for answer A.
q_b	int	Question number for answer B.
phase_a	text	Core or Shadow.
phase_b	text	Core or Shadow.
archetype_fit	numeric(6,4)	Archetype fit component.
comm_fit	numeric(6,4)	Communication fit component.
reg_fit	numeric(6,4)	Regulation fit component.
love_fit	numeric(6,4)	Love fit component.
readiness_fit	numeric(6,4)	Readiness alignment component.
shadow_fit	numeric(6,4)	Shadow interaction component.
integrity_fit	numeric(6,4)	Inverse integrity burden component.
cluster_fit	numeric(6,4)	Inverse cluster burden component.
pace_fit	numeric(6,4)	Inverse pace mismatch component.
polarity_fit	numeric(6,4)	Chemistry fit component.
pursuit_fit	numeric(6,4)	Pursuit/receptivity component.
forecast_fit	numeric(6,4)	Pair stability / forecast signal.
risk_avg	numeric(6,4)	Average pair risk.
risk_max	numeric(6,4)	Maximum pair risk.
gate	numeric(6,4)	Workbook gate output.
rule_penalty	numeric(6,4)	Rule penalty output.
behavioral_component	numeric(6,4)	Behavioral subtotal.
stability_component	numeric(6,4)	Stability subtotal.
chemistry_component	numeric(6,4)	Chemistry subtotal.
base_0_88	numeric(6,4)	Pre-gate subtotal.
pair_norm_0_1	numeric(6,4)	Canonical normalized pair result.

Appendix C. FeedbackData launch capture fields

Field	Type	Example	Why it matters
user_id	string	u_123	Join data

Field	Type	Example	Why it matters
match_id	string	m_987	Join data
match_score	number	724	Calibration target
tier_label	string	Strong potential	Calibration target
date_matched	date	2026-03-20	Time decay / cohort effects
message_count_7d	number	42	Engagement proxy
reply_latency_p50	number	12m	Responsiveness proxy
date_first_meet	date	2026-04-02	Conversion
relationship_status_30d	enum	talking/dating/ended	Outcome
relationship_status_90d	enum	dating/ended	Outcome
breakup_reason	enum	trust/conflict/pace/values/ unknown	Penalty tuning
conflict_events_count	number	3	Stability truth signal
repair_events_count	number	2	Repair truth signal
user_reported_satisfaction	1-10	8	Supervised label
user_reported_safety	1-10	7	Supervised label
user_reported_chemistry	1-10	9	Chemistry calibration
flag_harmful_behavior	bool	false	Hard penalty learning

Appendix D. Deferred or reserved tables after launch

Table family	Phase	Reason to defer
sanctuary_logs and state packets	Phase 1.5 / 2	Not required for launch scoring, matching, or gated messaging.
genome_features / feature_store	Phase 2+	Needed for richer AI, ML, and narrative enrichment after launch.
drift_states / shadow_activation	Phase 2+	Useful later; not required to implement the launch pair engine correctly.
qm_* tables	Phase 2+	Quantum Mirror memory, extraction, and moderation stack is beyond launch MVP.
relationship_mode / couple_rooms / symphonic artifacts	Later phases	Post-match relationship tooling is intentionally deferred.